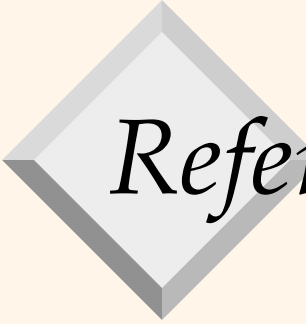


# *SQL Intro*



# *References*

- [RG] Sec 3.1-3.4, 5.1-5.2




# *Admin info*

☐ Please make sure you know how to find CMS and Piazza

- [cms.csuglab.cornell.edu](http://cms.csuglab.cornell.edu)
- [piazza.com](http://piazza.com)

☐ Quiz on course policies due 13th Feb



# *Relational model*

- ❑ Mathematical abstraction
- ❑ Abstraction provided by a database such as MySQL
- ❑ Closely related, used interchangeably, but technically not the same
  - Differences?
    - ❑ Duplicates
    - ❑ Ordering



# *Querying Relational Data*

- ❑ Several query abstractions/language
- ❑ **SQL** (Structured Query Language)
- ❑ **Relational Algebra**
  - Used as intermediate representation in your SQL engine/DBMS



# *Creating Relations in SQL*

- ▶ Creates Students relation
  - Type (domain) of each field is specified
  - Enforced by DBMS whenever tuples are added or modified
- ▶ Enrolled table holds information about courses that students take

```
CREATE TABLE Students  
  (sid CHAR(20),  
   name CHAR(20),  
   login CHAR(10),  
   age INT,  
   gpa REAL);
```

```
CREATE TABLE Enrolled  
  (sid CHAR(20),  
   cid CHAR(20),  
   grade CHAR(2));
```



# *Destroying and Altering Relations*

**DROP TABLE** Students;

Destroys the relation Students

- Schema information and tuples are deleted

**ALTER TABLE** Students

**ADD COLUMN** nationality CHAR(30);

Schema of Students is altered by adding a new field

- Every tuple in current instance is extended with a null value in the new field
- Type DESCRIBE Students to see schema



# *Adding and Deleting Tuples*

❑ Can insert a single tuple using:

```
INSERT INTO Students  
VALUES (12345, 'Smith', 'smith@ee', 18, 3.2, 'US');
```

❑ Can delete all tuples satisfying some condition  
(e.g., name = Smith):


```
DELETE FROM Students  
WHERE name = 'Smith';
```





# *Integrity Constraints (ICs)*

- ❑ Conditions that must be true for any instance of the database
- ❑ Specified when schema is defined
  - or with ALTER statement
- ❑ Enforced by the DBMS



# Primary Key Constraints

- ❑ A set of fields is a key for a relation if:
  1. No two distinct tuples can have same values in all key fields, and
  2. This is not true for any subset of the key
    - ❑ Part 2 false? A superkey
- ❑ If there is >1 key for a relation, one of the keys is chosen (by DBA) to be the primary key

ALTER TABLE STUDENTS ADD PRIMARY KEY (SID);



# Foreign Keys

- ❑ Foreign key: set of fields in one relation that is used to "refer" to a tuple in another relation
  - Usually must correspond to primary key of second relation
- ❑ Enrolled(sid: string, cid: string, grade: string)
  - Any sid in Enrolled must also be in Students
    - ❑ sid in Enrolled is a foreign key referencing sid in students
- ❑ If all foreign key constraints are enforced, referential integrity is achieved

# Foreign Keys in SQL

☐ Only students listed in the Students relation should be allowed to enroll for courses


```
CREATE TABLE Enrolled
(sid CHAR(20), cid CHAR(20), grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid) REFERENCES Students (sid) );
```


## Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

## Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8





# *Enforcing Referential Integrity*

- ❑ What should be done if an Enrolled tuple with a non-existent student id is inserted?
  - Reject it
- ❑ What should be done if a Students tuple is deleted?
  - Also delete all Enrolled tuples that refer to it
  - Disallow deletion of a Students tuple that is referred to
  - Set sid in Enrolled tuples that refer to it to a default sid.
  - Set sid in Enrolled tuples that refer to it to a special value null, denoting 'unknown' or 'inapplicable'
- ❑ Similar if primary key of Students tuple is updated



# *Referential Integrity in SQL*

► 4 options on deletes and updates

- Default is **NO ACTION** (delete/update is rejected)
- **CASCADE** (also delete all tuples that refer to deleted tuple)
- **SET NULL / SET DEFAULT** (sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE
ON UPDATE CASCADE);
```



☐ Relations

☐ Constraints:

- Attribute types
- Keys
- Foreign Keys

☐ CREATE, ALTER, INSERT, DELETE



# *Queries*

- ❑ Now, how do we get data out of the database?
- ❑ Using SQL queries
- ❑ Lots of very powerful features





# *Running example - Sailors and Boats*

- ❑ Sailors in a club reserve Boats
- ❑ Tables demo
- ❑ Example data in CMS

# *Basic SQL Query*

```
SELECT  S.sname  
FROM    Sailors S  
WHERE   S.age > 20;
```

```
SELECT DISTINCT S.sname  
FROM    Sailors S  
WHERE   S.age > 20;
```

- Default is that duplicates are not eliminated!
  - Need to explicitly say “DISTINCT”



# *Range variables/aliases*

- ❑ Not required but considered "good practice"
  - Especially once we get to multi-relation queries
- ❑ All three below queries are identical

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.age > 20;
```


```
SELECT  sname
FROM    Sailors
WHERE   age > 20;
```

```
SELECT  Sailors.sname
FROM    Sailors
WHERE   Sailors.age > 20;
```

# Basic SQL Query


```
SELECT      [DISTINCT] target-list  
[FROM      relation-list]  
[WHERE      condition]
```

- No FROM-clause often not supported
- SELECT \* returns all attributes



# SQL Query

```
SELECT S.sname  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid AND R.bid=101;
```



# Conceptual Evaluation Strategy

- Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
  - Compute the cross-product of *relation-list*
  - Discard resulting tuples if they fail *condition*.
  - Delete attributes that are not in *target-list*
  - If DISTINCT is specified, eliminate duplicate rows.
- This strategy is probably the least efficient way to compute a query!
  - An optimizer will find more efficient strategies to compute *the same answers*.

# JOINS

- ❑ A common operation: find all "matching" pairs of tuples from two relations

```
SELECT * FROM SAILORS S, RESERVES R WHERE R.sid = S.sid;
```

- ❑ Our example query can be thought of as a two step process:
  - JOIN Sailors and Reserves
  - Retain only tuples where bid = 101;



# *More JOIN syntax*

- ❑ Other ways to compute a JOIN between R and S in MySQL and PostgreSQL:


```
SELECT * FROM SAILORS S JOIN RESERVES R ON (S.sid = R.sid);
```

```
SELECT * FROM SAILORS JOIN RESERVES USING (sid);
```

```
SELECT * FROM SAILORS NATURAL JOIN RESERVES;
```

- ❑ But beware of the NATURAL JOIN
- Make sure you really know what the columns with the same names are!
  - And some systems (notably MS SQL Server) don't support it





# *Expressions and Strings*

```
SELECT  S.age, S.age > 30 AS isOver30, 2*S.age AS age2  
FROM    Sailors S  
WHERE   S.sname LIKE 'B_%B';
```

- *Find triples (of ages of sailors and two fields defined by expressions) for sailors whose names begin and end with B and contain at least three characters.*
- **AS** is used to name fields in result.
- **LIKE** is used for string matching
  - **`\_`** stands for any one character
  - **`%`** stands for 0 or more arbitrary characters.