

*SQL wrap-up; NULLs*  
*Relational Algebra intro*

# *Last time*

## ❖ Basic SQL features

- AND/OR, UNION/INTERSECT/EXCEPT

## ❖ Advanced SQL features

- Nested queries and correlation
- Aggregates (MAX, COUNT)
- GROUP BY and HAVING

# *Null Values*

- ❖ Represents values that are unknown (e.g., grade has not been assigned) or inapplicable (e.g., no SSN for a non-US-resident)
  - SQL provides a special value **NULL**
  - Can specify a field cannot be null using NOT NULL in definition

sid	sname	ssn
11	alice	123
22	bob	NULL

# *Null Values*

- ❖ Semantics of operators must be designed carefully
  - 3-valued logic (true, false, unknown)

SELECT 3 = NULL;

SELECT NULL = NULL;

SELECT NULL IS NULL;

SELECT NULL IS NOT NULL;

SELECT TRUE OR NULL;

SELECT TRUE AND NULL;

SELECT \* FROM Sailors2 S WHERE S.ssn IS NULL;

SELECT \* FROM Sailors2 S WHERE S.ssn > 124;

SELECT \* FROM Sailors2 S WHERE S.ssn <= 124;

# Outer Joins

- ❖ Suppose we join Sailors and Reserves, but some sailors have not made a reservation
- ❖ Can imagine wanting to retain all Sailors tuples in the result
- ❖ Can do this with an OUTER JOIN

```
SELECT * FROM SAILORS LEFT OUTER JOIN RESERVES USING (sid);
```

- ❖ Also available: right outer join, full outer join
  - What do you think they do?

# *Outer Join Example*


- ❖ For each sailor, display their id, name and how many reservations they have made in total
  - Including 0 if no reservations;

# *Outer Join Example*


```
SELECT sid, sname, count(bid)
FROM (SELECT * FROM SAILORS JOIN
      RESERVES USING (sid)) AS TEMP
GROUP BY sid, sname;
```

```
SELECT sid, sname, count(bid)
FROM (SELECT * FROM SAILORS LEFT OUTER JOIN
      RESERVES USING (sid)) AS TEMP
GROUP BY sid, sname;
```

## *Be careful with NULLs!*



```
SELECT sid, sname, count(bid)
FROM (SELECT * FROM SAILORS LEFT OUTER JOIN
      RESERVES USING (sid)) AS TEMP
GROUP BY sid, sname;
```



```
SELECT sid, sname, count(*)
FROM (SELECT * FROM SAILORS LEFT OUTER JOIN
      RESERVES USING (sid)) AS TEMP
GROUP BY sid, sname;
```



# *SQL Summary:*

- ❖ Basic SELECT/FROM/WHERE queries
- ❖ Expressions and strings
- ❖ Set operators
- ❖ Nested queries
- ❖ Aggregation
- ❖ GROUP BY/HAVING
- ❖ Null values and Outer Joins
- ❖ (ORDER BY and other features...)

# *How to write complex SQL queries?*

- ❖ Gets easier with practice
- ❖ Look at lots of examples (eg in your textbook)
  - Including the textbook exercises
  - Solutions to odd-numbered exercises freely available on textbook's website <http://pages.cs.wisc.edu/~dbbook/>

# *How to write complex SQL queries?*

## ❖ Bottom-up strategy

- Can I write anything that is even slightly related to the query I need?
- Maybe you can use that as a subquery/building block

## ❖ Top-down strategy

- Suppose I could ask for a "magic table" given to me by someone smart, and use that as a subquery
- What would that table be and how would I use it?

# *How to write complex SQL queries?*

- ❖ Can you compute the opposite (negation/complement) of what you want
- ❖ Sailors who have reserved all boats?
  - How about the sailors who have NOT reserved all boats?

# *How to write complex SQL queries?*


- ❖ Consider "instantiating" something in your query to make it simpler
- ❖ E.g. "find sailors who have reserved all boats"
  - If that's too hard, can you do "find out whether sailor with sid 1 has reserved all boats"?
  - I.e. write a query whose answer is empty if no, nonempty if yes (or vice versa)
- ❖ E.g. a GROUP BY query
  - If you don't know how to do something for every pair of sailorids, can you do it for a concrete pair, e.g. (1, 2)?

# Query Evaluation

- ❖ Your SQL query is parsed and translated into an intermediate representation
  - Based on relational algebra
- ❖ This allows the query to be optimized
  - May be several possible ways to compute the answer
  - And the best strategy may depend on features of the actual data
  - The DBMS is clever; can evaluate options and choose the best one

# *Relational algebra*

- ❖ A simple, powerful abstraction for representing SQL queries
- ❖ Based on the idea of operators




# *Preliminaries*

- ❖ A query is applied to relation instances, and the result of a query is also a relation instance.
  - So, operators are composable
- ❖ Start with set relational algebra – inputs and outputs are *sets* of tuples, no duplicates




# Operators

- ❖ A relational algebra operator takes as input one or more relations and outputs another relation



Unary  
operator



Binary  
operator

# *The unary RA operators*

- Select and Project

# *Selection operator*

- ❖ Input: a relation
- ❖ Output: a relation containing a subset of the tuples from the input relation
  - That satisfy a certain condition

bid	name	color
101	Misty	red
102	Pearl	blue
103	Speedy	blue

**selection**




bid = 101



bid	name	color
101	Misty	red

# *Projection operator*

- Input: a relation
- Output: a relation containing a subset of the columns from the input relation



- Relations are sets, so duplicates removed!!

*We can already use these*

- ❖ Suppose you didn't know SQL, but knew selection and projection
  - Query: find the color of boat number 101

# *The color of boat 101*

bid	name	color
101	Misty	red
102	Pearl	blue
103	Speedy	blue

**selection**

bid = 101

bid	name	color
101	Misty	red

**projection**


color

color

red

# *The color of boat 101*

projection



```
SELECT B.color  
FROM Boats B  
WHERE B.bid= '101';
```

selection

## *In more formal notation*

- ❖ Special symbols for operators
- ❖ Selection:  $\sigma$
- ❖ Projection:  $\pi$
- ❖ Our example query:

$$\pi_{color}(\sigma_{bid=101}(Boats))$$



# *Cross Product*

- ❖ Input: two relations
- ❖ Output: a relation containing all pairs of tuples from both relations

# Cross Product

bid	name	color
101	Misty	red
102	Pearl	blue



???

sid	bid	day
22	101	9/10/13
58	102	9/11/13



# *From SQL to RA*

- ❖ Projection, selection, cross product allow us to express most SELECT-FROM-WHERE queries
- ❖ Let's try it on our running example:

```
SELECT S.sname  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid AND R.bid='101';
```

## *In mathematical notation*

```
SELECT    S.sname  
FROM      Sailors S, Reserves R  
WHERE     S.sid=R.sid AND R.bid='101';
```

$$\pi_{S.sname}(\sigma_{R.bid=101 \wedge R.sid=S.sid}(R \times S))$$