

Relational Algebra Continued

Where we are

- ❖ Last time: Relational algebra – formal definitions and examples of operators

$$\sigma, \pi, \times, \bowtie, \cap, \cup, -$$

- ❖ Today:
 - More advanced RA examples
 - Relational algebra equivalences and how to prove them
 - Bag relational algebra

Harder RA queries

- ❖ Find sailors who have reserved all boats
 - Remember the double "NOT EXISTS" in SQL?
- ❖ A two-step strategy:
 - think of a "magical operator" that would help us
 - figure out how to implement the "magical operator" using $\sigma, \pi, \times, \bowtie, \cap, \cup, -$

Sailors and Boats

| sid | bid |
|-----|-----|
| s1 | b1 |
| s1 | b2 |
| s1 | b3 |
| s1 | b4 |
| s2 | b1 |
| s2 | b2 |
| s3 | b2 |
| s4 | b2 |
| s4 | b4 |

R

| bid |
|-----|
| b1 |
| b2 |
| b3 |
| b4 |

B

| sid |
|-----|
| s1 |
| s2 |
| s3 |
| s4 |

S

| sid |
|-----|
| s1 |

Property of answer relation

- ❖ Answer is a subset of Sailors
- ❖ $\text{Answer} \times B \subseteq R$
- ❖ No larger subset of S has the above property

Division

- ❖ Integer division: $16/3$ is the largest integer z such that $z * 3 \leq 16$
- ❖ Relational division: R/B is the largest relation T such that

$$T \times B \subseteq R$$

Division, Sailors and Boats

| sid | bid |
|-----|-----|
| s1 | b1 |
| s1 | b2 |
| s1 | b3 |
| s1 | b4 |
| s2 | b1 |
| s2 | b2 |
| s3 | b2 |
| s4 | b2 |
| s4 | b4 |

R

| bid |
|-----|
| b2 |

B1

| sid |
|-----|
| s1 |
| s2 |
| s3 |
| s4 |

R/B1

| bid |
|-----|
| b2 |
| b4 |

B2

| sid |
|-----|
| s1 |
| s4 |

R/B2

| bid |
|-----|
| b1 |
| b2 |
| b4 |

B3

| sid |
|-----|
| s1 |

R/B3

Division – formal definition

$$R/B = \{t \mid \forall b \in B \ (t \cdot b) \in R\}$$

Find the names of sailors who've reserved all boats

$$\rho(Tempsids, (\pi_{sid,bid} Reserves) / (\pi_{bid} Boats))$$

$$\pi_{sname}(Tempids \bowtie Sailors)$$

Expressing R/B Using Basic Operators

- ❖ Now, need to express division using basic operators
- ❖ *Idea*: For R/B , compute all values that are not 'disqualified'.
 - t value is *disqualified* if by attaching b value from B , we obtain a tb tuple that is not in R .
 - Let A be all the attributes that are in R but not in B

Disqualified r values: $\pi_A((\pi_A(R) \times B) - R)$

R/B : $\pi_A(R) -$ all disqualified tuples

More extensions

- ❖ In classical RA, no support for aggregation (MIN/MAX/COUNT) or GROUP BY
- ❖ But there are extensions that introduce such operators
 - Beyond the scope of this class

Relational Algebra Equivalences

- ❖ Selections:
 - $\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots \sigma_{cn}(R))$ (Cascade)
 - $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$ (Commute)
- ❖ Joins:
 - $R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T$ (Associative)
 - $(R \bowtie S) \equiv (S \bowtie R)$ (Commute)

More equivalences

- ❖ Do projections and selections commute?
- ❖ What about selection and join?
- ❖ Projection and join?
- ❖ See textbook for more equivalences (15.3)

Proving equivalences

- ❖ Because we have formally defined our RA operators, can PROVE that the equivalences are true.

Proving equivalences

- ❖ Some equivalences may seem obvious, but others do not
- ❖ Eg our two definitions of division

$$R/S = \{r \mid \forall s \in S \ (r \cdot s) \in R\}$$

$$\pi_A(R) - \pi_A((\pi_A(R) \times S) - R)$$

Proving equivalences

- ❖ Let's see how to prove a simple equivalence

$$\pi_A(R) \equiv \pi_A(\pi_B(R)), \quad A \subseteq B$$

- ❖ Two queries are equivalent if return same set of tuples on every database
 - Otherwise, the queries have a different meaning!

Some NON-equivalent queries

```
SELECT DISTINCT S.sname
FROM Sailors S, Boats B1, Boats B2, Reserves R1,
Reserves R2
WHERE S.sid=R1.sid AND R1.bid=B1.bid
AND S.sid=R2.sid AND R2.bid=B2.bid
AND (B1.color='red' AND B2.color='blue');
```

- ▶ These return a different answer on a DB with two sailors called Bob, one of whom reserves red boat and other reserves blue boat

```
SELECT S.sname
FROM Sailors S, Boats B, Reserves R WHERE
R.sid=S.sid AND R.bid=B.bid
AND B.color='red'
INTERSECT
SELECT S.sname
FROM Sailors S, Boats B, Reserves R WHERE
R.sid=S.sid AND R.bid=B.bid
AND B.color='blue';
```

Proving equivalences

- ❖ Let's see how to prove a simple equivalence

$$\pi_A(R) \equiv \pi_A(\pi_B(R)), \quad A \subseteq B$$

- ❖ Two queries are equivalent if return same set of tuples on every database
 - Otherwise, the queries have a different meaning!
- ❖ So, need to prove on any DB the two sets above are equal.
 - How?

$$\pi_A(R) \equiv \pi_A(\pi_B(R)), \quad A \subseteq B$$

- ❖ Show every element in LHS set is in RHS set
- ❖ And vice versa
- ❖ So need to show 2 things for arbitrary R:
- ❖ If $t \in \pi_A(R)$ then $t \in \pi_A(\pi_B(R))$
- ❖ If $t \in \pi_A(\pi_B(R))$ then $t \in \pi_A(R)$

$$\pi_A(R) \equiv \pi_A(\pi_B(R)), \quad A \subseteq B$$

- ❖ If $t \in \pi_A(R)$ then $t = s[A]$ for some $s \in R$
 - No magic here, just applying def of projection!
- ❖ Now, consider tuple s as it is fed through our other query $\pi_A(\pi_B(R))$
 - Yields a result tuple $(s[B])[A]$ (again from def)
 - If we can show $(s[B])[A] = t$ we are done!
 - Because it will follow that $t \in \pi_A(\pi_B(R))$

$$\pi_A(R) \equiv \pi_A(\pi_B(R)), \quad A \subseteq B$$

- ❖ Need to show $(s[B])[A] = t$
- ❖ But know $t = s[A]$
- ❖ So need to show $(s[B])[A] = s[A]$

- ❖ This is clear from the fact that $A \subseteq B$
- ❖ And from the definition of []

Reminder from last lecture

- ❖ For a tuple t , let $t[A]$ denote the restriction of tuple t to exactly the attributes in A
- ❖ Suppose $A = \{a_1, a_2, \dots, a_k\}$
- ❖ Suppose $t(a)$ denotes value of tuple t for attribute a
- ❖ Then $t[A] = t(a_1) \cdot t(a_2) \cdot \dots \cdot t(a_k)$
- ❖ The dot represents concatenation

$$\pi_A(R) \equiv \pi_A(\pi_B(R)), \quad A \subseteq B$$

- ❖ So far we have proved one direction (the "forward" direction) of this equivalence
- ❖ Other direction – pretty similar
- ❖ Start with tuple t in result of $\pi_A(\pi_B(R))$
- ❖ Show that it also appears in result of $\pi_A(R)$
 - Will need to use definition of projection
 - Will need to trace back the origin ("provenance") of t to some $s \in R$

Proving equivalences

- ❖ Remember to prove in both directions
- ❖ And make it clear which direction you are proving at all times
- ❖ Will be using definitions of the operators
- ❖ Will be tracing a tuple back through the operators to original relation
 - And then back down "on the other side"
- ❖ Some proofs will be more "exciting" than others

Bag relational algebra

- ❖ In the Relational Algebra we saw, all inputs and outputs to the operators are sets
 - No duplicates
- ❖ But in practice want to be able to reason about *bags*
 - A bag is a set that may contain repetitions/multiplicities
 - $\{1, 1, 2, 3\}$ is a bag but not a set
 - $\{1, 2, 3\}$ is a set and also a bag

Why bags?

- ❖ Basic relations you use are still (most likely) sets
 - If they have a primary key they certainly are
- ❖ But it can be handy to reason about bags as intermediate steps in evaluation
 - The π operator removes duplicates
 - But this may be an expensive operation
 - So why do it if we don't care about unique values?
 - Would be handy to have a "projection that keeps duplicates" operator

Bag relational algebra

- ❖ Can define an algebra on relations that are bags

R

| a | b |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 3 | 3 |

- ❖ Will have all the same operators, but need to explain how they work on bags

Selection

- ❖ Works the same way as before, it can just output a bag this time

R

| a | b |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 3 | 3 |

$$\sigma_{a=1}(R)$$

| a | b |
|---|---|
| 1 | 2 |
| 1 | 2 |

Projection

- ❖ Same as before, just doesn't eliminate duplicates

R

| a | b |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 3 | 3 |

$\pi_a(R)$

| a |
|---|
| 1 |
| 1 |
| 3 |

Cross product and join

- ❖ Same as before, treat any duplicates as though they were separate tuples

R

| a | b |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 3 | 3 |

S

| c | d |
|---|---|
| 7 | 8 |
| 9 | 2 |

$R \times S$

| a | b | c | d |
|---|---|---|---|
| 1 | 2 | 7 | 8 |
| 1 | 2 | 9 | 2 |
| 1 | 2 | 7 | 8 |
| 1 | 2 | 9 | 2 |
| 3 | 3 | 7 | 8 |
| 3 | 3 | 9 | 2 |

Union, intersection, difference

❖ $R \cup_B S$

- Each element appears as many times as it appears in R plus as many times as it appears in S

❖ $R \cap_B S$

- Each element appears the minimum number of times it appears either in R or S

❖ $R -_B S$

- Each element appears the number of times it appears in R, minus # times appears in S, but never <0 times.

Examples

R

| a | b |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 3 | 4 |

S

| a | b |
|---|---|
| 1 | 2 |
| 3 | 4 |
| 3 | 4 |

| a | b |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 1 | 2 |
| 3 | 4 |
| 3 | 4 |
| 3 | 4 |

$$R \cup_B S$$

| a | b |
|---|---|
| 1 | 2 |
| 3 | 4 |

$$R \cap_B S$$

| a | b |
|---|---|
| 1 | 2 |

$$R -_B S$$

Union, intersection, set difference

- ❖ If you want to play with those, try UNION ALL, INTERSECT ALL, EXCEPT ALL
- ❖ In Postgres, because MySQL doesn't support INTERSECT and EXCEPT...
- ❖ Basically: real SQL implements a combination of set and bag semantics

Equivalences

- ❖ Some classical equivalences no longer apply in bag relational algebra
 - E.g. $R \cup R = R$
- ❖ To prove an equivalence, now need to show:
 - For every tuple t , if appears with multiplicity k (i.e. k times) on RHS, appears k times on LHS
 - And vice versa

Clarification

- ❖ Bag relational algebra is useful in practice, but some courses (and your textbook) don't talk about it
- ❖ In the remainder of this course, when we ask you to write something in "relational algebra", assume we mean set RA
 - If we don't, we will mention the bag RA explicitly
- ❖ Practicum folks – you are working in the bag RA

Relational Algebra Summary

- ❖ Basic operators and how to use them
 - Selection, projection, cross product, union, intersection, difference
- ❖ Derived operators
 - E.g. join, division (defined in terms of basic ones)
- ❖ Equivalences and how to prove them
- ❖ Bag RA