

Implementing Relational Operators: Selection, Projection, Join



Readings

❖ [RG] Sec. 14.1-14.4



Last time

- ❖ Started discussion on how to implement RA operators (selection)
- ❖ Metric: #I/Os required
 - don't include the cost of writing out results
 - same for all implementations



No index


- ❖ Unsorted file (or sorted on "wrong" attribute)
 - Only option is to scan whole thing linearly
 - Cost: # pages in the file
- ❖ File is sorted on selection attribute
 - Binary search, then retrieve all qualifying tuples
 - Cost: $\log(\text{\#pages in file}) + \text{\#pages with qualifying tuples}$




Using indexes for selection


❖ Tree index cost:

- Traverse tree from root
 - ◆ Done only once
 - ◆ 2 to 3 I/Os (trees are short)
- Scan leaf level pages to find all data entries
- Retrieve actual data records




Cost Components – Tree Index





Cost Components – Tree Index

Index




Component 2: Traversing
subset of data entries in index
(may be ok to skip if clustered)

File



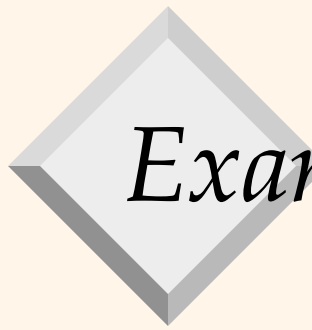
Cost Components – Tree Index





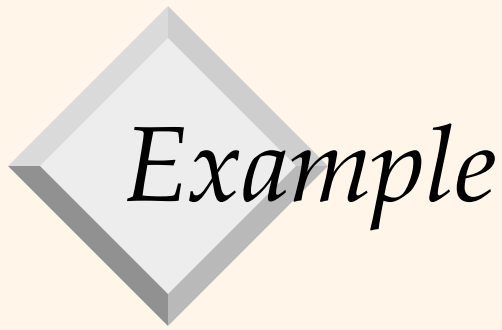
What about hash indexes?

- ❖ Depends on implementation (linear, extendible etc)
- ❖ But typically small
- ❖ Reasonable assumption: 1 or 2 I/Os to find right bucket
 - Plus cost of retrieving actual data records (depends on number of matching records)




Example

- ❖ Selection on Reserves, condition is "R.rname < 'C%'"
- ❖ Assume uniform distribution of names, so about 10% of relation should be retrieved
 - 10K tuples, 100 pages



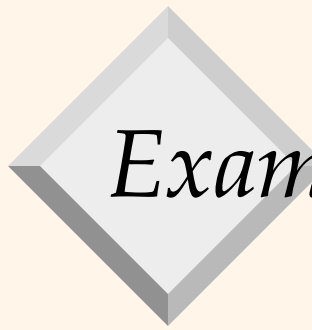
Example

- ❖ Have clustered B+ tree index on rname
 - Traverse index to find first leaf page - 1 or 2 I/Os
 - Start scan and retrieve tuples - 100 I/Os
- ❖ Have unclustered B+ tree index on rname
 - Worst case, retrieving each tuple requires a separate I/O so 10K I/Os
 - Much cheaper to just scan all 1000 original pages!
- ❖ Heuristic: scan cheaper than unclustered index if expect to retrieve more than 5% of tuples



Optimizations are possible

- ❖ Alternative 2 or 3, unclustered index
- ❖ Find qualifying data entries from index
- ❖ Sort the rids of the data entries to be retrieved
 - Remember rid = (page ID, slot #)
- ❖ Fetch rids in order
 - Ensures each data page is read from disk just once!
 - Although number of data pages retrieved still likely to be more than with clustering



Example

```
SELECT *  
FROM   Sailor S  
WHERE  S.Age = 25 AND S.Salary > 100K
```

❖ Have Hash index on Age




Evaluation Options

❖ Option 1

- Use available index (on Age) to get **superset** of relevant data entries
- Retrieve the tuples corresponding to the set of data entries
- Apply remaining predicates on retrieved tuples
- Return those tuples that satisfy all predicates

❖ Option 2

- Sequential scan! (always available)
- May be better depending on selectivity



And another example...

```
SELECT *  
FROM   Sailor S  
WHERE  S.Age = 25 AND S.Salary > 100K
```

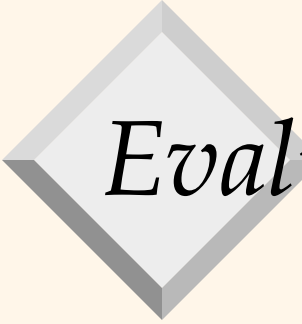
- ❖ Have Hash index on Age
- ❖ Have B+ tree index on Salary



Evaluation Options

❖ Option 1

- Choose **most selective** access path (index)
 - ◆ Could be index on Age or Salary, depending on selectivity of the corresponding predicates
- Use this index to get **superset** of relevant data entries
- Retrieve the tuples corresponding to the set of data entries
- Apply remaining predicates on retrieved tuples
- Return those tuples that satisfy all predicates



Evaluation Alternatives

❖ Option 2

- Get rids of data records using each index
 - ◆ Use index on Age and index on Salary
- **Intersect** the rids
 - ◆ We'll discuss intersection soon
- Retrieve the tuples corresponding to the rids

❖ Option 3

- Sequential scan!



More complex selection conditions

- ❖ When can we use an index?
- ❖ When it "matches" at least some of our selection condition.
- ❖ E.g. suppose we have a tree index for R on $\langle \text{sid}, \text{bid}, \text{day} \rangle$
 - Will help for for "sid > 10"
 - Will help for "sid > 10 AND bid = 100"
 - But not helpful for "bid = 100"



Using indexes for selection

- ❖ A hash index matches (a conjunction of) terms that has a term *attribute = value* for every attribute in the search key of the index.

E.g., Hash index on $\langle a, b, c \rangle$ matches $a=5 \text{ AND } b=3 \text{ AND } c=5$; but it does not match $b=3$, or $a=5 \text{ AND } b=3$, or $a>5 \text{ AND } b=3 \text{ AND } c=5$.



Using indexes for selection

- ❖ A tree index matches (a conjunction of) terms that involve only attributes in a *prefix* of the search key.

E.g., Tree index on $\langle a, b, c \rangle$ matches the selection $a=5$ AND $b=3$, and $a=5$ AND $b>6$, but not $b=3$.



Complex Selections

$(day < 8/9/94 \text{ AND } rname = 'Paul') \text{ OR } bid = 5 \text{ OR } sid = 3$

- ❖ Selection conditions are first converted to conjunctive normal form (CNF):

$(day < 8/9/94 \text{ OR } bid = 5 \text{ OR } sid = 3) \text{ AND } (rname = 'Paul' \text{ OR } bid = 5 \text{ OR } sid = 3)$



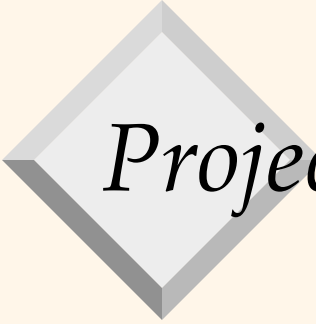
Complex Selections

- ❖ Combination of techniques seen so far
 - If have index for one (or more) of the conjuncts, can use it to filter tuples and apply remaining condition to only those
 - Can use union of retrieved tuples (or RIDs) for "OR"
 - But really, at some point sequential scan becomes best bet



Selection summary


- ❖ Depends on what is available
 - File sorted on selection attribute(s)
 - Indexes (clustered or not)
- ❖ Options:
 - Sequential scan (not always totally stupid)
 - Binary search
 - Use index or combination of indexes



Projection

```
SELECT S.Name, S. Age  
FROM   Sailors S
```

```
SELECT DISTINCT S.Name, S. Age  
FROM   Sailors S
```

Implementing Projection

- ❖ In the general case, will need to scan the whole file
 - Because we want something from each tuple in the result
- ❖ The expensive part is eliminating duplicates if desired




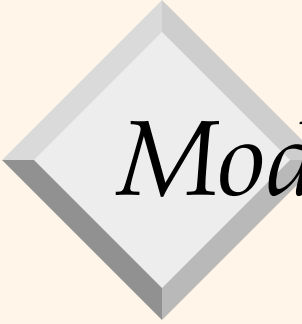
Sorting-based projection

- ❖ Basic idea:
 - make a pass through tuples and retain only desired attributes
 - sort result of the above
 - go through resulting file and output only non-duplicates
- ❖ Fortunately, we know how to do external sorting for good performance!

General External Merge Sort

- ❖ Make multiple passes to merge runs
 - Pass 1: Produce runs of length $B(B-1)$ pages
 - Pass 2: Produce runs of length $B(B-1)^2$ pages
 - ...
 - Pass P : Produce runs of length $B(B-1)^P$ pages





Modifications to External Sorting

❖ Pass 0

- Project out unwanted columns here (so don't need a separate pass before)
- Still produce runs of length B pages
- Tuples in runs are smaller than input tuples

❖ Merge passes

- Eliminate duplicates during merge



Hashing-based projection

- ❖ Idea: hash all tuples into buckets based on attributes we want
- ❖ Bucket = partition (subset) of the input
- ❖ All duplicates will go into the same bucket

- ❖ Problem: hash table may not fit in memory!




Hashing-based projection

- ❖ If whole hashtable won't fit into memory, maybe a single bucket will
- ❖ All (pairs of) duplicates will be in the same bucket
- ❖ So we can eliminate duplicates bucket by bucket


Projection Based on Hashing

- ❖ Assume relation does not fit in memory
- ❖ First pass
 - Divide relation into partitions
 - Eliminate unwanted fields as you go
 - No duplicate elimination yet!



Now process buckets one at a time

- ❖ Use a different hash function h_2

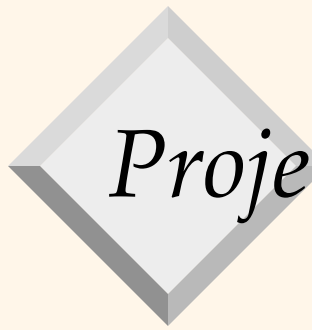


- ❖ If hash table for bucket does not fit in memory, can recursively apply hashing algorithm from previous phase.



Comments on Projection

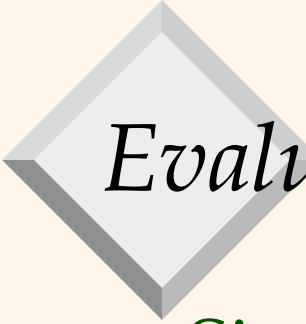
- ❖ Advantages of sort-based projection
 - Better handling of skew
 - Results in sorted order
- ❖ Hash-based projection lends itself better to parallelizing



Projection

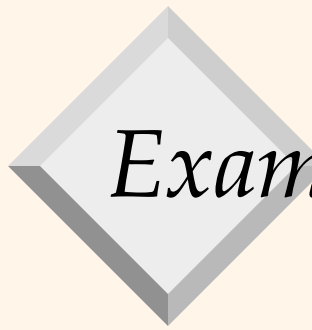
```
SELECT DISTINCT S.Name, S. Age  
FROM   Sailors S
```

❖ Have B+ tree index on (Name, Age)



Evaluation Using “Covering” Index

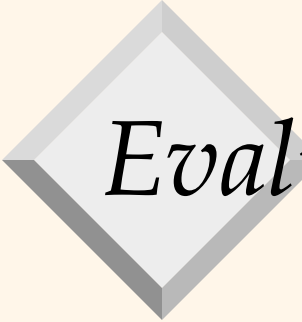
- ❖ Simply scan leaf levels of index structure
 - No need to retrieve actual data records
 - Index-only scan
- ❖ Works so long as the index search key includes all the projection attributes
 - Extra attributes in search key are okay
 - Best if projection attributes are prefix of search key
 - ◆ Can eliminate duplicates in single pass of index-only scan



Example


```
SELECT  DISTINCT S.Name, S. Age  
FROM    Sailors S
```

- ❖ Have Hash index on Name
- ❖ Have B+ tree index on Age
- ❖ Sailors relation has 100 other attributes!



Evaluation Using RID Joins

- ❖ Retrieve (SearchKey1, RID) pairs from first index
- ❖ Retrieve (SearchKey2, RID) pairs from second index
- ❖ **Join** these based on RID to get (SearchKey1, SearchKey2, RID) triples
 - We will discuss joins soon!
- ❖ Project out the third column to get the desired result



Projection Summary


- ❖ General case: scan and eliminate duplicates
 - Sorting and hashing approaches
- ❖ Indexes: can use those if available
 - Index-only scans
 - RID joins



Next up: Joins!

```
SELECT *  
FROM   Reserves R, Sailors S,  
WHERE  R.sid = S.sid
```

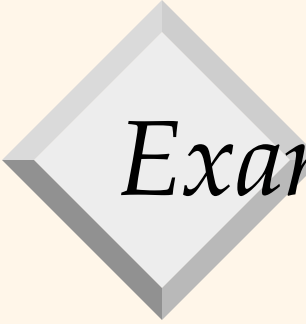
❖ No indices on Sailors or Reserves



Tuple Nested Loop Join

```
foreach tuple r in R do
  foreach tuple s in S do
    if r.sid == s.sid then add <r, s> to result
```

- ❖ R is “outer” relation
- ❖ S is “inner” relation



Example relations

- ❖ Sailors (sid, sname, rating, age)
- ❖ Reserves(sid, bid, day, rname)
- ❖ Each tuple of Sailors 50 bytes long (80 tuples per page) and have 500 pages
- ❖ Each tuple of Reserves 40 bytes long (100 tuples per page) and have 1000 pages

Analysis

❖ Assume

- M pages in R , p_R tuples per page
 - ♦ $M = 1000$, $p_R = 100$
- N pages in S , p_S tuples per page
 - ♦ $N = 500$, $p_S = 80$

❖ Total cost = $M + p_R * M * N$

- Ignore cost of writing out result
- Same for all join methods

❖ This is 50,001,000 I/O's in our example!



Page Nested Loop Join

```
foreach page p1 in R do
  foreach page p2 in S do
    foreach r in p1 do
      foreach s in p2 do
        if r.sid == s.sid then add <r, s> to result
```

- ❖ R is “outer” relation
- ❖ S is “inner” relation

Analysis

- ❖ Assume
 - M pages in R, p_R tuples per page
 - ◆ $M = 1000, p_R = 100$
 - N pages in S, p_S tuples per page Select
 - ◆ $N = 500, p_S = 80$
- ❖ Total cost = $M + M * N$
 - Which is 501,000 I/O s
- ❖ Note: Smaller relation should be “outer”
 - Better for S to be “outer” in this case!