

Operator Implementation Wrap-Up

Query Optimization

Last time:

❖ Nested loop join algorithms:

- TNLJ
- PNLJ
- BNLJ
- INLJ

❖ Sort Merge Join

❖ Hash Join

General Join Conditions


- ❖ Equalities over several attributes (e.g., *R.sid=S.sid AND R.rname=S.sname*):
 - Index NL works if we have an index on both sid and sname (together or separately)
 - For Sort-Merge and Hash Join, sort/partition on combination of the two join columns.

General Join Conditions

- ❖ Inequality conditions (e.g., *R.rname* < *S.sname*):
 - For Index NL, need (clustered!) B+ tree index.
 - Range probes on inner; # matches likely to be much higher than for equality joins.
 - Hash Join, Sort Merge Join not applicable.
 - Block NLJ quite likely to be the best join method here.


Blocking vs non blocking algorithms

- ❖ Suppose your join is not evaluated in isolation, but sits within a bigger RA tree
- ❖ Interesting to think about overall evaluation




Blocking vs non blocking algorithms

- ❖ Selection produces R tuples one at a time
- ❖ Some join algorithms can get started right away, others can't



Blocking vs non blocking algorithms

- ❖ Hash join has to wait for all of the $\sigma(R)$ tuples - **blocking**
- ❖ Nested loops joins can get started right away - **non blocking**
- ❖ Sort-merge join?
- ❖ Note: not the same usage of "blocking" as in BNLJ!!
 - Here: blocking = needs to read at least one input completely before producing output



Other implementations

- ❖ The implementations you have seen are geared to specific requirements
 - Compute **entire** result **as fast as possible**
- ❖ Sometimes you may have a setting with different desiderata
 - Which may call for totally different implementations!

Case Study: Online Aggregation

- ❖ Setting: interactive exploration of large data sets to discover general trends
- ❖ Example:

```
SELECT P.zipcode , AVG(E.salary)
FROM EmploymentData E, PersonalData P
WHERE E.ssn = P.ssn
GROUP BY P.zipcode;
```

Case Study: Online Aggregation

- ❖ This query may take a lot of time to compute
- ❖ But we often don't need precise results
- ❖ Some estimate (with a confidence interval) is enough
- ❖ So, goal is to compute **partial results fast**
 - And give a **running confidence interval** so user can stop evaluation once satisfied
 - Means we want to **sample** from DB (definitely don't want sorted order!)

Joins for online aggregation

- ❖ Definitely do not want to have to read both relations before starting to output result
 - So the blocking algorithms (hash join and sort merge) are out

Joins for online aggregation

- ❖ Nested loops might work:
 - Pick *random* tuple from R (not trivial!)
 - Join it with all of S, update confidence interval
 - Get another random tuple and repeat until enough
 - Note would want smaller relation as inner because can update confidence intervals more often

Ripple joins

- ❖ Even better
- ❖ Avoid scanning all of the inner relation each time
- ❖ Basic idea:
 - Pick random tuple from R and random tuple from S
 - Join them
 - Pick two more random tuples, one from each
 - Join all 4 tuples together
 - Continue until confidence interval acceptable

Ripple Join: Square Two-Table Join

R
S X

Ripple Join: Square Two-Table Join

R
S X X
X X

Ripple Join: Square Two-Table Join

R
S X X X
X X X
X X X

Ripple Join: Square Two-Table Join

R
S X X X X
X X X X
X X X X
X X X X

Pseudocode

```
for (max = 1 to infty)
    for (i = 1 to max - 1)
        if (match (R[i], S[max]))
            output tuple
    for (i=1 to max)
        if (match (R[max], S[i]))
            output tuple
```

Ripple Joins

- ❖ Can have non-square aspect ratios if desired due to statistical properties of data
 - i.e. if want to sample one relation more often than the other
- ❖ This method of join evaluation allows confidence intervals to shrink quite fast
 - Theory and experimental results in research paper if you're interested ([link in CMS](#))

Ripple Joins

- ❖ Extremely inefficient if we wanted to compute the whole join
- ❖ Even worse than tuple nested loops join
 - Because would need to keep track of the already-seen tuples from both relations (eventually won't fit in memory)
- ❖ Can use indexes, blocking or hashing to help
- ❖ But the main reason this works is that we almost always stop computation very early

Set Operations

- ❖ Intersection and cross-product special cases of join.
 - Intersection: equality on all fields is join condition
 - Cross product: no equality condition

Set Operations - Union

- ❖ Main challenge: eliminating duplicates
- ❖ Sorting based approach
 - Sort both relations (on combination of all attributes).
 - Scan sorted relations and merge them.

Set Operations - Union

- ❖ Hash based approach to union:
 - Partition R and S using hash function h .
 - For each S partition, build in-memory hash table (using h_2), scan corresponding R partition and add tuples to hash table while discarding duplicates
 - When done with partition, write out hashtable as (part of) result

Set Operations

- ❖ Set difference ($R - S$) similar to union
- ❖ Sorting-based approach
 - During merge pass, write out tuples in R after checking that do not appear in S
- ❖ Hashing-based approach
 - For each tuple of R , probe hashtable for partition of S and only write tuple (to output) if *not* found in hashtable.

Aggregate Operations (AVG, MIN, etc.)

❖ Without grouping:

- In general, requires scanning the relation.
- Keep track of some "running information"
 - SUM?
 - MAX/MIN?
 - AVG?

Aggregate Operations (AVG, MIN, etc.)

❖ With grouping:

- Sort on group-by attributes, then scan relation and compute aggregate for each group.
 - Can improve upon this by combining sorting and aggregate computation (total cost = cost of sort in this case)
- Hashing approach: build a hash table with entries $\langle \text{grouping-value}, \text{running-info} \rangle$ and update it as you scan the entire relation

Aggregate Operations (AVG, MIN, etc.)

- ❖ May be able to use indexes
 - Index-only scan
 - Retrieve records in sorted order instead of having to sort

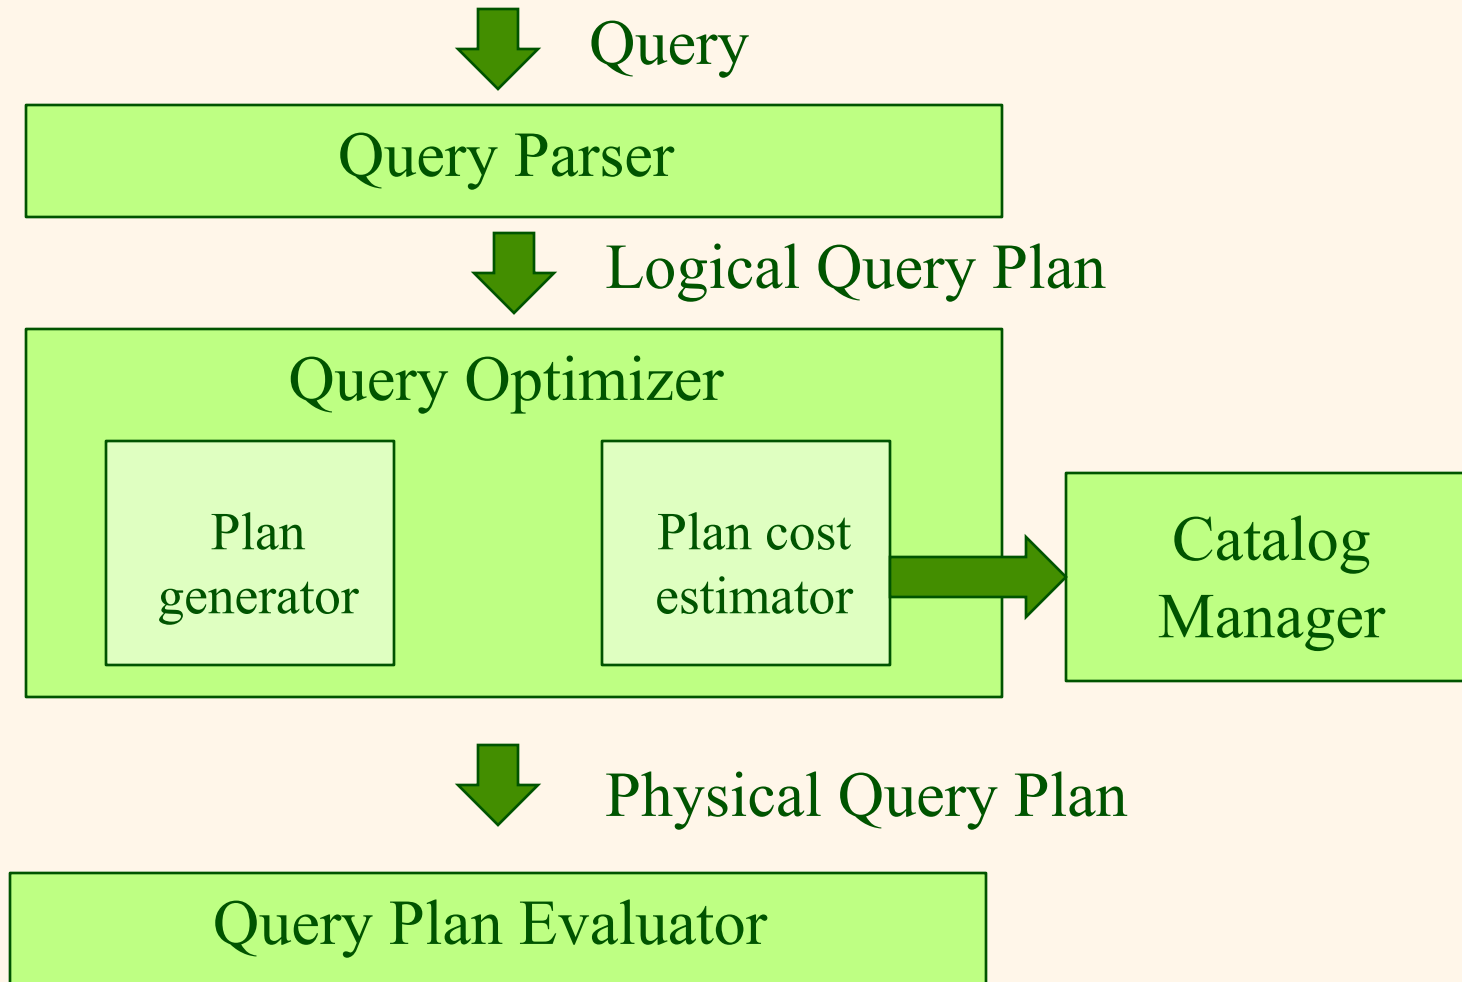
Summary so far

- ❖ Understand how to implement basic Relational Algebra operators
 - Select
 - Project
 - Join
 - Set operators
 - Aggregation/GROUP BY
- ❖ Understand that other implementations may be appropriate if setting/requirements are different

Putting it all together

- ❖ How to use these implementation algorithms to process your SQL queries efficiently?
- ❖ Query evaluation and optimization

Putting it all together



Parsing and decomposition

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.age =  
      (SELECT MAX(S2.age)  
       FROM Sailors S2);
```

- ❖ This query will generate two **blocks**
- ❖ Blocks correspond to a single SELECT-FROM-WHERE clause
- ❖ Blocks optimized one at a time

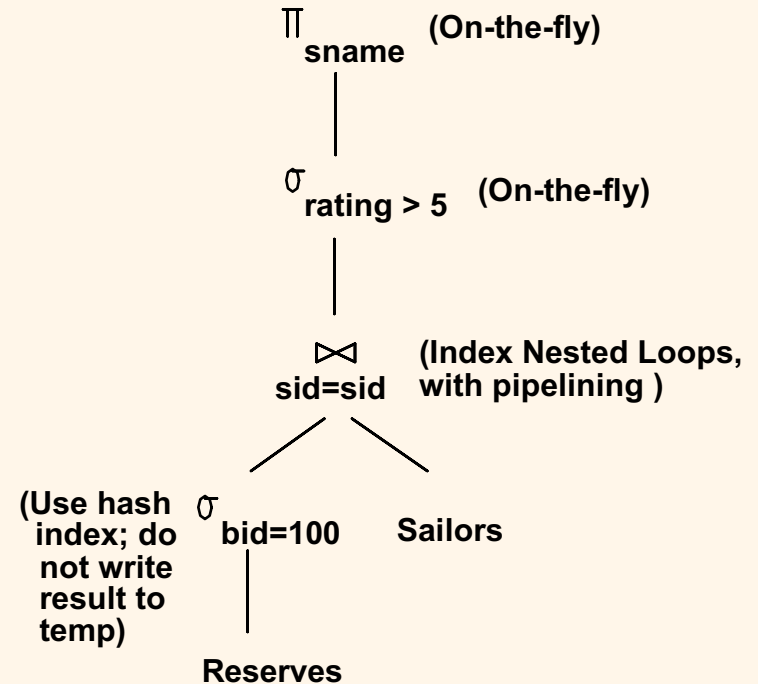

Optimizing a block

- ❖ A block is basically a Relational Algebra select-project-join ($\sigma\pi\bowtie$) expression
- ❖ With additional "operators"/annotations to handle features like
 - Aggregation
 - GROUP BY
 - ORDER BY
 - Etc
- ❖ Core of the optimizer's work: find the best **physical query plan** for the SPJ part

Query & logical and physical plans

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

Logical
query plan:



- ❖ **Physical query plan** = RA tree annotated with info on access methods and operator implementation

The work of an optimizer

- ❖ Generate some different physical plans
 - Reorder operators (using our handy RA equivalences)
 - Experiment with different implementations for each operator
- ❖ For each plan, estimate the cost
- ❖ Choose the lowest-cost plan

The work of an optimizer

- ❖ Ideally: Want to find best plan. Practically: Avoid worst plans!
- ❖ Optimization can't take too long, otherwise might defeat the purpose (if takes longer to optimize than to run a "dumb" plan)

Two main questions

- ❖ How to generate (a good subset of) the possible plans?
- ❖ How to compute the cost of each plan?
- ❖ Let's start with the second question....
 - Basically it's about putting together the per-operator calculations we have been doing already

Let's look at some examples

Sailors (sid: integer, sname: string, rating: integer, age: real)

Reserves (sid: integer, bid: integer, day: date, rname: string)

❖ Reserves:

- Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.

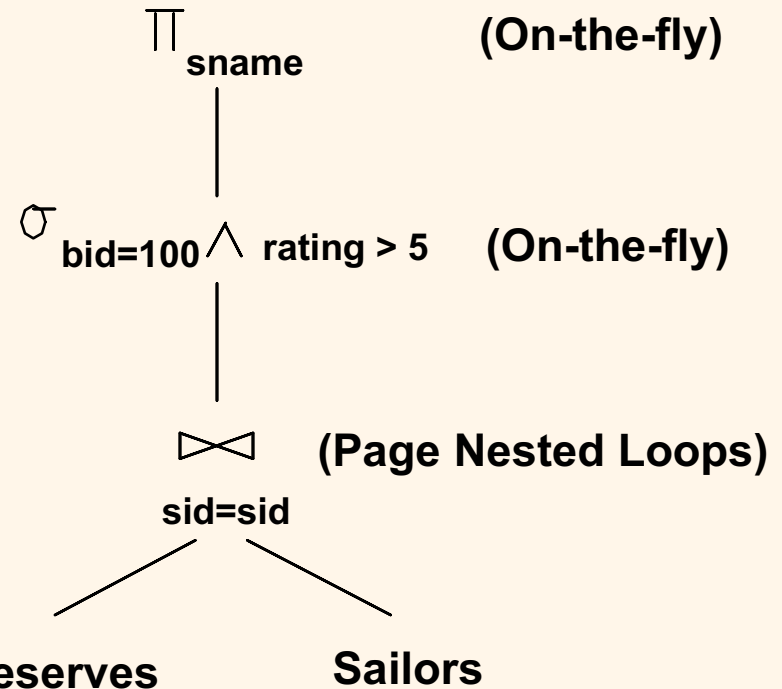

❖ Sailors:

- Each tuple is 50 bytes long, 80 tuples per page, 500 pages.

A first physical query plan

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

Logical
plan:



❖ Cost: $1000 + 500 * 1000 = 501,000$ page I/Os

❖ Convention: **left child** of join = **outer relation**

FAQ: what does "on the fly" mean?

- ❖ Just means that we can apply the operation while the tuple is already in memory
- ❖ So no extra I/O cost