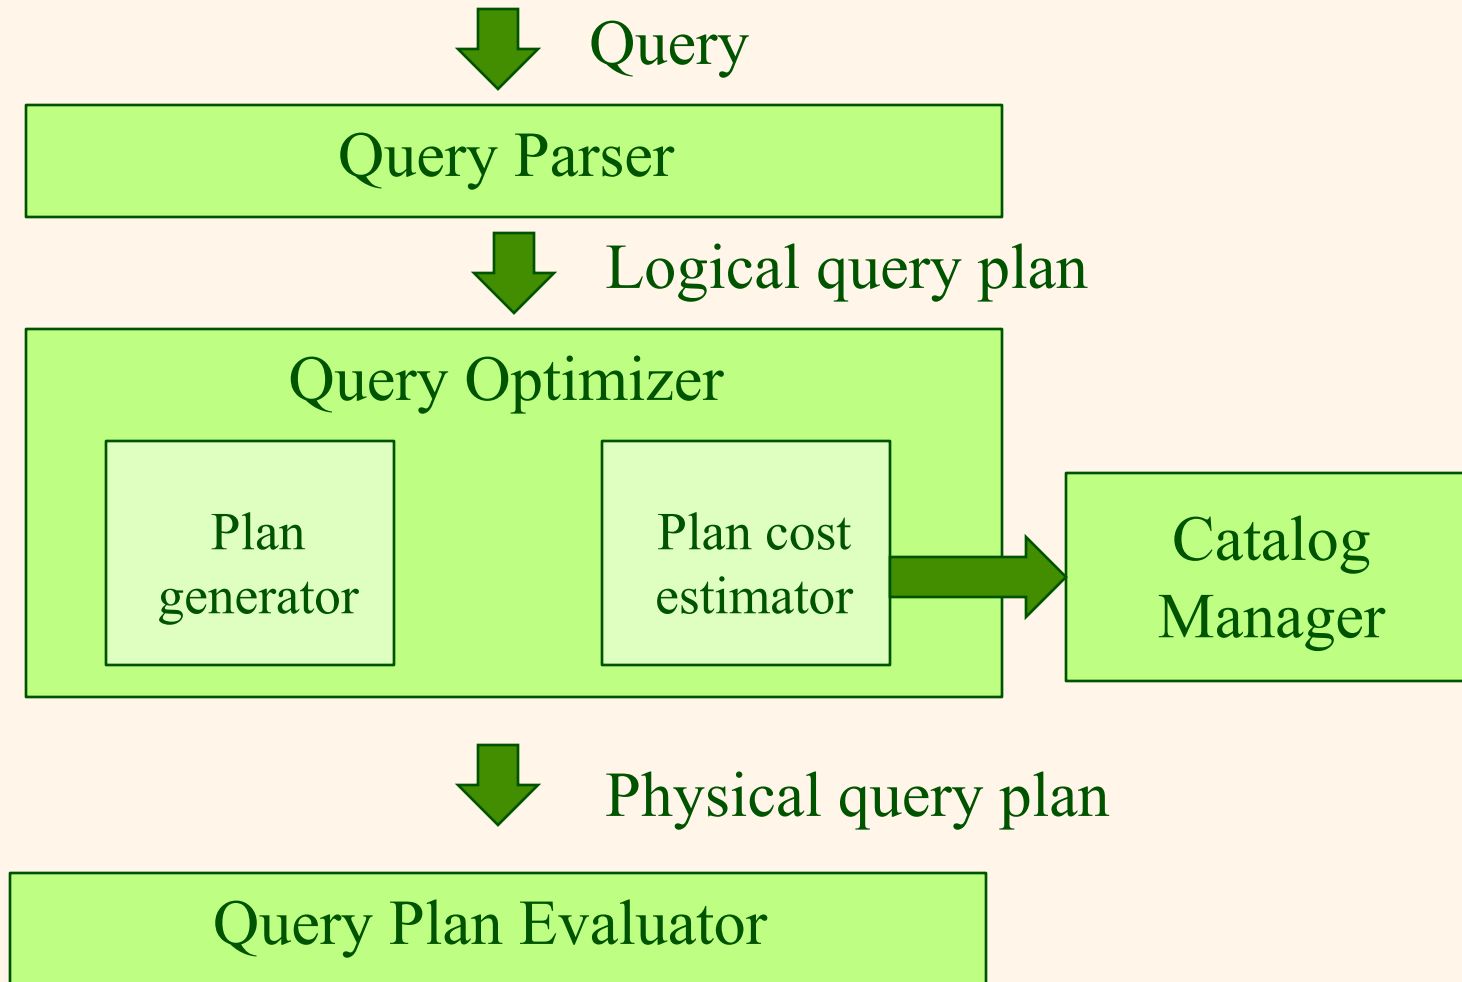


Query Optimization

Query optimization



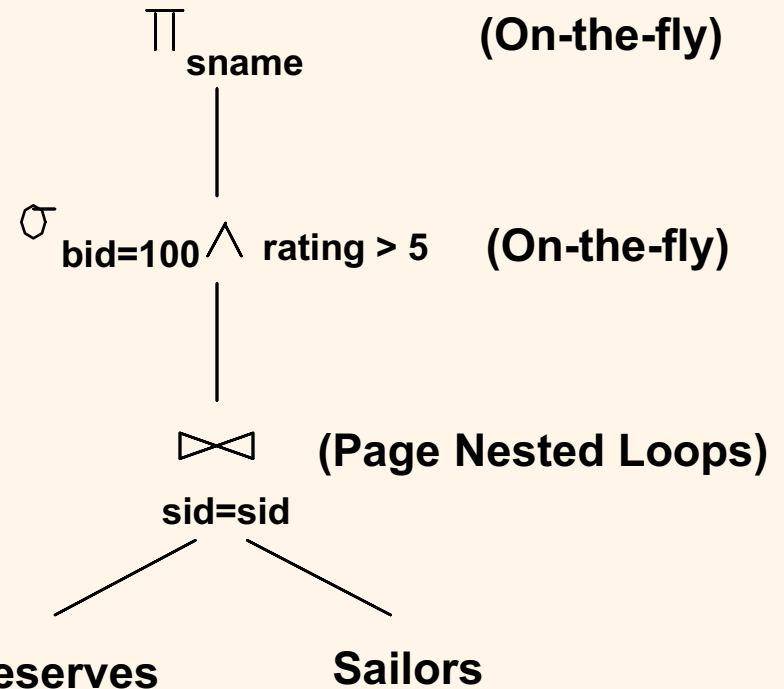

Two main questions

- ❖ How to generate (a good subset of) the possible plans?
- ❖ How to compute the cost of each plan?
- ❖ Let's start with the second question....
 - Basically it's about putting together the per-operator calculations we have been doing already

A first physical query plan

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

Logical
plan:




- ❖ Cost: $1000 + 500 * 1000 = 501,000$ page I/Os
- ❖ Convention: **left child** of join = **outer relation**

FAQ: what does "on the fly" mean?

- ❖ Just means that we can apply the operation while the tuple is already in memory
- ❖ So no extra I/O cost


Alternative Plans (No Indexes)



❖ Sort-merge join with 5 buffers:


- Scan Reserves (1000) + write temp T1 (10 pages, if we have 100 boats, uniform distribution).
- Scan Sailors (500) + write temp T2 (250 pages, if we have 10 ratings).
- Sort T1 ($2 \times 2 \times 10$), sort T2 ($2 \times 4 \times 250$), merge (10+250)
- Total: 4060 page I/Os.

Alternative Plans (No Indexes)



- ❖ Suppose we do BNLJ instead (3-pg blocks for T1) join cost = $10+4*250$, total cost = 2770.
- ❖ If we push projections, T1 has only *sid*, T2 only *sid* and *sname*:
 - T1 fits in 3 pages, cost of BNL drops substantially, total < 2000.


Alternative Plans (With Indexes)



- ❖ Now suppose have hash index on bid of Reserves and a hash index on sid of Sailors
- ❖ INLJ with pipelining
 - Outer relation in the join is never materialized
- ❖ Join column *sid* is a key for Sailors.
 - At most one matching tuple

Alternative Plans 2

With Indexes



- ❖ With clustered index on *bid* of Reserves, we get $100,000/100 = 1000$ tuples on $1000/100 = 10$ pages.
- ❖ Decision not to push *rating*>5 before the join is based on availability of *sid* index on Sailors.
- ❖ **Cost:** Selection of Reserves tuples (10 I/Os); for each, must get matching Sailors tuple (1000×1.2); total **1210 I/Os**.

Calculating cost of a plan

❖ Cost components:

- Reading input tables
- Cost of each node/operator in the plan
 - Including writing intermediate tables if appropriate
- Sorting result at the end if needed

Calculating cost of a plan

- ❖ Need to make assumptions about data to estimate size of intermediate tables and results
 - What fraction of the tuples will pass this selection condition?
 - How many tuples from relation R1 will join with each tuple from relation R2?

Estimating sizes

- ❖ How many tuples pass the selection in "SELECT * FROM R WHERE R.A < 10"?
- ❖ We want the reduction factor for the selection

result tuples = # tuples in R * reduction factor

- ❖ Several approaches depending on precision desired

Reduction factors

- ❖ Can pick an arbitrary reduction factor e.g. 0.1 (0.3 for inequality constraints like $R.A < 42$).
 - This may be enough! Remember we just want to compare plans to each other, not compute the true costs
- ❖ Can use a more precise reduction factor based on statistics/histograms

Data statistics

- ❖ DB catalogs typically contain at least:
 - # tuples and # pages for each relation.
 - # distinct key values and low/high key values for each index.
- ❖ Catalogs updated periodically.
 - Updating whenever data changes is too expensive; lots of approximation anyway, so slight inconsistency ok.
- ❖ More detailed information (e.g., histograms of the values in some field) sometimes stored.

Reduction factors

- ❖ In our homework/exam questions we usually give you some relevant info about data distribution
 - If we do, use that info.
 - If we don't, assume some sensible reduction factor (and tell us what assumption you are making)

Reduction factors for joins

- ❖ `SELECT * FROM R, S WHERE R.A = S.B`
- ❖ How big is the join result?
 - Could vary from 0 to the product of $|R|$ and $|S|$
- ❖ Various heuristics with various levels of precision
 - Your textbook discusses some of them (Section 15.2.1)
 - You'll explore some others in the Practicum P5
 - In an exam/homework question, watch for relevant info such as primary/foreign keys.

Enumeration of Physical Plans


- ❖ There are two main cases:
 - Single-relation plans
 - Multiple-relation plans

Single relation queries

- ❖ No joins (by definition)
- ❖ Need to access the relation somehow
 - file scan or index
 - consider each possible access path and pick cheapest one
- ❖ Tuples are then pipelined to remaining selections, projections, aggregates.


Queries Over Multiple Relations

- ❖ Core of the problem: how to evaluate the joins?
 - Can't consider all possible orderings
 - (Remember, if it takes you longer to optimize than to run the unoptimized query, that defeats the purpose....)



Queries Over Multiple Relations

- ❖ Decision: *only left-deep join trees are considered.*
 - Left-deep means the **right child of each node is a base table** (a real DB table, not a join)
 - (pushed selections/projections on base tables ok)



Queries Over Multiple Relations

- ❖ Considering only left-deep trees cuts down on search space
- ❖ Left-deep trees allow us to generate all fully pipelined plans (which are particularly desirable)
 - Outer input to join is never materialized
 - Not all left-deep plans are fully pipelined (e.g., plans that use SMJ).

Enumeration of Plans

- ❖ Left-deep plans differ only in the order of relations, the access method for each relation, and the join method for each join.

Finding the best plan

- ❖ Approach 1: exhaustive recursive search

Finding the best plan

- ❖ Consider the best plan for joining k relations
- ❖ There are k options for the last relation in the join
 - For each choice of the last relation, we want to join it with the optimal plan for the remaining $k-1$ relations
- ❖ How to compute optimal plan for $k-1$ relations?
 - Pick one as the last relation and recurse...

Finding the best plan

- ❖ Better approach: avoid recomputation through dynamic programming
- ❖ Save intermediate results that will be reused later
- ❖ Compute bottom-up from subsets of size 1, 2, 3, etc.

Enumeration of Plans

- Pass 1: Find best 1-relation plan for each relation
 - includes any selects/projects just on this relation.
- Pass 2: Find best way to join result of each 1-relation plan (as outer) to another relation. (*All 2-relation plans.*)
- Pass k: Find best way to join result of a (k-1)-relation plan (as outer) to the kth relation. (*All k-relation plans.*)

Some practical remarks

- ❖ Cost function for intermediate plans may be number of I/Os or something else
 - Simpler e.g. intermediate relation sizes
 - More complex e.g. include whether plan produces tuples in a sorted order (could be useful!!)
 - Your textbook presents an algorithm where we retain both the cheapest plan and any sorted-order plans

Enumeration of Plans (Contd.)

- ❖ ORDER BY, GROUP BY, aggregates etc. handled as a final step
 - Use a plan that gives result in sorted order
 - Or use an additional sorting operator

Example

Sailors:

Hash, B+ on *sid*


Reserves:

Clustered B+ tree on *bid*

B+ on *sid*

Boats

B+, Hash on *color*



```
SELECT S.sid, COUNT(*)
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
GROUP BY S.sid;
```

Pass 1

- ❖ Best plan for accessing each relation regarded as the first relation in an execution plan
 - Sailors: File Scan
 - going through B+ index gives tuples in sorted order but index is unclustered – not a good plan (file scan + sort likely better!!)
 - Reserves: File Scan
 - clustered B+ index also works but why pay the overhead of root-to-leaf navigation?
 - Boats: Hash index on color

Pass 2

- ❖ For each of the plans in pass 1, generate plans joining another relation as the inner, using all join methods
 - File Scan *Sailors* (outer) with *Reserves* (inner)
 - File Scan *Reserves* (outer) with *Sailors* (inner)
 - *Boats* hash on color (outer) with *Reserves* (inner)
 - File Scan *Reserves* (outer) with *Boats* (inner)
 - ... etc
- ❖ Retain cheapest plan for each pair of relations
 - Also sorted-order plans even if they are not cheapest

A pruning heuristic

- ❖ Do not combine a partial plan with a relation unless there is a nontrivial join condition between them
 - i.e., avoid Cartesian products if possible.
 - in our case, don't bother with the Sailors/Boats pair of relations
 - note may not always be possible (e.g. if query requires cross product)

Pass 3

- ❖ For each of the plans retained from Pass 2, taken as the outer, generate plans for join with the last table
 - E.g.:
 - Outer: Boats hash on color with Reserves (bid) (sort-merge)
 - Inner: Sailors (file scan)
 - Join algorithm: sort-merge

Add cost of aggregate/GROUP BY

- ❖ Cost to sort the result by sid, if not returned sorted

Optimization summary

- ❖ Parse query into RA tree
- ❖ Optimize one block at a time
- ❖ Generate a subset of the possible evaluation plans
 - Dynamic programming approach
 - Investigates only left-deep join plans
- ❖ Cost calculations based on statistics maintained in the catalog