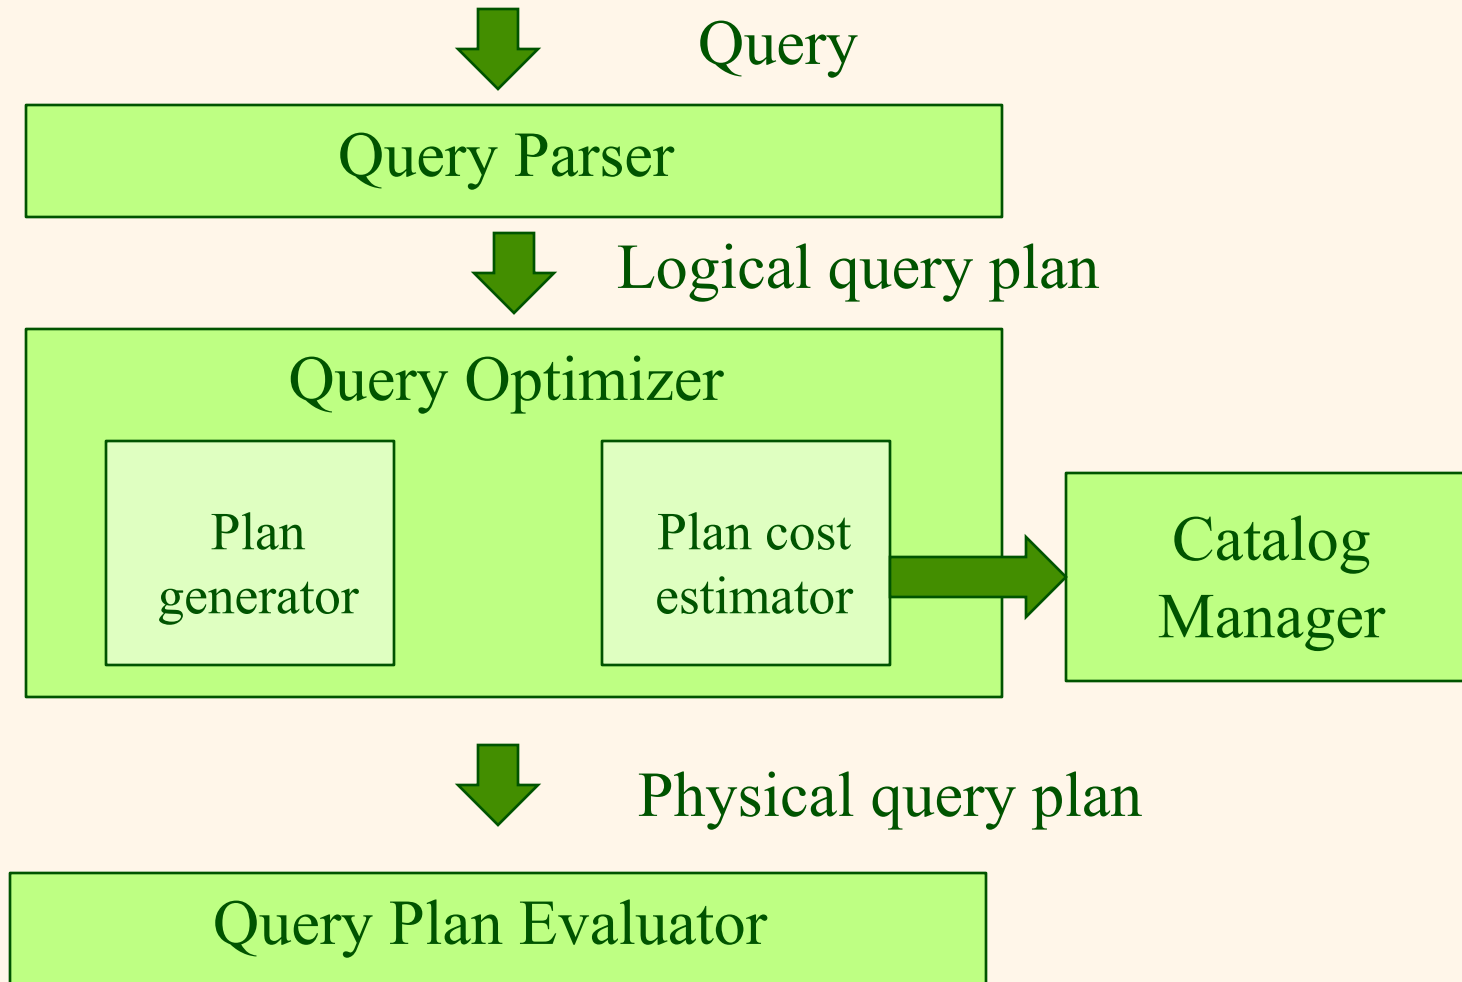



Query Optimization Wrap-up *(end of Part 2 of course)*

Query optimization





Parsing and decomposition

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.age =  
      (SELECT MAX(S2.age)  
       FROM Sailors S2);
```

- ❖ This query will generate two **blocks**
- ❖ Blocks correspond to a single SELECT-FROM-WHERE clause
- ❖ Blocks optimized one at a time





Optimizing a block


- ❖ A block is basically a Relational Algebra select-project-join ($\sigma\pi\bowtie$) expression
- ❖ With additional "operators"/annotations to handle features like
 - Aggregation
 - GROUP BY
 - ORDER BY
 - Etc
- ❖ Core of the optimizer's work: find the best **physical query plan** for the SPJ part

Query & logical and physical plans

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```




- ❖ **Physical query plan** = RA tree annotated with info on access methods and operator implementation



The work of an optimizer

- ❖ Generate some different physical plans
 - Reorder operators (using our handy RA equivalences)
 - Experiment with different implementations for each operator
- ❖ For each plan, estimate the cost
- ❖ Choose the lowest-cost plan



The work of an optimizer

- ❖ Ideally: Want to find best plan. Practically: Avoid worst plans!
- ❖ Optimization can't take too long, otherwise might defeat the purpose (if takes longer to optimize than to run a "dumb" plan)



Two main questions

- ❖ How to generate (a good subset of) the possible plans?
- ❖ How to compute the cost of each plan?
- ❖ Let's start with the second question....
 - Basically it's about putting together the per-operator calculations we have been doing already



Let's look at some examples

Sailors (sid: integer, sname: string, rating: integer, age: real)

Reserves (sid: integer, bid: integer, day: date, rname: string)

❖ Reserves:

- Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.

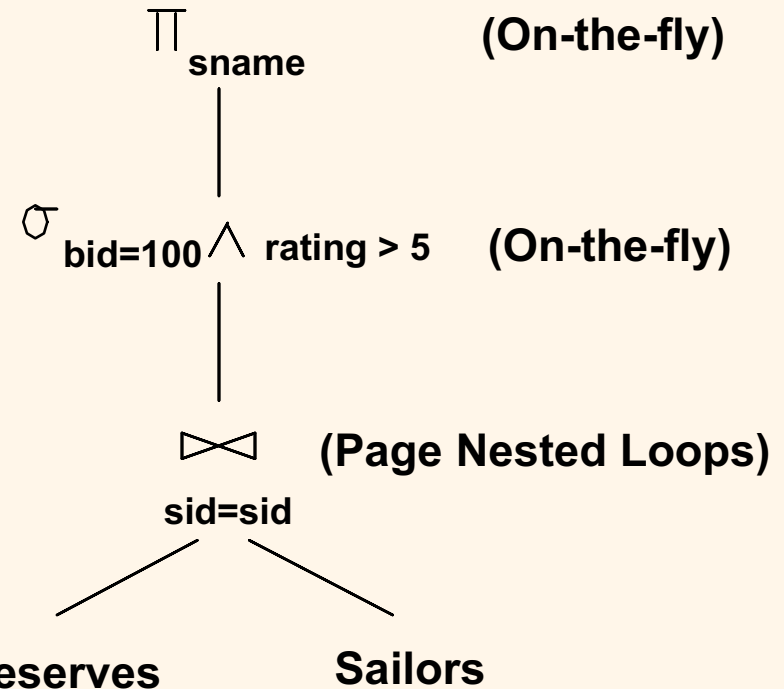

❖ Sailors:

- Each tuple is 50 bytes long, 80 tuples per page, 500 pages.

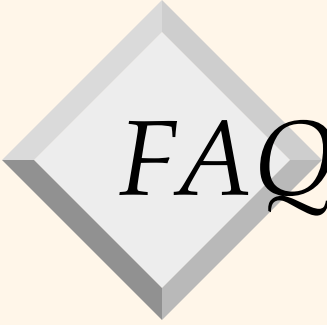
A first physical query plan

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
R.bid=100 AND S.rating>5
```

Logical
plan:




- ❖ Cost: $1000 + 500 * 1000 = 501,000$ page I/Os
- ❖ Convention: **left child** of join = **outer relation**



FAQ: what does "on the fly" mean?

- ❖ Just means that we can apply the operation while the tuple is already in memory
- ❖ So no extra I/O cost


Alternative Plans (No Indexes)



❖ Sort-merge join with 5 buffers:

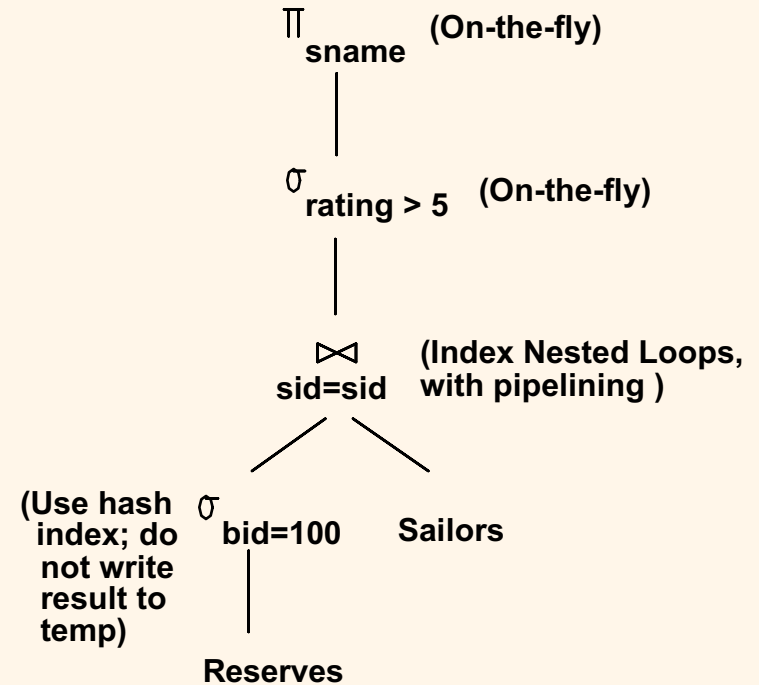
- Scan Reserves (1000) + write temp T1 (10 pages, if we have 100 boats, uniform distribution).
- Scan Sailors (500) + write temp T2 (250 pages, if we have 10 ratings).
- Sort T1 ($2 \times 2 \times 10$), sort T2 ($2 \times 4 \times 250$), merge (10+250)
- Total: 4060 page I/Os.

Alternative Plans (No Indexes)




- ❖ Suppose we do BNLJ instead (3-pg blocks for T1) join cost = $10+4*250$, total cost = 2770.
- ❖ If we push projections, T1 has only *sid*, T2 only *sid* and *sname*:
 - T1 fits in 3 pages, cost of BNL drops substantially, total < 2000.

Alternative Plans (With Indexes)



- ❖ Now suppose have hash index on `bid` of `Reserves` and a hash index on `sid` of `Sailors`
- ❖ INLJ with pipelining
 - Outer relation in the join is never materialized
- ❖ Join column `sid` is a key for `Sailors`.
 - At most one matching tuple

Alternative Plans 2 With Indexes



- ❖ With clustered index on *bid* of Reserves, we get $100,000/100 = 1000$ tuples on $1000/100 = 10$ pages.
- ❖ Decision not to push *rating*>5 before the join is based on availability of *sid* index on Sailors.
- ❖ **Cost:** Selection of Reserves tuples (10 I/Os); for each, must get matching Sailors tuple (1000×1.2); total **1210 I/Os**.



Calculating cost of a plan

- ❖ Cost components:
 - Reading input tables
 - Cost of each node/operator in the plan
 - ◆ Including writing intermediate tables if needed
 - Sorting result at the end if needed



Calculating cost of a plan

- ❖ Need to make assumptions about data to estimate size of intermediate tables and results
 - What fraction of the tuples will pass this selection condition?
 - How many tuples from relation R1 will join with each tuple from relation R2?




Estimating sizes

- ❖ How many tuples pass the selection in "SELECT * FROM R WHERE R.A < 10"?
- ❖ We want the reduction factor for the selection

result tuples = # tuples in R * reduction factor

- ❖ Several approaches depending on precision desired




Reduction factors

- ❖ Can pick an arbitrary reduction factor e.g. 0.1 (0.3 for inequality constraints like $R.A < 42$).
 - This may be enough! Remember we just want to compare plans to each other, not compute the true costs
- ❖ Can use a more precise reduction factor based on statistics/histograms




Data statistics

- ❖ DB catalogs typically contain at least:
 - # tuples and # pages for each relation.
 - # distinct key values and low/high key values for each index.
- ❖ Catalogs updated periodically.
 - Updating whenever data changes is too expensive; lots of approximation anyway, so slight inconsistency ok.
- ❖ More detailed information (e.g., histograms of the values in some field) sometimes stored.




Reduction factors

- ❖ In our homework/exam questions we usually give you some relevant info about data distribution
 - If we do, use that info.
 - If we don't, assume some sensible reduction factor (and tell us what assumption you are making)



Reduction factors for joins

- ❖ `SELECT * FROM R, S WHERE R.A = S.B`
- ❖ How big is the join result?
 - Could vary from 0 to the product of $|R|$ and $|S|$
- ❖ Various heuristics with various levels of precision
 - Your textbook discusses some of them (Section 15.2.1)
 - You'll explore some others in the Practicum P5
 - In an exam/homework question, watch for relevant info such as primary/foreign keys.



Enumeration of Physical Plans

- ❖ There are two main cases:
 - Single-relation plans
 - Multiple-relation plans






Single relation queries

- ❖ No joins (by definition)
- ❖ Need to access the relation somehow
 - file scan or index
 - consider each possible access path and pick cheapest one
- ❖ Tuples are then pipelined to remaining selections, projections, aggregates.


Queries Over Multiple Relations

- ❖ Core of the problem: how to evaluate the joins?
 - Can't consider all possible orderings
 - (Remember, if it takes you longer to optimize than to run the unoptimized query, that defeats the purpose ...)



Queries Over Multiple Relations

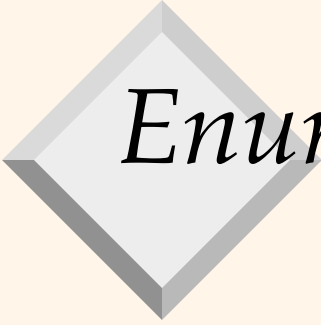
- ❖ Decision: *only left-deep join trees are considered.*
 - Left-deep means the **right child of each node is a base table** (a real DB table, not a join)
 - ◆ (pushed selections/projections on base tables ok)





Queries Over Multiple Relations

- ❖ Considering only left-deep trees cuts down on search space
- ❖ Left-deep trees allow us to generate all fully pipelined plans (which are particularly desirable)
 - ◆ Outer input to join is never materialized
 - ◆ Not all left-deep plans are fully pipelined (e.g., plans that use SMJ).



Enumeration of Plans

- ❖ Left-deep plans differ only in the order of relations, the access method for each relation, and the join method for each join.



Finding the best plan

- ❖ Approach 1: exhaustive recursive search



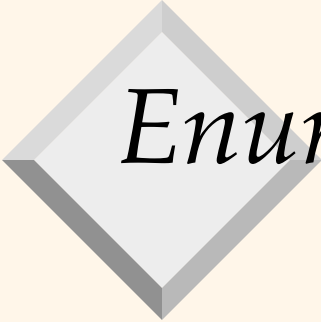
Finding the best plan

- ❖ Consider the best plan for joining k relations
- ❖ There are k options for the last relation in the join
 - For each choice of the last relation, we want to join it with the optimal plan for the remaining $k-1$ relations
- ❖ How to compute optimal plan for $k-1$ relations?
 - Pick one as the last relation and recurse...



Finding the best plan

- ❖ Better approach: avoid recomputation through dynamic programming
- ❖ Save intermediate results that will be reused later
- ❖ Compute bottom-up from subsets of size 1, 2, 3, etc.



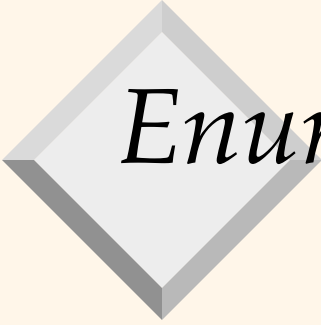
Enumeration of Plans

- Pass 1: Find best 1-relation plan for each relation
 - ◆ includes any selects/projects just on this relation.
- Pass 2: Find best way to join result of each 1-relation plan (as outer) to another relation. (*All 2-relation plans.*)
- Pass k: Find best way to join result of a (k-1)-relation plan (as outer) to the kth relation. (*All k-relation plans.*)



Some practical remarks

- ❖ Cost function for intermediate plans may be number of I/Os or something else
 - Simpler e.g. intermediate relation sizes
 - More complex e.g. include whether plan produces tuples in a sorted order (could be useful!!)
 - Your textbook presents an algorithm where we retain both the cheapest plan and any sorted-order plans



Enumeration of Plans (Contd.)

- ❖ ORDER BY, GROUP BY, aggregates etc. handled as a final step
 - Use a plan that gives result in sorted order
 - Or use an additional sorting operator

Example

Sailors:

Hash, B+ on *sid*


Reserves:

Clustered B+ tree on *bid*

B+ on *sid*

Boats

B+, Hash on *color*



```
SELECT S.sid, COUNT(*)
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
GROUP BY S.sid;
```

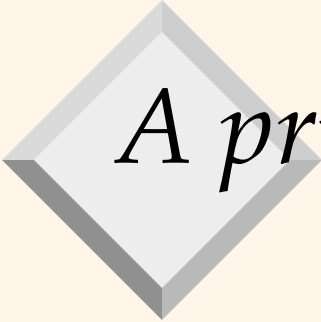
Pass 1

- ❖ Best plan for accessing each relation regarded as the first relation in an execution plan
 - Sailors: File Scan
 - ◆ going through B+ index gives tuples in sorted order but index is unclustered – not a good plan (file scan + sort likely better!!)
 - Reserves: File Scan
 - ◆ clustered B+ index also works but why pay the overhead of root-to-leaf navigation?
 - Boats: Hash index on color



Pass 2

- ❖ For each of the plans in pass 1, generate plans joining another relation as the inner, using all join methods
 - File Scan *Sailors* (outer) with *Reserves* (inner)
 - File Scan *Reserves* (outer) with *Sailors* (inner)
 - *Boats* hash on color (outer) with *Reserves* (inner)
 - File Scan *Reserves* (outer) with *Boats* (inner)
 - ... etc
- ❖ Retain cheapest plan for each pair of relations
 - Also sorted-order plans even if they are not cheapest



A pruning heuristic

- ❖ Do not combine a partial plan with a relation unless there is a nontrivial join condition between them
 - i.e., avoid Cartesian products if possible.
 - in our case, don't bother with the Sailors/Boats pair of relations
 - note may not always be possible (e.g. if query requires cross product)



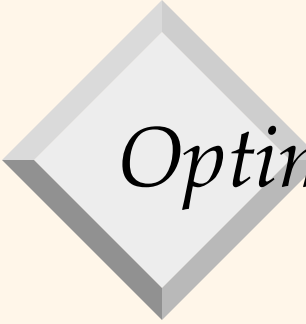
Pass 3

- ❖ For each of the plans retained from Pass 2, taken as the outer, generate plans for join with the last table
 - E.g.:
 - Outer: Boats hash on color with Reserves (bid) (sort-merge)
 - Inner: Sailors (file scan)
 - Join algorithm: sort-merge



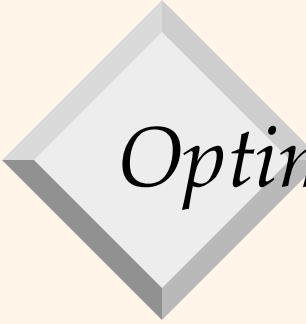
Add cost of aggregate/GROUP BY

- ❖ Cost to sort the result by sid, if not returned sorted



Optimization summary

- ❖ Parse query into RA tree
- ❖ Optimize one block at a time
- ❖ Generate a subset of the possible evaluation plans
 - Dynamic programming approach
 - Investigates only left-deep join plans
- ❖ Cost calculations based on statistics maintained in the catalog



Optimization summary

- ❖ Parse query into RA tree
- ❖ Optimize one block at a time
- ❖ Generate a subset of the possible evaluation plans
 - Dynamic programming approach
 - Investigates only left-deep join plans
- ❖ Cost calculations based on statistics maintained in the catalog



Nested Queries

- ❖ Nested block optimized independently
- ❖ Outer block optimized with cost of "calling" nested block computation taken into account.
- ❖ Ordering of blocks means some good strategies are not considered. *The non-nested version of the query is typically optimized better.*

```
SELECT S.sname  
FROM Sailors S  
WHERE S.sid IN  
      (SELECT R.sid  
       FROM Reserves R  
       WHERE R.bid=103);
```

Nested block to optimize:

```
SELECT R.sid  
FROM Reserves R  
WHERE R.bid=103
```

Equivalent non-nested query:

```
SELECT S.sname  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid  
      AND R.bid=103;
```



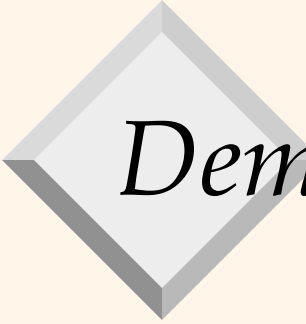
Nested Queries

- ❖ If the query is correlated may have to compute each nested block for every outer value
 - Algorithms for *decorrelation*

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS
  (SELECT *
   FROM Reserves R
   WHERE R.bid=103
   AND R.sid=S.sid);
```

Equivalent non-nested query:

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid
AND R.bid=103;
```




Demo time!

- ❖ In a real DBMS, can see the optimizer "at work"
- ❖ Let's see what indexes we have
 - SHOW INDEX FROM Sailors \G
 - \d Sailors;
- ❖ Let's add another index
 - ❖ CREATE INDEX sailors_rating ON Sailors(rating) USING BTREE;
 - ❖ CREATE INDEX sailors_rating ON Sailors USING BTREE (rating);



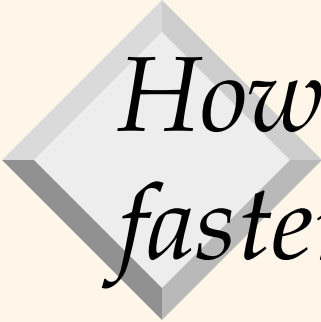
EXPLAIN statement

- ❖ `EXPLAIN SELECT S.sname FROM Sailors S WHERE S.rating > 5;`
- ❖ `EXPLAIN SELECT S.sname FROM Sailors S, Reserves R WHERE S.sid=R.sid;`
- ❖ In PostgreSQL, handy features:
 - `EXPLAIN (FORMAT JSON)`
 - `EXPLAIN ANALYZE`



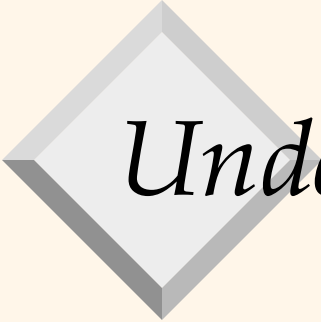
Now with more joins

- ❖ `EXPLAIN SELECT S.sname FROM Sailors S,
Boats B1, Boats B2, Reserves R1, Reserves R2
WHERE S.sid=R1.sid AND R1.bid=B1.bid
AND S.sid=R2.sid AND R2.bid=B2.bid
AND (B1.color='red' AND B2.color='blue');`



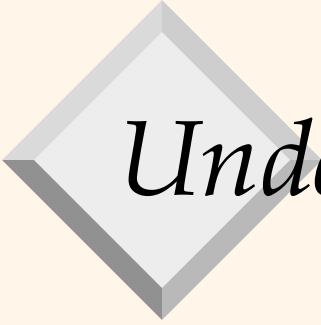
How to make your queries run faster?

- ❖ A very complex and difficult problem
- ❖ Analyze your workload and use EXPLAIN to understand what is happening
- ❖ Create indexes to help
- ❖ Understand and use the performance tuning tools your system provides
 - http://docs.oracle.com/cd/E11882_01/server.112/e41573.pdf for Oracle's Performance Tuning guide (500+ pages!!)




Understanding the Workload

- ❖ For each query in the workload:
 - Which relations does it access?
 - Which attributes are retrieved?
 - Which attributes are involved in selection/join conditions? How selective are these conditions likely to be?



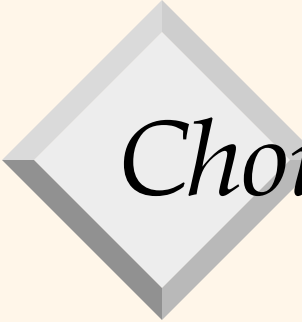
Understanding the Workload

- ❖ For each update in the workload:
 - Which attributes are involved in selection/join conditions?
 - How selective are these conditions likely to be?
 - What is the type of update (INSERT/DELETE/UPDATE), and what attributes are affected?



Choice of Indexes

- ❖ What indexes should we create?
 - Which relations should have indexes? What field(s) should be the search key? Should we build several indexes?
- ❖ For each index, what kind of an index should it be?
 - Clustered? Hash/tree?



Choice of Indexes (Contd.)

- ❖ **One approach:** Consider the most important queries in turn. Consider the best plan using the current indexes, and see if a better plan is possible with an additional index. If so, create it.
- ❖ Before creating an index, must also consider the impact on updates in the workload!
 - **Trade-off:** Indexes can make queries go faster, updates slower. Require disk space, too.




Index Selection Guidelines

- ❖ Attributes in WHERE clause are candidates for index keys.
 - Exact match condition suggests hash index.
 - Range query suggests tree index.
 - ◆ Clustering is especially useful for range queries; can also help on equality queries if there are many duplicates.



Index Selection Guidelines

- ❖ Multi-attribute search keys should be considered when a WHERE clause contains several conditions.
 - Order of attributes is important for range queries.
 - Such indexes can sometimes enable **index-only** strategies for important queries.
 - ◆ For index-only strategies, clustering is not important!



Index Selection Guidelines

- ❖ Try to choose indexes that benefit as many queries as possible.
- ❖ Since only one index can be clustered per relation, choose it based on important queries that would benefit the most from clustering.




Optimizer Hints

- ❖ Various DBMS will provide you with "hooks" to control how your query is evaluated
- ❖ Optimizer hints can be helpful but use them judiciously




Optimizer Hints - Caveats

- ❖ What works for one query may not work for another
- ❖ Or even for the same query when DB changes
- ❖ If you're using a lot of optimizer hints, something may be wrong
 - Are your statistics out of date?
 - Are you writing really bad SQL?
- ❖ BUT very useful for testing/experimentation purposes




MySQL index hints - example

- ❖ `EXPLAIN SELECT S.sname FROM Sailors S WHERE S.rating > 5 \G`
- ❖ `EXPLAIN SELECT S.sname FROM Sailors S USE INDEX (sailors_rating) WHERE S.rating > 5 \G`
- ❖ `EXPLAIN SELECT S.sname FROM Sailors S FORCE INDEX (sailors_rating) WHERE S.rating > 5 \G`



PostgreSQL

- ❖ Configuration parameters, eg. *enable_hashjoin*, *enable_indexonlyscan*
- ❖ Other params such as *random_page_cost* – estimate of cost of fetching random page from disk
- ❖ See documentation for more details



End of Part 2 of the course

- ❖ Query implementation and optimization
- ❖ Reached the end of material for **prelim**