





# *Prelim discussion*

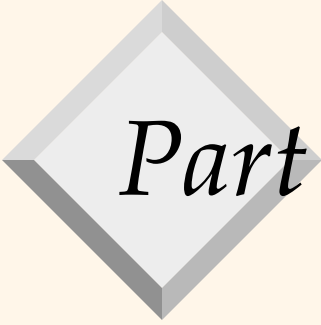
- ❖ A sample (from last Spring) is on CMS
- ❖ A 12-hour take-home
  - Noon to midnight
  - Should take *much* less than 12 hours
- ❖ Two (different) exams
  - Thurs 13 Oct
  - Sun 16 Oct
- ❖ Email me *real soon now* if you have problems with both days



# *Database Design*




*Readings: [RG] Sec. 2.1-2.5, 3.5*



## *Part 3 of the course – DB design*

- ❖ How to **apply** the relational DB abstraction to a real world problem?
- ❖ How do I create a schema that **matches reality** and has **good mathematical properties** (e.g. no redundancy?)
  - A lot of what we do applies to nonrelational abstractions too
- ❖ Shorter unit (about 1 week)
  - H3 material



## *Part 3 of the course – DB design*

- ❖ Designing a database is a complex, multi-step process
  - And often iterative
  - Once you finally "get it right", real-world scenario changes and you have to start again!




# *The Database Design Process*

- ❖ Requirements analysis
  - based on use cases, description of business processes, etc.
- ❖ Conceptual design (create data model)
  - What is the DB actually *about*?
  - May use tools such as the ER model
- ❖ Schema refinement/normalization
  - Does the schema have any "bad properties" (e.g. redundancy) that should be removed?
- ❖ Physical tuning
  - Based on factors like known query workload
  - Choose indexes etc.



# *The Database Design Process*

- ❖ Requirements analysis
  - based on use cases, description of business processes, etc.
- ❖ Conceptual design (create data model)
  - What is the DB actually *about*?
  - May use tools such as the ER model
- ❖ Schema refinement/normalization
  - Does the schema have any "bad properties" (e.g. redundancy) that should be removed?
- ❖ Physical tuning
  - Based on factors like known query workload
  - Choose indexes etc.



# *Conceptual design*

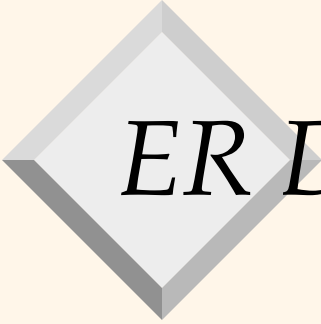
- ❖ What are the entities and relationships in the real world?
- ❖ What information about these entities and relationships should we store in the database?
- ❖ What are the integrity constraints or business rules that hold?
- ❖ Handy modeling tool: ER (entity-relationship) diagrams





# *Entity Sets*


- ❖ *Entity* = “thing” or object.
- ❖ *Entity set* = collection of similar entities.
  - Similar to a class in OO languages
  - **Instance of an entity set** = set of actual entities
- ❖ *Attribute* = property of (the entities of) an entity set.
  - Attributes are simple values, e.g. integers or character strings, not structs, sets, etc.



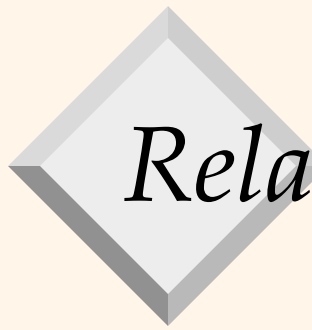
# *ER Diagrams*

- ❖ In an entity-relationship diagram:
  - Entity set = rectangle.
  - Attribute = oval, with a line to the rectangle representing its entity set.

# *Example:*



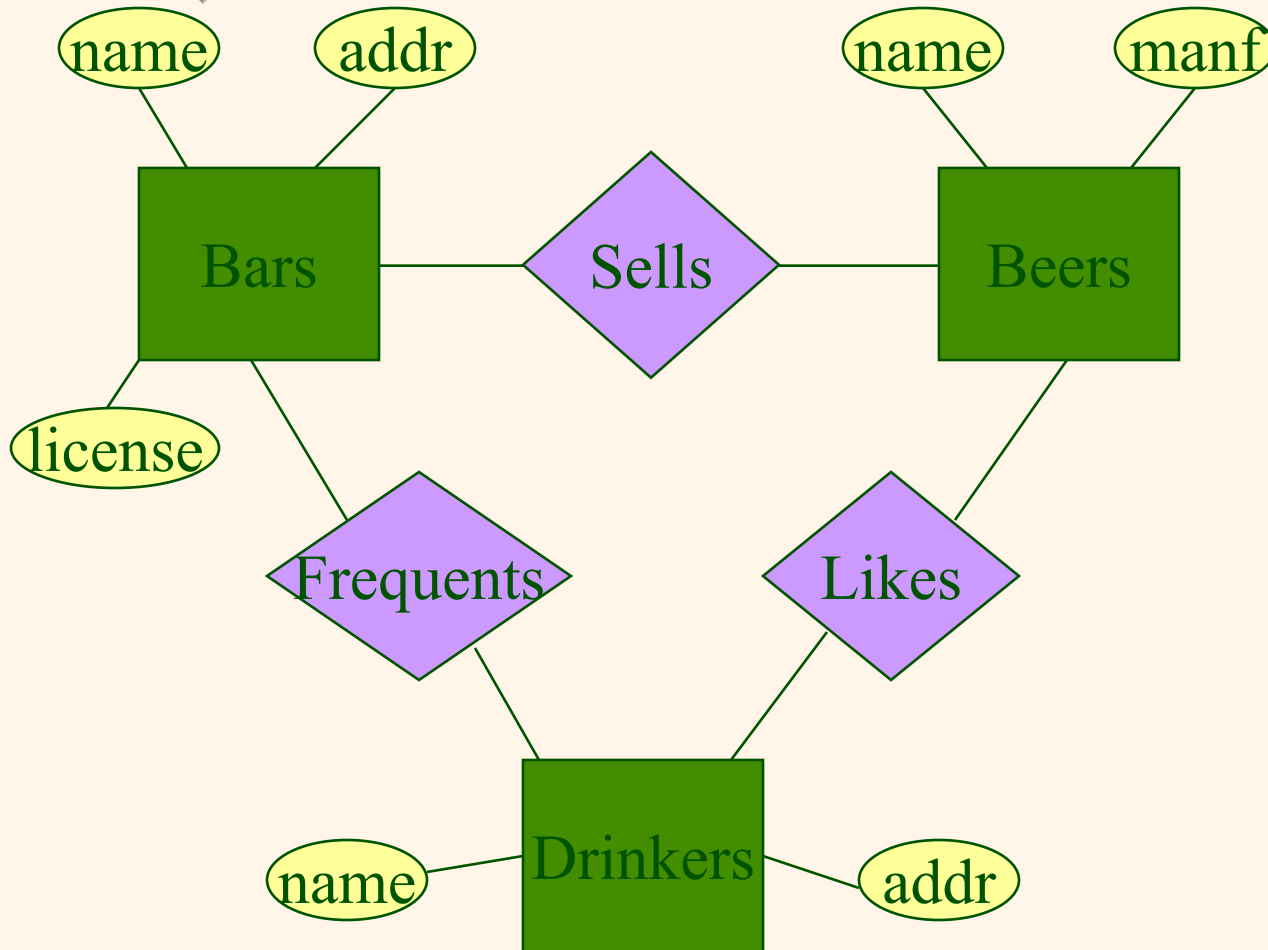
- ❖ Entity set **Beers** has two attributes, **name** and **manf** (manufacturer).
- ❖ Each **Beers** entity has values for these two attributes, e.g. (Bud, Anheuser-Busch)



# *Relationships*

- ❖ A relationship connects two or more entities.
- ❖ Analogous to entity sets, there is a notion of relationship sets.
- ❖ A relationship set is represented by a diamond, with lines to each of the entity sets involved.

# Example



Bars sell some beers.

Drinkers like some beers.

Drinkers frequent some bars.



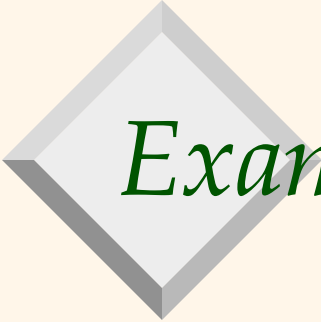
# *Entity Sets, Relationship Sets*

- ❖ The current “value” of an entity set is the set of entities that belong to it.
  - Example: the set of all bars in our database.
  - *Instance* of an entity set



# *Entity Sets, Relationship Sets*

- ❖ The “value” of a relationship is a *relationship set*, a set of tuples with one component for each related entity set
  - *Instance* of a relationship set
  - Two entities can only participate in each relationship together once (set, not multiset)



## *Example: Relationship Set Instance*

- ❖ For **Sells**, we might have a relationship set instance like this:
  - Notice no duplicates (since it's a *set*)

bar	beer
Joe's Bar	Bud
Joe's Bar	Miller
Sue's Bar	Bud
Sue's Bar	Pete's Ale
Sue's Bar	Heineken

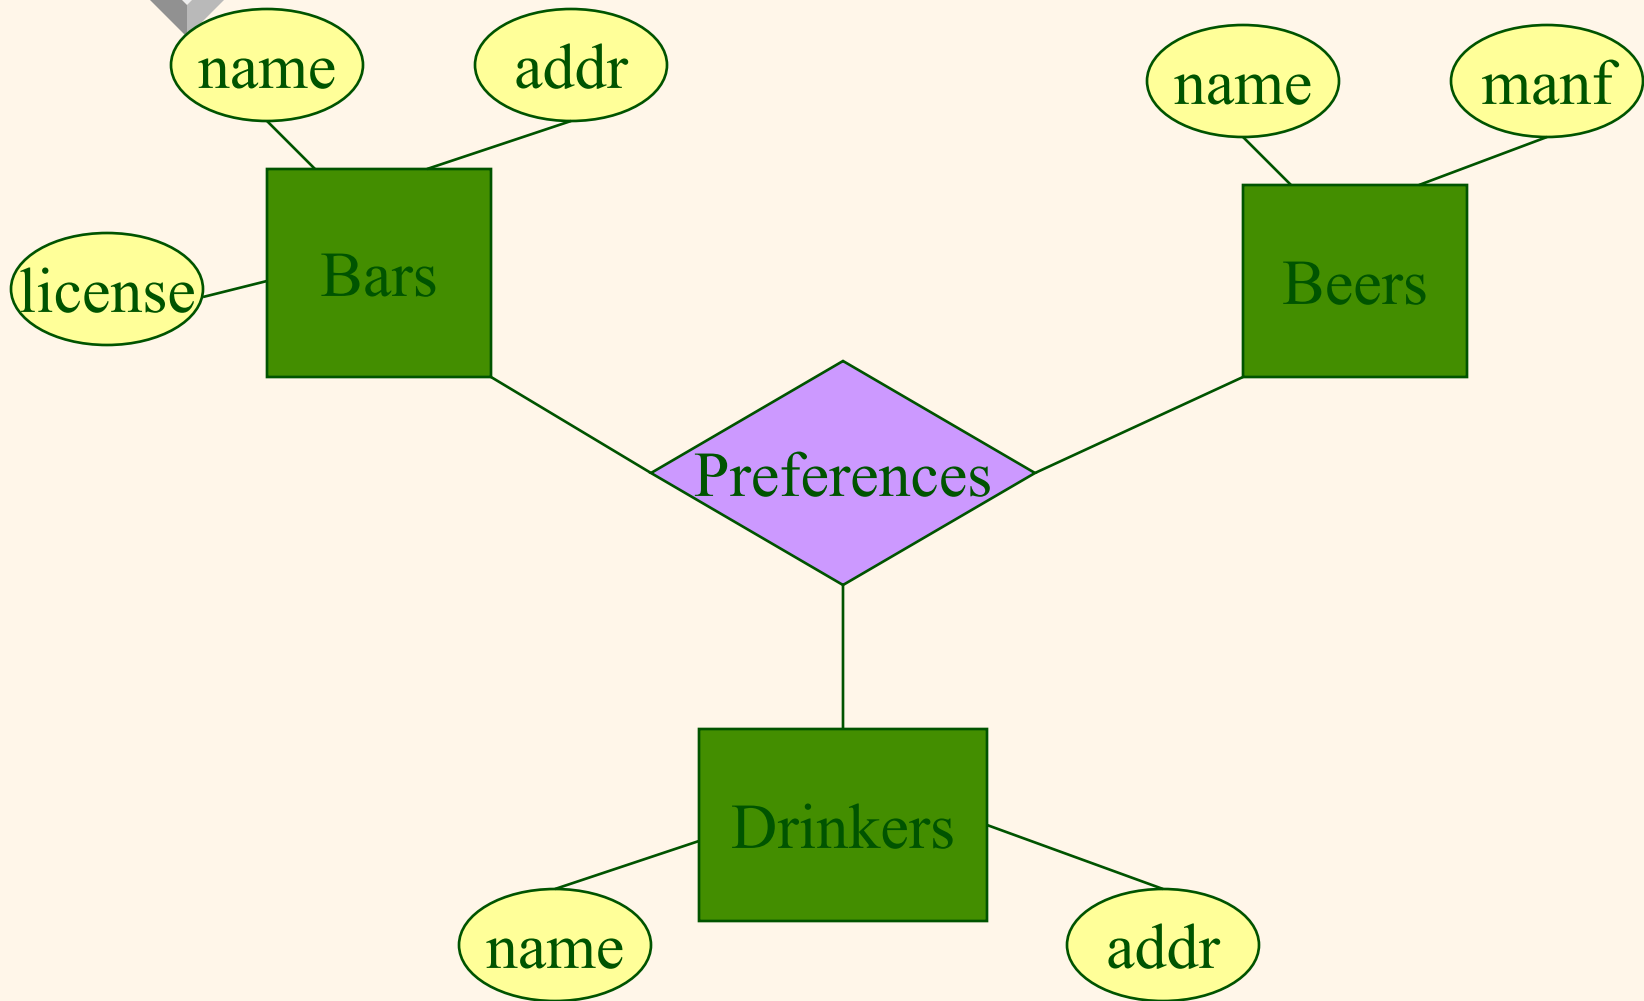




# *Multiway Relationships*

- ❖ Sometimes, we need a relationship that connects more than two entity sets.
- ❖ Suppose that drinkers will only drink certain beers at certain bars.
  - Our three binary relationships **Likes**, **Sells**, and **Frequents** do not allow us to make this distinction.
  - But a 3-way relationship would.

# *Example: 3-Way Relationship*





# *A Relationship Set Instance*

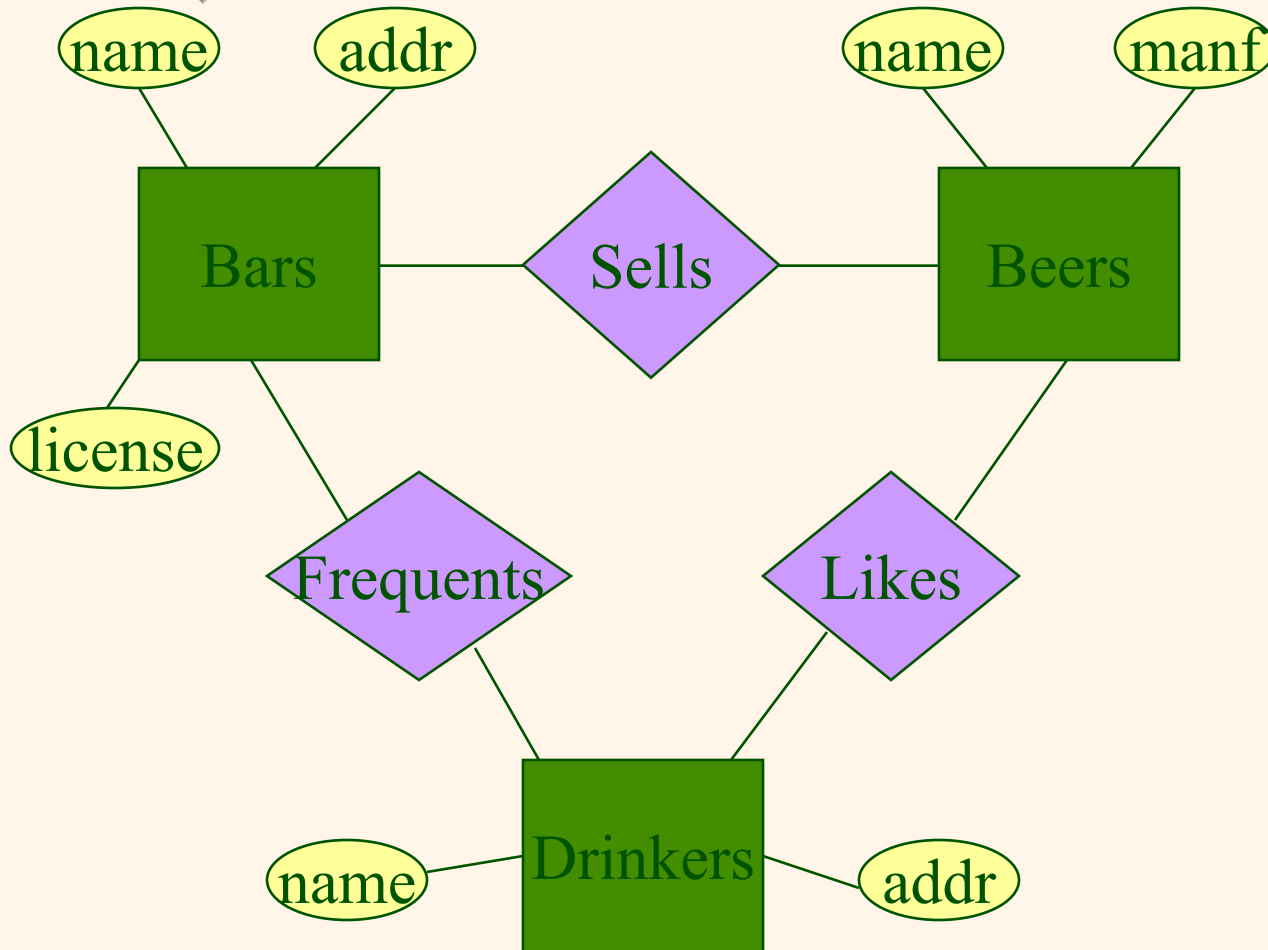
bar	drinker	beer
Joe's Bar	Ann	Miller
Sue's Bar	Ann	Bud
Sue's Bar	Ann	Pete's Ale
Joe's Bar	Bob	Miller
Joe's Bar	Bob	Bud
Sue's Bar	Cathy	Pete's Ale



# *Recap*

- ❖ Conceptual design = process of translating requirements into a formal(ish) model
- ❖ ER (Entity-Relationship) diagrams

# *Example*



Bars sell some beers.

Drinkers like some beers.

Drinkers frequent some bars.




# *Many-Many Relationships*

- ❖ Focus: **binary** relationships, such as **Sells** between **Bars** and **Beers**.
- ❖ In a many-many relationship, an entity of either set can be connected to many entities of the other set.
  - a bar sells many beers; a beer is sold by many bars.



# *Many-One Relationships*

- ❖ Some binary relationships are many -one from one entity set to another.
- ❖ Each entity of the first set is connected to at most one entity of the second set.
- ❖ But an entity of the second set can be connected to zero, one, or many entities of the first set.




## *Example: Many-One Relationship*

- ❖ Favorite, from Drinkers to Beers is many-one.
- ❖ A drinker has at most one favorite beer.
- ❖ But a beer can be the favorite of any number of drinkers, including zero.




# Example: Many-One Relationship





# One-One Relationships

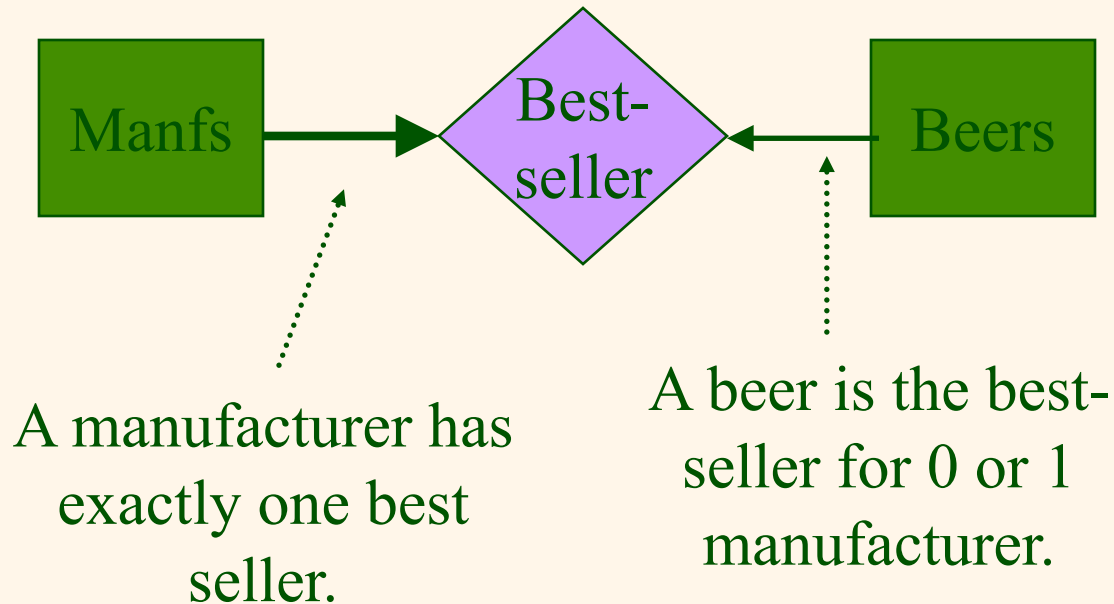
- ❖ In a one-one relationship, each entity of either entity set is related to at most one entity of the other set.
- ❖ **Example:** Relationship **Best-seller** between entity sets **Manfs** (manufacturer) and **Beers**.
  - A beer cannot be made by more than one manufacturer, and no manufacturer can have more than one best-seller (assume no ties).



# *Participation constraints*

- ❖ Consider **Best-seller** between **Manfs** and **Beers**.
- ❖ Some beers are not the best-seller of any manufacturer
- ❖ But a beer manufacturer has to have a best-seller
  - Participation constraint – thick line

# *In the ER Diagram*






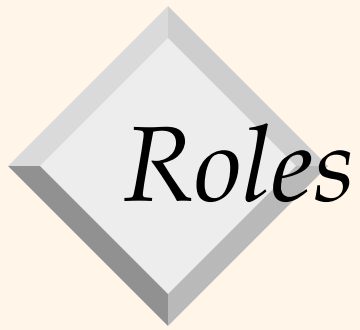
# *Attributes on Relationships*

- ❖ Sometimes it is useful to attach an attribute to a relationship.
- ❖ Think of this attribute as a property of tuples in the relationship set.

## *Example: Attribute on Relationship*



Price is a function of both the bar and the beer,  
not of one alone.




# *Roles*

- ❖ Sometimes an entity set appears more than once in a relationship.
- ❖ Label the edges between the relationship and the entity set with names called roles.

# *Example: Roles*

supervisor	subordinate
Bob	Ann
Joe	Sue








# Keys

- ❖ A key is a set of attributes for one entity set such that no two entities in this set agree on all the attributes of the key.
  - It is allowed for two entities to agree on some, but not all, of the key attributes.
- ❖ We must designate a key for every entity set

# Keys in ER diagrams

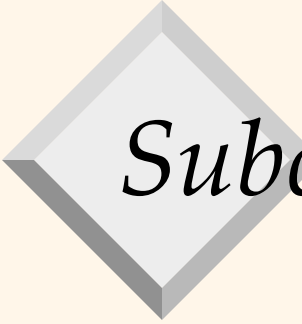


- Note that time and room might also be a **candidate** key, but we must select only one **primary** key.



# Keys

- ❖ In a relationship set, the keys of all the entities involved form a superkey
- ❖ This means that any specific pair (triple, etc) of entities can only appear in the instance of a relationship set once



# Subclasses


- ❖ *Subclasses* allow us to further distinguish between entities within an entity set.
- ❖ Example: Ales are a kind of beer.
  - Not every beer is an ale, but some are.
  - Let us suppose that in addition to all the *properties* (attributes and relationships) of beers, ales also have the attribute color.



# *Subclasses in E/R Diagrams*

- ❖ Assume subclasses form a tree.
  - I.e., no multiple inheritance.
- ❖ Isa triangles indicate the subclass relationship.
  - Point to the superclass.

# *Example*






# *Reasons to have subclasses*

- ❖ Add attributes that only make sense in some cases
  - E.g. the ale example
- ❖ “Group” entities together because they participate in the same relationship
  - E.g. Motorboats and Cars have very different attributes, but both are licensed to Owners, so want them to participate in same Licensed-To relationship.

# Weak Entities

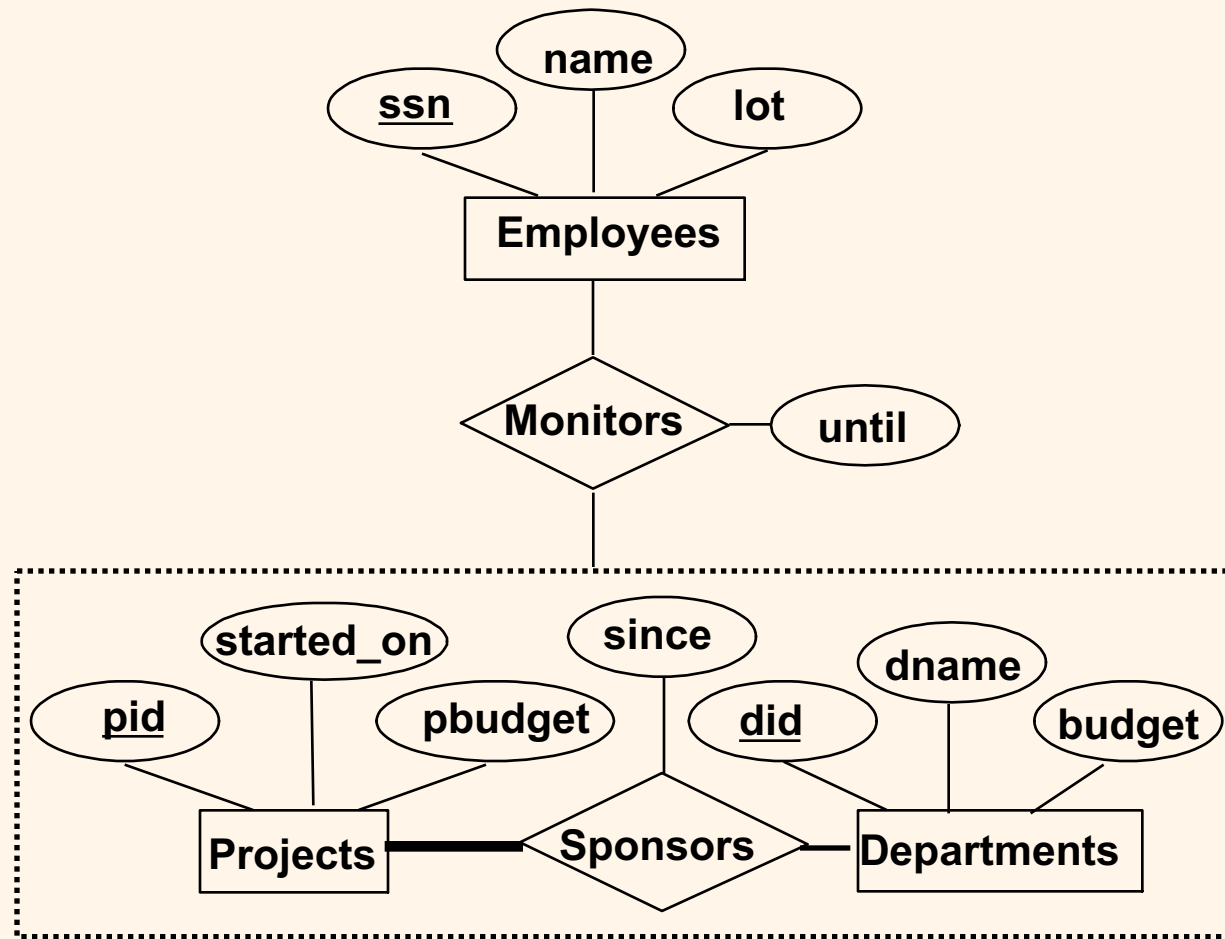
- ❖ A weak entity can be identified uniquely only by considering the primary key of another (*owner*) entity.
  - Owner entity set and weak entity set must participate in a one-to-many relationship set (one owner, many weak entities).
  - Weak entity set must have total participation in this identifying relationship set.
  - *linenum* is a partial key for LineItem






# Aggregation

- ❖ Used when we have to model a relationship involving (entity sets and) a *relationship set*.
  - Aggregation allows us to treat a relationship set as an entity set for purposes of participation in (other) relationships.



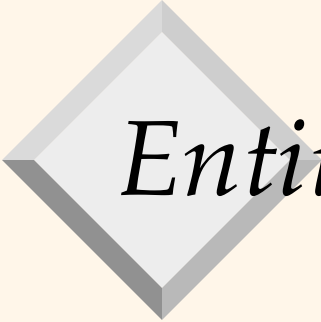
## ➡ Aggregation vs. ternary relationship:

- ❖ Monitors is a distinct relationship, with a descriptive attribute.
- ❖ Also, can say that each sponsorship is monitored by at most one employee.



# *Design choices*


- ❖ Should a concept be modeled as an entity or an attribute?
- ❖ Should a concept be modeled as an entity or a relationship?



# *Entity vs. Attribute*

- ❖ Should address be an attribute of Employees or an entity (connected to Employees by a relationship)?
- ❖ Depends:
  - If we have several addresses per employee, address must be an entity (since attributes cannot be set-valued)
  - If the structure (city, street, etc.) is important, e.g., we want to retrieve employees in a given city, address must be modeled as an entity (since attribute values are atomic)


# *More design choices*



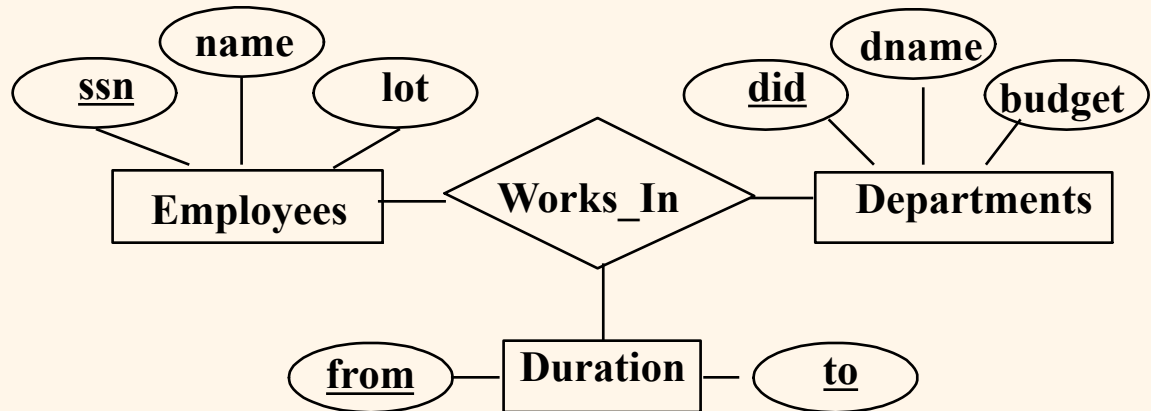
- Does not allow an employee to work in a department for two or more periods

# More design choices


- ▶ Want to record several values of the descriptive attributes for each instance of this relationship



- ▶ Accomplished by introducing new entity set, **Duration**

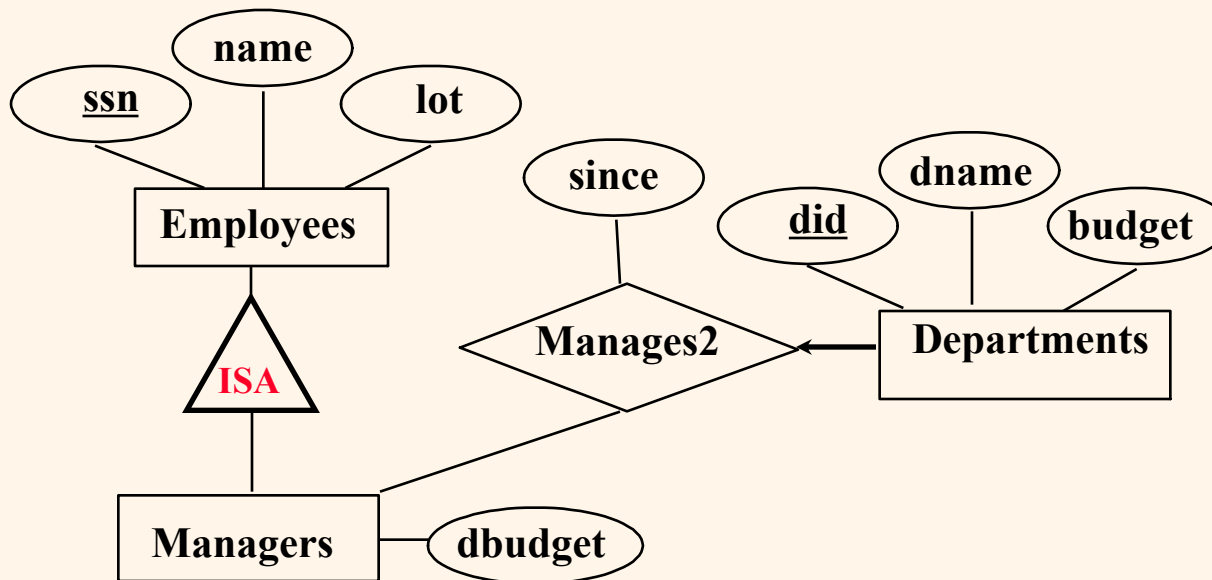




# More design choices



- ▶ What if a manager gets a discretionary budget that covers *all* managed depts?
  - **Redundancy:** dbudget stored for each dept managed by manager
  - **Misleading:** Suggests dbudget associated with department-mgr combination

# More design choices






# *From ER Diagrams to Relations*

- ❖ Entity set  $\rightarrow$  relation.
  - Attributes  $\rightarrow$  attributes.
- ❖ Relationships  $\rightarrow$  relations whose attributes are only:
  - The keys of the connected entity sets.
  - Attributes of the relationship itself.
  - Though we can do better sometimes




# *Entity Set -> Relation*



Relation: Beers(name, manf)

# Relationship -> Relation



Likes(drinker, beer)

Favorite(drinker, beer)

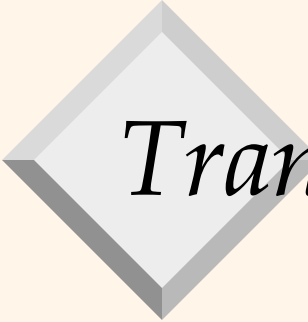
Buddies(name1, name2)



# *Translating relationships to SQL*

- ❖ Likes(drinker,beer)
  - Many-many
  - Optional on both sides

```
CREATE TABLE LIKES (drinker VARCHAR(15),  
                    beer VARCHAR(15),  
                    PRIMARY KEY (drinker, beer),  
                    FOREIGN KEY drinker references DRINKERS,  
                    FOREIGN KEY beer references BEERS);
```




# *Translating constraints to SQL*

## ❖ Favorite(drinker,beer)

- One-many
- Optional on both sides


```
CREATE TABLE FAVORITES (drinker VARCHAR(15),  
                           beer VARCHAR(15),  
                           PRIMARY KEY drinker,  
                           FOREIGN KEY drinker references DRINKERS,  
                           FOREIGN KEY beer references BEERS);
```



# *Translating constraints to SQL*

- ❖ Favorite(drinker,beer)
  - One-many
  - Optional on both sides
- ❖ Better: add favorite beer name to drinker table

```
CREATE TABLE DRINKER(name VARCHAR(15),  
                      address VARCHAR(20),  
                      fav_beer VARCHAR(15),  
                      PRIMARY KEY name,  
                      FOREIGN KEY fav_beer references BEERS);
```




# *Translating constraints to SQL*


## ❖ Favorite(drinker,beer)

- One-many
- What if every drinker must have a favorite?

```
CREATE TABLE DRINKER(name VARCHAR(15),  
                      address VARCHAR(20),  
                      fav_beer VARCHAR(15) NOT NULL;  
                      PRIMARY KEY name,  
                      FOREIGN KEY fav_beer references BEERS  
                      ON DELETE NO ACTION);
```

# *Translating subclasses*





# *Object-Oriented*

## Beers

name	manf
Bud	Anheuser-Busch

## Ales

name	manf	color
Summerbrew	Pete's	dark

Good for queries like “find the color of ales made by Pete’s.”

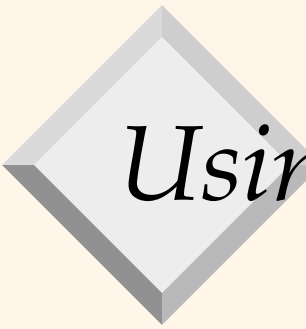


# *ER style*

name	manf
Bud	Anheuser-Busch
Summerbrew	Pete's

name	color
Summerbrew	dark

Good for queries like “find all beers  
(including ales) made by Pete's.”




# Using NULLs

name	manf	color
Bud	Anheuser-Busch	NULL
Summerbrew	Pete's	dark

Saves space unless there are *lots* of attributes that are usually NULL.

## ER to Relational (contd.)

- ❖ A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.





# *Translating Weak Entity Sets*


- ❖ Weak entity set and identifying relationship set are translated into a single table.
  - When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Dep_Policy ( pname CHAR(20),  
                           age INTEGER,  
                           cost REAL,  
                           ssn CHAR(11) NOT NULL,  
                           PRIMARY KEY (pname, ssn),  
                           FOREIGN KEY (ssn)  
                               REFERENCES Employees,  
                           ON DELETE CASCADE);
```



## *In practice*

- ❖ You might not use ER diagrams, but you will be using some sort of diagrams as your data gets more complex
- ❖ Any nontrivial application will have a moderate number of entities and relationships
- ❖ GUI tools to help you create diagrams and translate between diagrams and relational schema (both directions)
  - E.g. MySQL Workbench



## *Summary so far*

- ❖ Modeling your data with ER diagrams
- ❖ Translating ER diagrams to the relational model and to SQL CREATE statements
- ❖ Next: improving the relational tables you have generated to remove redundancy