

Concurrency Control - Formal Foundations



Last time

- ❖ Intro to ACID transactions
- ❖ Focus on Isolation
 - Every transaction has the illusion of having the DB to itself
- ❖ Isolation anomalies
 - bad things that can happen due to interleaving



Transaction 1

UPDATE accounts


SET amount= amount - 100

WHERE name = 'Alice';

UPDATE accounts

SET amount= amount + 100

WHERE name = 'Bob';



Transaction 2

UPDATE accounts

SET amount= amount*1.1

WHERE name = 'Alice';

UPDATE accounts

SET amount= amount*1.1

WHERE name = 'Bob';

Bad interleaving #1

Transaction 1	Transaction 2
UPDATE accounts SET amount= amount- 100 WHERE name = 'Alice';	
	UPDATE accounts SET amount= amount*1.1 WHERE name = 'Alice';
	UPDATE accounts SET amount= amount*1.1 WHERE name = 'Bob';
UPDATE accounts SET amount= amount+ 100 WHERE name = 'Bob';	



Bad interleaving #1

- ❖ Let's introduce some notation to make it easier to write these down
- ❖ Transactions work on objects A, B, C
- ❖ Transactions read (R) and Write (W) objects
 - UPDATE accounts
 - SET amount= amount- 100
 - WHERE name = 'Alice';
- ❖ This becomes R1(A) W1(A)
 - A = tuple with Alice's info
 - R because of reference to amount on RHS
 - W because we change amount
 - 1 because we're in transaction 1

Bad interleaving #1

Transaction 1	Transaction 2
UPDATE accounts SET amount= amount- 100 WHERE name = 'Alice';	
	UPDATE accounts SET amount= amount*1.1 WHERE name = 'Alice';
	UPDATE accounts SET amount= amount*1.1 WHERE name = 'Bob';
UPDATE accounts SET amount= amount+ 100 WHERE name = 'Bob';	

Bad interleaving #1

Transaction 1	Transaction 2
R1(A) W1(A)	
	R2(A) W2(A)
	R2(B) W2(B)
R1(B) W1(B)	

- More shorthand: R1(A) W1(A) becomes RW1(A)



Reading uncommitted data

- ❖ What is the essence of the problem?
 - RW1(A) RW2(A) RW2(B) RW1(B)
 - Transaction 2 has read the value of A written by Transaction 1 before Transaction 1 committed
 - Who knows whether Transaction 1 was done modifying A?
 - In this case it was, but it needed to make a corresponding change to B which 2 didn't see
- ❖ Another variant – read from someone that later aborts
 - RW1(A) RW2(A) Abort1 RW2(B) Commit2



Dirty read

- ❖ Reading a value modified by an uncommitted transaction
 - Not **guaranteed** to cause a problem in all cases, but
 - Particularly bad if the transaction you read from aborts after your read
 - But you see it can be bad even with no aborts



Dirty read

- ❖ Why bad:
 - Don't know if transaction is done changing the value
 - Or other values that are related to it by an invariant
 - If transaction aborts, we are really in trouble.




Unrepeatable read

- ❖ The second isolation anomaly
- ❖ Even if we never read uncommitted data, can still have a different problem
- ❖ R1(A) W2(A) C2 R1(A)
- ❖ Read the same data item twice, get different answer
 - Real world example: first read checks if enough inventory available, second read is part of decrement operation on inventory



Lost updates


- ❖ The third isolation anomaly
- ❖ Fundamentally related to writes (not reads)



Transaction 3 – $W(A)$ $W(B)$

```
UPDATE accounts  
SET amount = 100  
WHERE name = 'Alice';
```

```
UPDATE accounts  
SET amount = 100  
WHERE name = 'Bob';
```



Transaction 4 – W(A) W(B)

UPDATE accounts

SET amount = 200

WHERE name = 'Alice';

UPDATE accounts

SET amount = 200

WHERE name = 'Bob';

Lost updates

- ❖ Suppose transactions interleave as follows:
- ❖ $W3(A)$ $W4(A)$ $W4(B)$ $W3(B)$
 - At the end Alice has \$200 and Bob has \$100, so the amounts are not equal...
 - If we had equality as a consistency constraint, it has just been violated
- ❖ Although no-one read an inconsistent state, somehow some writes got mixed up giving an inconsistent DB at the end



The three isolation anomalies

❖ Dirty read

- Read data someone else is not finished with
- RW1(A) RW2(A) RW2(B) RW1(B)

❖ Unrepeatable read

- Data is changing under you
- R1(A) W2(A) C2 R1(A)

❖ Lost Update

- Writes are otherwise “correct” but destroy each other in a bad way
- W3(A) W4(A) W4(B) W3(B)



Avoiding Anomalies

- ❖ How do we know what all the possible anomalies are?
 - So that we can design an algorithm to prevent them...
- ❖ Are some of them more serious than others?



A more principled approach

❖ Definition: **transaction schedule**

- A sequence of operations performed by a set of transactions on the DB
- e.g. W1(A) W2(A) W2(B) W1(B) Commit2 Commit1

❖ A transaction schedule is a concrete interleaving

- We can identify “good” and “bad” schedules
- E.g. if exhibits one of our anomalies, bad!!
- Once we know what's "good" and "bad", we can start understanding how to avoid "bad" schedules

Serializability

- ❖ Intuition: suppose transactions in the schedule had executed serially
 - e.g. $W1(A)$ $W1(B)$ $C1$ $W2(A)$ $W2(B)$ $C2$
- ❖ Then nothing bad could have happened!
 - Alice and Bob banking example

Bad interleaving

Transaction 1	Transaction 2
UPDATE accounts SET amount= amount- 100 WHERE name = 'Alice';	
	UPDATE accounts SET amount= amount*1.1 WHERE name = 'Alice';
	UPDATE accounts SET amount= amount*1.1 WHERE name = 'Bob';
UPDATE accounts SET amount= amount+ 100 WHERE name = 'Bob';	




Serializability

- ❖ Idea: compare our schedule of interest to a **serial schedule** with the same transactions
 - We want to "simulate serial execution"



Serializability

- ❖ First requirement: a good schedule is one that produce the same final DB as some serial schedule (with same transactions)
 - if doesn't produce same final DB as any serial schedule, this is pretty iffy!
 - ◆ How can we have any guarantee the DB is consistent in this case?



Final-state serializability

- ❖ A schedule S is **final-state serializable** if there is some serial schedule S' involving the same transactions such that the **final DB is the same after executing S and S'**
 - No matter what DB you start with
- ❖ Example: $W1(A)$ **$W2(A)$** $W1(B)$ **$W2(B)$** $C1$ **$C2$**



Do we need something stronger?

- ❖ R1(A) W2(A) R1(A) C1 C2
- ❖ Unrepeatable read...
- ❖ But final DB produced is the same as after serial execution where 2 executes before 1!
- ❖ Apparently requiring same final DB may not be enough!
 - Reads matter too
 - T1 should see the same value of A if it reads it twice (without making any changes of its own)



View Equivalence

- ❖ Intuition: If transaction **i** sees a value written by transaction **j** in schedule S1 (it is "tainted" or "contaminated" by **j**), it should also see value written by transaction **j** in Schedule S2
- ❖ There is a **dependency** (an **ordering**) between **j** and **i**
- ❖ Two schedules are "similar" if they have all the same dependencies between transactions

View Equivalence

- ❖ Schedules S1 and S2 are **view equivalent** if:
 - If T_i reads initial value of A in S1, then T_i also reads initial value of A in S2
 - If T_i reads value of A written by T_j in S1, then T_i also reads value of A written by T_j in S2
 - If T_i writes final value of A in S1, then T_i also writes final value of A in S2

T1: R(A)	W(A)
T2: W(A)	
T3:	W(A)

T1: R(A),W(A)	
T2: W(A)	
T3:	W(A)



View Serializability

- ❖ A schedule S is view serializable if it is view-equivalent to some serial schedule
- ❖ Let's see why our "offending" schedule is not view serializable...
- ❖ $R1(A)$ $W2(A)$ $R1(A)$ $C1$ $C2$



So what now?

- ❖ Want to allow only **view-serializable** schedules (interleavings) in our system
- ❖ Will need some algorithm to achieve this
- ❖ Useful to define a third kind of serializability
 - **conflict-serializability**
- ❖ It is NP-complete to test whether a schedule is view-serializable
 - But conflict-serializability very easy to check efficiently