


Concurrency Control




Last time

- ❖ Isolation anomalies
 - Dirty read, unrepeatable read, lost update
- ❖ Towards defining "correct" execution
 - Transaction schedules
 - ◆ e.g. R1(A) W2(B) R2(C)
 - "Good" schedules have some equivalence to a serial schedule with same transaction

Big Picture (all inclusions are proper)

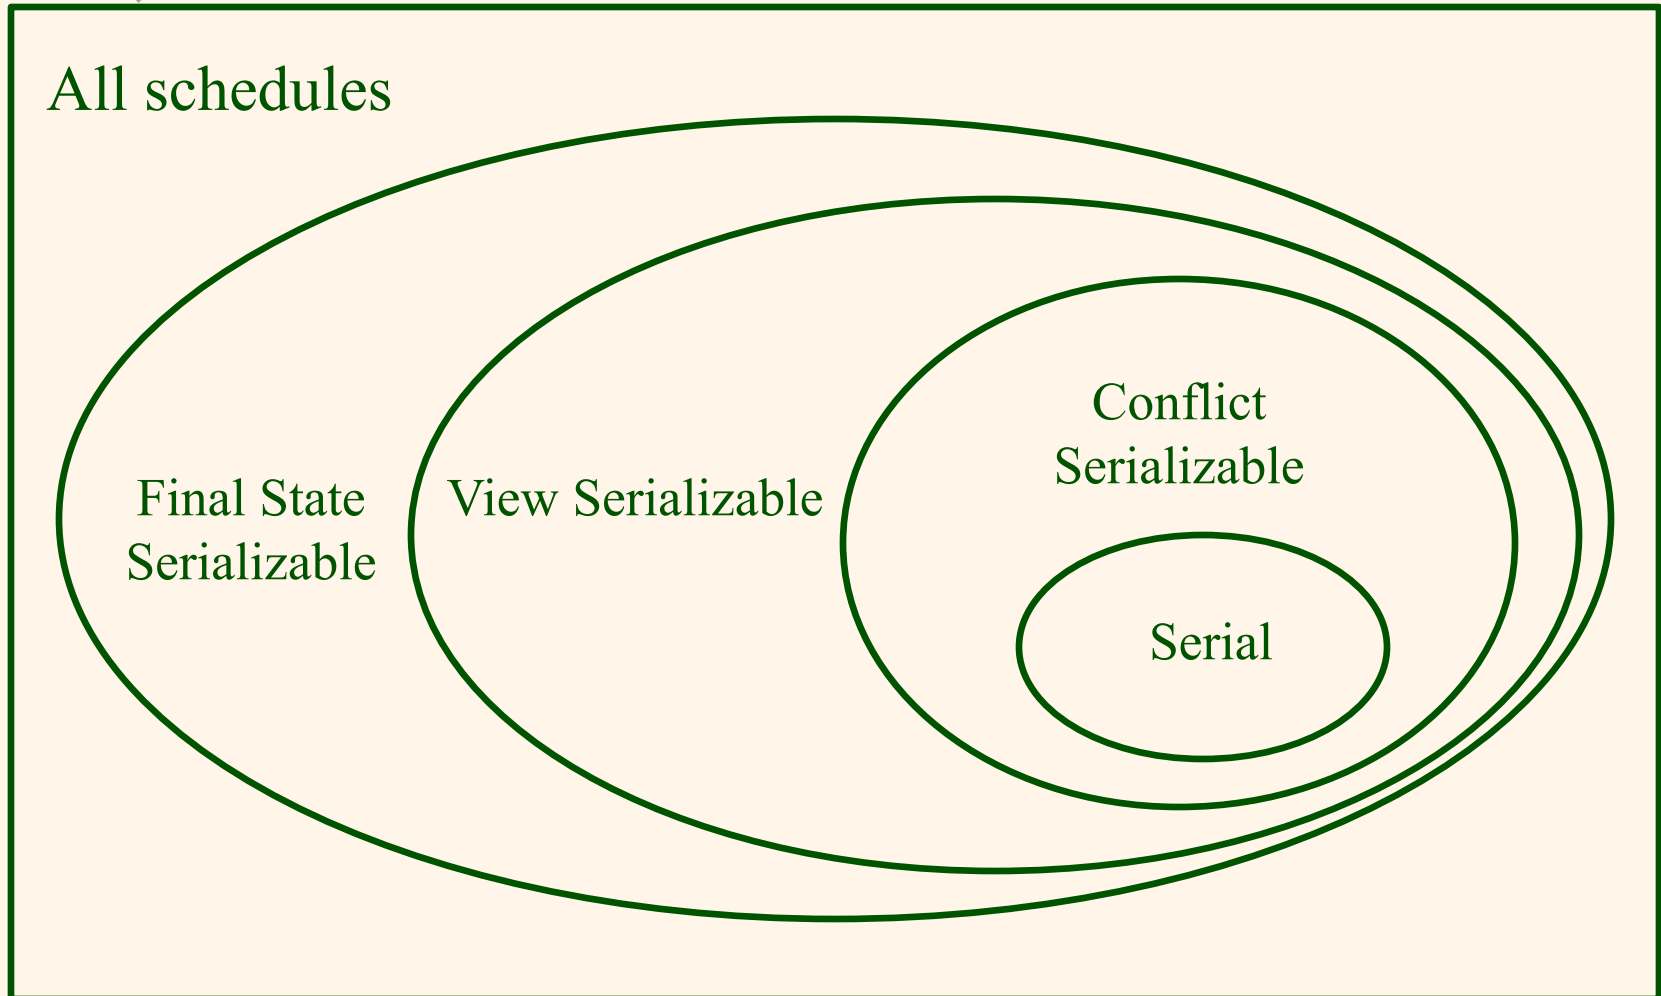




Another notion of serializability

- ❖ View-serializability is nice but difficult to check/enforce in practice
- ❖ A third notion: **conflict-serializability**
- ❖ This can be enforced efficiently

Big Picture (all inclusions are proper)





Conflicting Operations

- ❖ Based on identifying **conflicting pairs of operations** between transactions
- ❖ Observation: read-only transactions cannot interfere with each other in any bad way!
- ❖ The only problems are due to writes
 - All three of our anomalies involve a write somewhere



Conflicting Operations

- ❖ Two operations by different transactions on the same object **conflict if at least one is a write**
 - WW
 - WR
 - RW
- ❖ Example: if T1 and T2 both write Alice's bank account balance, they conflict
- ❖ But if T1 only writes Alice's balance and T2 only writes Bob's, no conflict!
- ❖ Also, two writes by the same transaction never conflict with each other



Conflicting Operations

- ❖ Conflicting operations **do not commute**
 - $W1(A) W2(A)$
 - $R1(A) W2(A)$
 - $W1(A) R2(A)$
- ❖ So, they induce some notion of precedence or ordering




Conflict Graphs

- ❖ Given a schedule, can identify all conflicting pairs of operations and represent them as a graph
- ❖ Nodes are transactions
- ❖ Edge from i to j if transaction i contains an operation that **conflicts with and precedes** (in the schedule) an operation by transaction j
- ❖ Example: $R1(A)$ **$W2(A)$** $R1(A)$

Conflict Graphs


❖ $R1(A)$ $W2(A)$ $R1(A)$






A conflict graph exercise

R1(A) R2(A) R1(C) W1(A) R3(C) W2(B) W3(B) W3(C)



A conflict graph exercise


R1(A) R2(A) R1(C) W1(A) R3(C) W2(B) W3(B) W3(C)






Conflict Serializability

- ❖ A schedule is conflict serializable if its conflict graph contains no cycle
- ❖ Alternative (equivalent) statement: it is conflict serializable if it has the same conflict graph as some serial schedule
 - Why are these equivalent?
- ❖ Topological sort on the conflict graph gives us equivalent serial execution



A conflict graph exercise

R1(A) R2(A) R1(C) W1(A) R3(C) W2(B) W3(B) W3(C)





An equivalent serial schedule

R1(A) R2(A) R1(C) W1(A) R3(C) W2(B) W3(B) W3(C)

- ❖ Executing the transactions in the order 2, 1, 3 produces the same DB
- ❖ Let's see if we believe it...

R2(A) W2(B) R1(A) R1(C) W1(A) R3(C) W3(B) W3(C)




A slightly different schedule

R1(A) R2(A) R1(C) W1(A) R3(C) W3(B) W2(B) W3(C)


Compare with previous:

R1(A) R2(A) R1(C) W1(A) R3(C) W2(B) W3(B) W3(C)




A conflict graph exercise

R1(A) R2(A) R1(C) W1(A) R3(C) W3(B) W2(B) W3(C)




Big Picture (all inclusions are proper)





What about aborts?


- ❖ Convention: if a schedule contains transactions that abort, we ignore the operations (R/W) of these transactions when determining if a schedule is serializable
- ❖ Rationale: aborted transactions' operations should not have happened and will be cleaned up
- ❖ Although, the "cleanup" can be nontrivial!



Dealing with aborts

❖ Cleanup issues

- W1(A) R2(A) W2(B) C2 A1 ... <= not good!
- 2 has read a value of A that "should not have been there"
- 2 has already committed, so extremely unclear how to clean up in a sensible way



Recoverability

- ❖ A schedule is **recoverable** if a transaction commits only after the commit of all transactions whose changes it has read
 - Never have to undo an *already committed* transaction due to someone else's abort
 - W1(A) R2(A) W2(B) C2 ... \leq not recoverable



Recoverability

- ❖ Enforcing this will be a problem - need to keep track of who has read from whom
 - So, more (stronger) criteria



ACA schedules

- ❖ A schedule avoids cascading aborts if a transaction never reads uncommitted writes
 - Now that is easier to enforce
 - $W1(A) \text{ } R2(A) \text{ } W2(B) \text{ } C1 \text{ } C2$ is recoverable but does not avoid cascading aborts
- ❖ But there could still be problems even if we enforce ACA



Strict schedules

- ❖ Consider schedule $W1(A) W2(A) C2 A1$
- ❖ Recoverable and avoids cascading aborts because no-one ever read anything
- ❖ But "cleanup" could be a pain
 - Need to retain write by T2 while erasing write by T1
 - Depending on what the objects are and what records you keep of changes, could be a lot of work
 - May need to roll back to initial state and redo T2



Strict schedule

- ❖ If we don't want this problem, can place an even stronger restriction
- ❖ A schedule is **strict** if whenever a transaction writes to a data item, no other transaction can read or write to that item until T has committed or aborted
- ❖ Avoids our problem $W1(A) W2(A) C2 A1$

Summary

All schedules

Recoverable

Avoid Cascading
Aborts

Strict

Serial



Summary

- ❖ Restrictions on schedules for serializability:
 - Final-state serializable
 - View-serializable
 - Conflict-serializable
- ❖ Restrictions on schedules for handling aborts:
 - Recoverable
 - Avoids Cascading Aborts
 - Strict



What was all this for?

- ❖ We now know how to tell **good** interleavings from **bad** ones
 - Want either view-serializability or conflict-serializability
 - And possibly something like ACA or strictness too
- ❖ The system needs to allow only **good** interleavings
- ❖ Let's see this in action in MySQL



Transaction 1

START TRANSACTION;

UPDATE accounts

SET amount= amount-100

WHERE name = 'Alice';

UPDATE accounts

SET amount= amount+100

WHERE name = 'Bob';

COMMIT;



Transaction 2

START TRANSACTION;

UPDATE accounts

SET amount= amount*1.1

WHERE name = 'Alice';

UPDATE accounts

SET amount= amount*1.1

WHERE name = 'Bob';

COMMIT;


Bad interleaving

Transaction 1	Transaction 2
UPDATE accounts SET amount= amount- 100 WHERE name = 'Alice';	
	UPDATE accounts SET amount= amount*1.1 WHERE name = 'Alice';
	UPDATE accounts SET amount= amount*1.1 WHERE name = 'Bob';
UPDATE accounts SET amount= amount+ 100 WHERE name = 'Bob';	




Enforcing conflict serializability

- ❖ Goal: system must only allow interleavings (schedules) that are (conflict) serializable.
- ❖ Still don't know **how** to do it
 - Specification vs implementation
 - There may be (and are!) different ways to do it



Protocols for enforcing a Property

- ❖ General systems discussion
- ❖ Common problem: need to enforce a **property**
 - All interleavings of transactions are conflict-serializable
 - No customer receives a product without having paid for it
 - If someone modifies your 4320 grade, that person is one of the 4320 course staff
- ❖ To enforce it, system uses an algorithm or **protocol**
 - E.g. require me to log in before I modify your grade



Protocols and properties

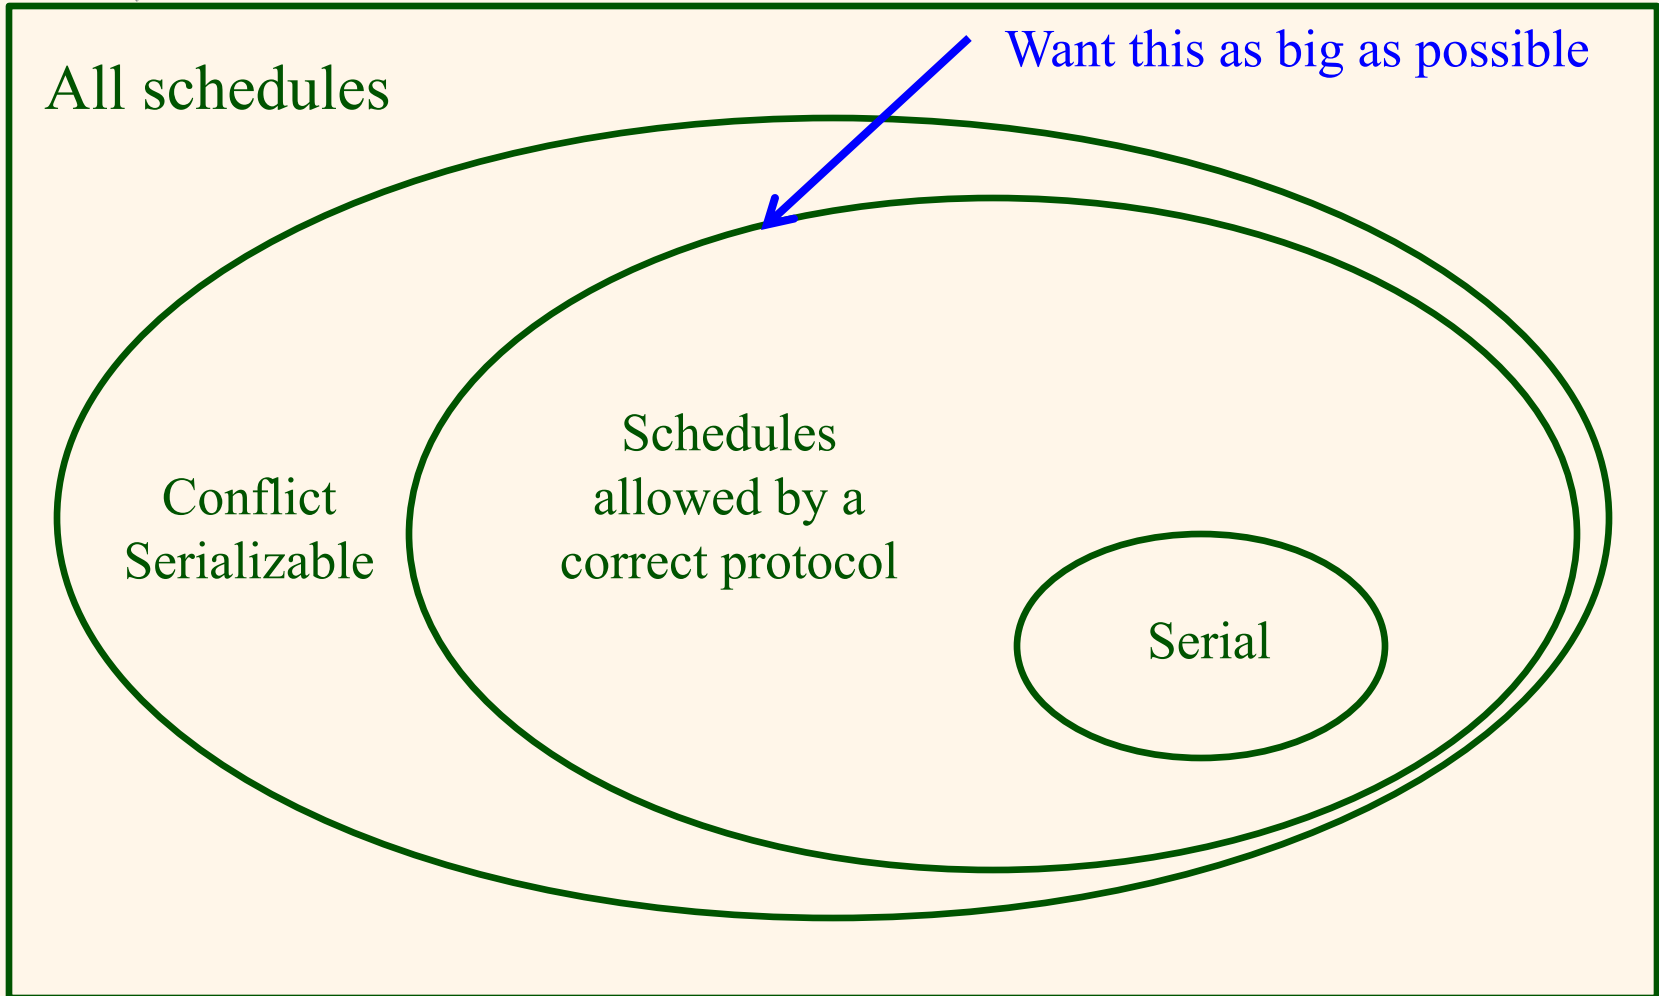
- ❖ There can be many different protocols for enforcing a property
- ❖ Some protocols can be stricter than necessary, but still correct
 - I.e. forbid all "bad" behavior
 - But forbid some "good" behavior too
 - Ok as long as **never allow "bad" behavior**
 - This often happens in practice



Back to transactions

- ❖ **Property:** conflict-serializability
- ❖ A variety of **protocols**
- ❖ Most of which are conservative (disallow some conflict serializable schedules)

Big Picture





Concurrency control protocols

- ❖ A very simple **protocol**: force all transactions to execute serially
 - No operations by different transactions may interleave
 - If another transaction is not done reading/writing to DB, you must wait
- ❖ Does this permit **only** conflict-serializable schedules?
- ❖ Does this permit **all** conflict-serializable schedules?



Reasonable protocols

- ❖ Should allow as many conflict-serializable schedules as possible
- ❖ Should be simple and not impose too much overhead of their own
 - Otherwise the performance boost from allowing interleavings is lost
- ❖ Several (families of) solutions