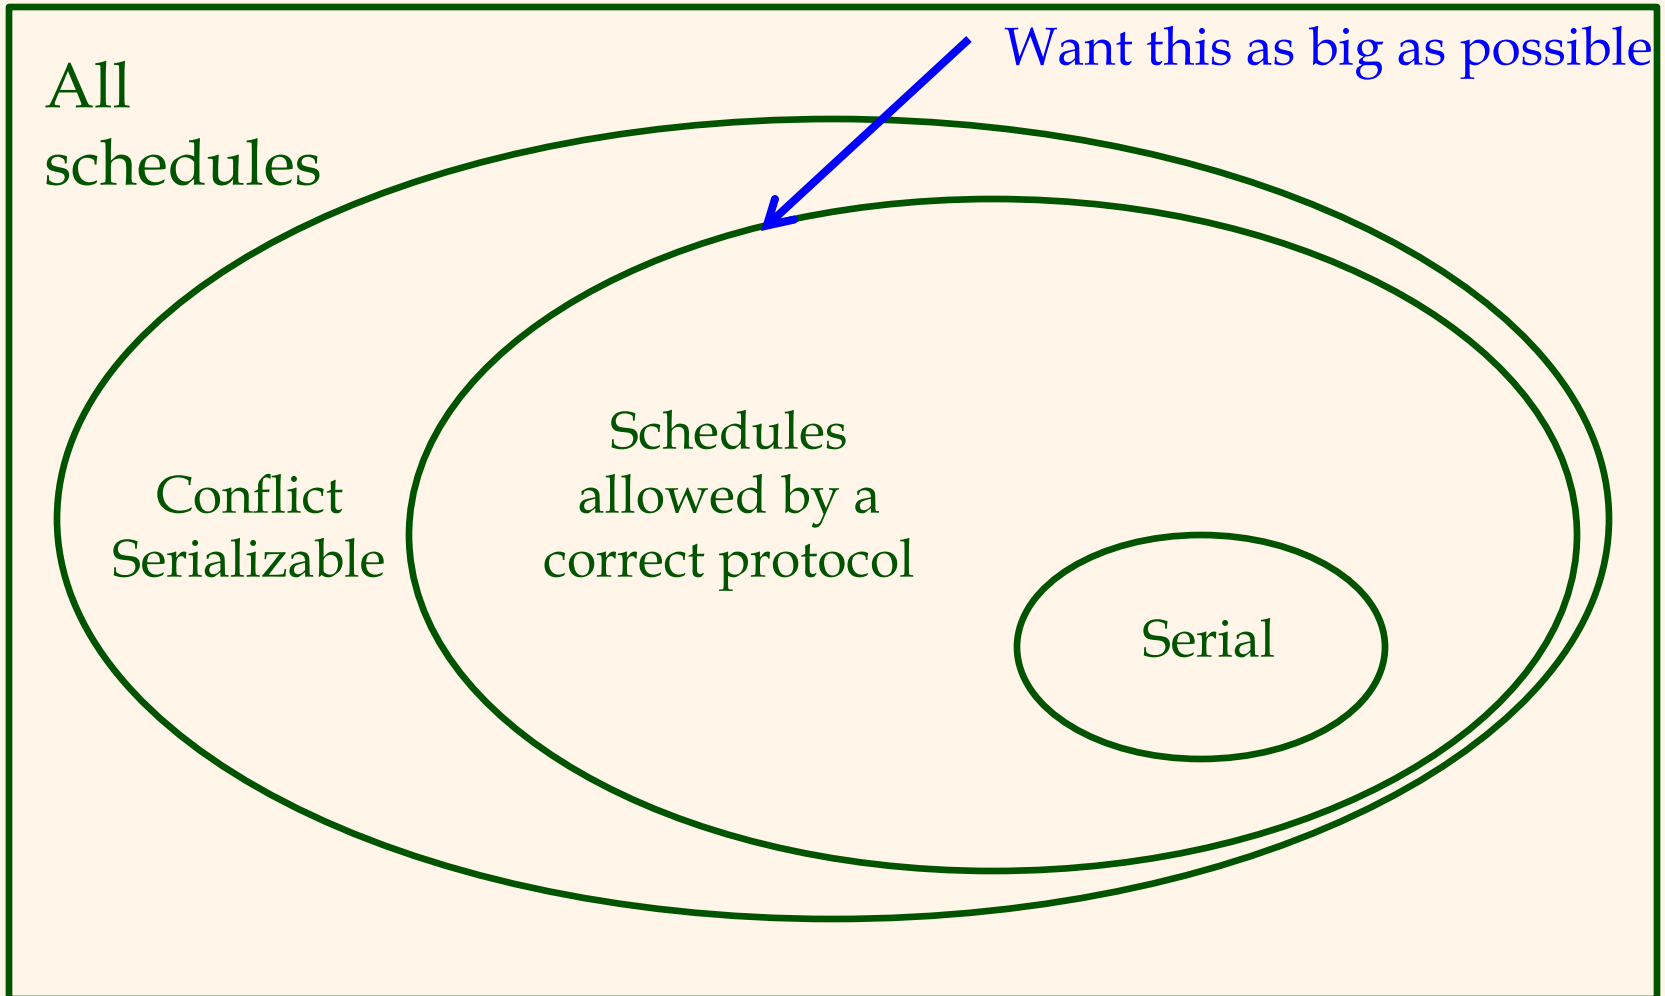
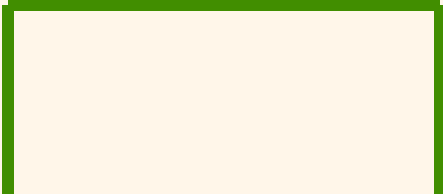
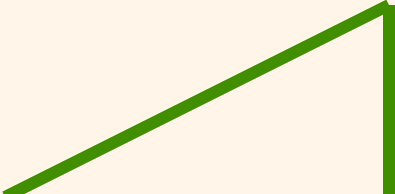
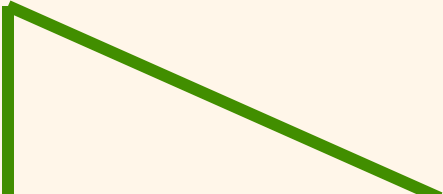
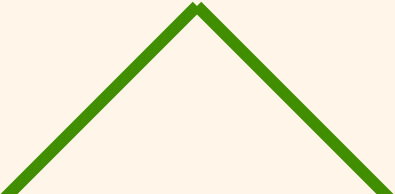


Locking continued

Concurrency control protocols



Last time: 2PL

<div>Conservative</div> <div>Strict</div>	Yes	No
Yes		
No		

A final theoretical note on 2PL

- ❖ Even non-strict, non-conservative 2PL is “too much” (forbids some interleavings that are ok)
 - Every interleaving/schedule permitted by 2PL is conflict-serializable
 - But not necessarily vice versa!
 - W1(A) R2(A) Commit2 R3(B) Commit3 W1(B) Commit1

In practice:

- ❖ Strict 2PL is the most commonly used locking protocol
 - Locks acquired as needed
 - Locks held until commit
- ❖ Means deadlock is an issue and must be dealt with

Deadlocks

- ❖ Deadlock: Cycle of transactions waiting for locks to be released by each other.
- ❖ Two ways of dealing with deadlocks:
 - Deadlock prevention
 - Deadlock detection

Deadlock Detection

- ❖ Create a **waits-for graph**:
 - Nodes are transactions
 - There is an edge from T_i to T_j if T_i is waiting for T_j to release a lock
- ❖ Periodically check for cycles in the waits-for graph
- ❖ If cycle found, abort one of the transactions (its locks get released allowing others to proceed)

Deadlock Detection

Example: (S = request shared lock, X = request exclusive lock)

T1: S(A), R(A), S(B)



T2: X(B), W(B)

T3: S(C), R(C)

T4: X(B)

X(C)

X(A)



Deadlock Prevention

- ❖ Assign priorities based on timestamps.
 - Lower timestamp (older transaction) => higher priority
- ❖ Assume T_i wants a lock that T_j holds. Two policies are possible:
 - **Wound-wait**: If T_i has higher priority, T_j aborts; otherwise T_i waits
 - **Wait-die**: If T_i has higher priority, T_i waits for T_j ; otherwise T_i aborts
- ❖ If a transaction re-starts, make sure it has its original timestamp to avoid starvation (note the higher-priority transaction is never aborted in either scheme)

Deadlock Prevention

❖ Advantage of wait-die:

- Non-preemptive (if I have all my locks will never be aborted for deadlock reasons)

❖ Disadvantage of wait-die:

- A younger transaction that conflicts with an older transaction may be repeatedly aborted for the same reason

Next

- ❖ So far have assumed
 - all operations are on a single object
 - all objects are independent
 - no objects are created or deleted
- ❖ If these assumptions fail
 - this may be bad - problems may arise
 - or good - there may be opportunities for more efficient locking protocols.


The Phantom Problem

- ❖ Suppose we have Strict 2PL and transaction 1 performs the following query **twice**:
 - `SELECT * FROM Widgets WHERE color = 'red';`
- ❖ In the interim, transaction 2 inserts a new red widget
 - Can do this – T1 only has locks on pre-existing red widgets!
- ❖ Loss of isolation even though we used 2PL!

Avoiding phantoms

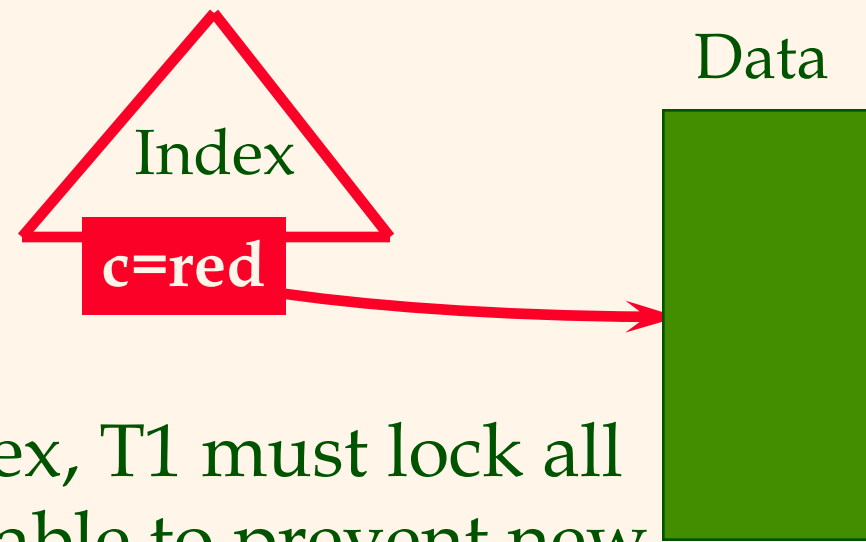
- ❖ The problem is not that 2PL is incorrect, the problem is that T1 did not "really" lock all red widgets
- ❖ Need way to lock a set of tuples by specifying all possible tuples that could be in the set
 - Predicate locking
 - Can be computationally expensive with arbitrary predicates
 - In practice, could be done with index locking

Index Locking




- ❖ If there is a index on the *color* field, T1 can lock the index page containing the data entries with *color* = red.
 - If there are no entries with *color* = red, T1 must lock the index page where such a data entry *would* be, if it existed!

Index Locking



- ❖ If there is no suitable index, T1 must lock all pages, and lock the file/table to prevent new pages from being added, to ensure that no new records with *color* = red are added.

Locking in B+ trees



Locking in B+ Trees

- ❖ How can we enable “safe” concurrent access to index structures?
- ❖ Note: this is useful even apart from index locking
 - If you have multiple transactions that need to search/insert/delete the tree

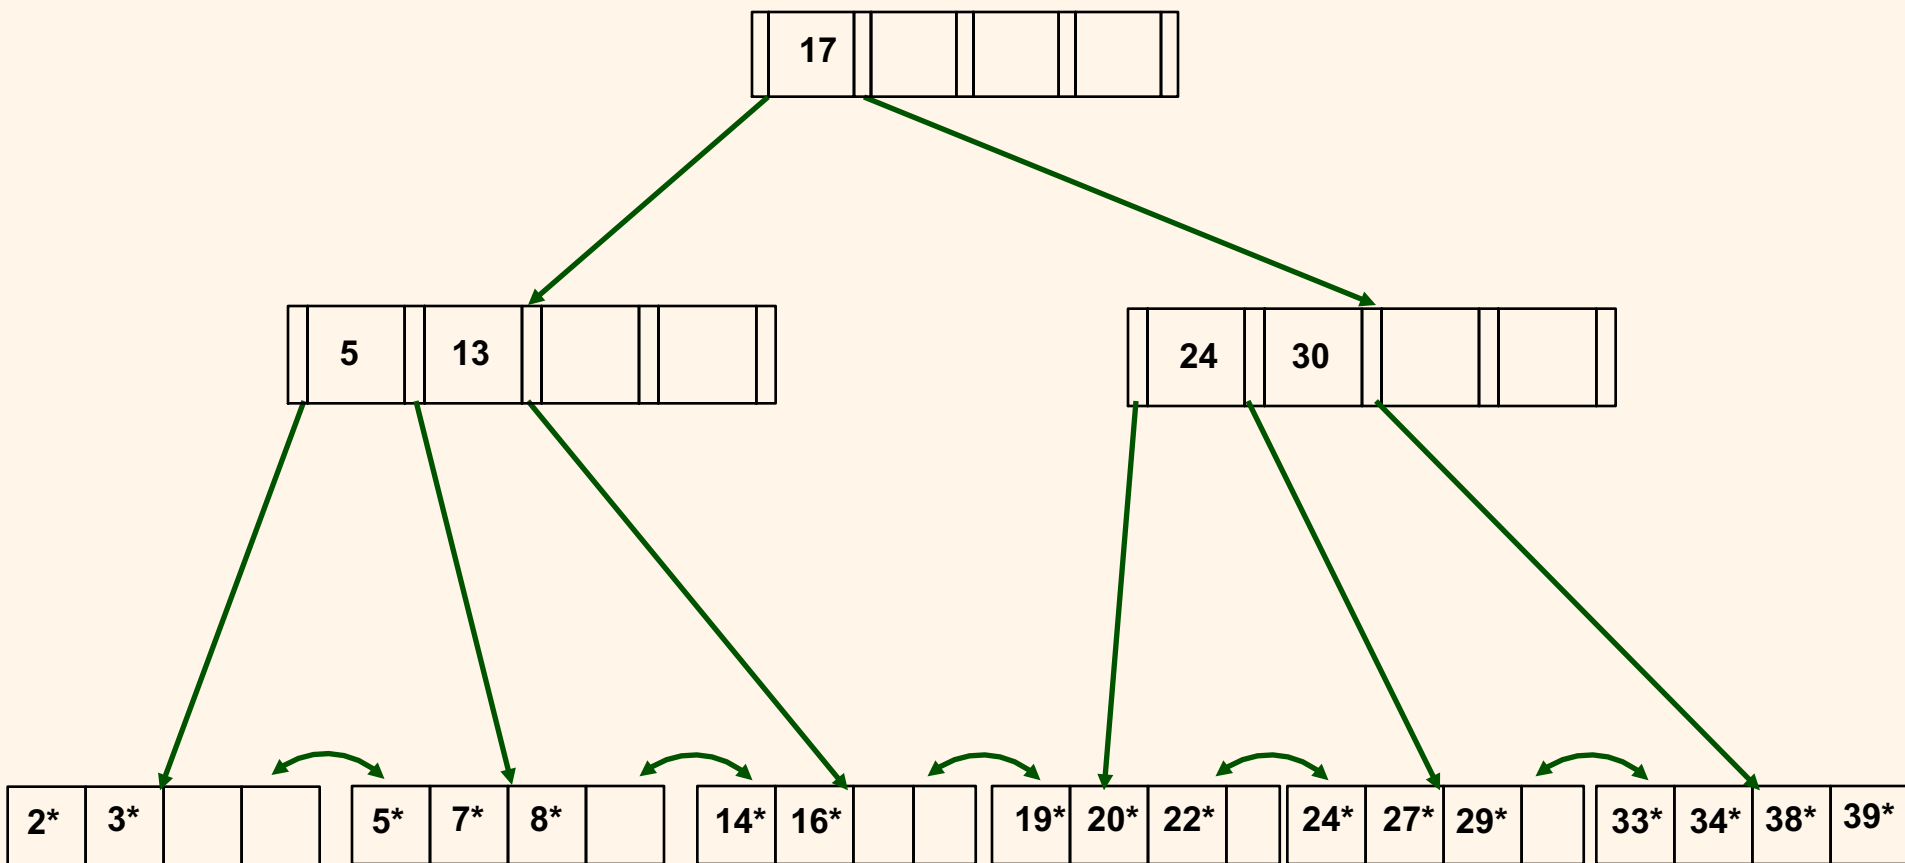
Locking in B+ Trees

- ❖ One solution: Ignore the tree structure, just lock pages while traversing the tree, following 2PL (objects to be locked = pages).
- ❖ Problem:
 - This has terrible performance!
 - Root node (and nodes close to root) become bottlenecks because every tree access begins at root.

Useful Observations

- ❖ Higher levels of the tree only direct searches for leaf pages.
 - Searches never go back to the root
 - So could release lock on root (and other nodes close to the root) early
 - OK even if we violate 2PL

Locking in B+ trees



Two Useful Observations

- ❖ For inserts, a node on a path from root to modified leaf must be locked (in X mode, of course), only if a split can propagate up to it from the modified leaf. (Similar point holds w.r.t. deletes.)
- ❖ So could also unlock some nodes early

A Simple Tree Locking Algorithm


- ❖ **Search:** Start at root and go down;
repeatedly, S lock child then unlock parent.

Example

ROOT

Consider:

- 1) Search 38*
- 2) Insert 45*



A Simple Tree Locking Algorithm

- ❖ **Insert/Delete:** Start at root and go down, obtaining X locks as needed. Once child is locked, check if it is safe:
 - If child is safe, release all locks on ancestors.
- ❖ **Safe node:** Node such that changes will not propagate up beyond this node.
 - Inserts: Node is not full.
 - Deletes: Node is not half-empty.

Improvements

- ❖ *Insert/Delete* sets a lot of exclusive locks which may not be necessary especially at the top levels
- ❖ Refinements based on requesting S locks until get to the leaf, then locking leaf only in X mode
 - Gambles that leaf won't need to split (common case)
 - If we are unlucky and leaf does split, need to handle it somehow
 - E.g. release all locks and restart using previous algorithm
 - Or try to request lock *upgrades* from S to X

Multiple-Granularity Locks

- ❖ Suppose we have a hierarchy on data "containers"
- ❖ Can lock at different levels

