

2PC (conclusion)



Last time: Two-Phase Commit

Coordinator

Send prepare

Wait for all responses

Decide **abort** or **commit**

Send abort or commit

Wait for all ACKs

Subordinate

Make local decision

Send yes or no


Perform **abort** or **commit**

Send ACK



State at coordinator

- ❖ Coordinator keeps some state (in memory) while running 2PC in a **transaction table**
- ❖ For each transaction
 - who are the subordinates
 - where we are in the protocol (which messages coordinator has sent/received)
- ❖ Ack messages from subordinates allow coordinator to garbage collect this state



2PC and failure

- ❖ Now let us talk about communication and site failures



2PC and comm failures

- ❖ If subordinate loses contact with a coordinator before receiving prepare?
 - Subordinate can decide to abort, since it hasn't voted so transaction cannot have committed




2PC and comm failures

- ❖ If coordinator loses contact with a subordinate before receiving all yes/no votes?
 - Coordinator can decide to abort since no-one has committed yet; must notify all subordinates who voted yes that we are aborting




2PC and comm failures

- ❖ If subordinate loses contact with a coordinator after voting but before receiving a commit/abort?
 - If voted no, can abort and is done
 - If voted yes, **can't unilaterally decide what to do**
 - Comm was lost during the **uncertainty period**




2PC and comm failures

- ❖ If subordinate loses contact with a coordinator after voting but before receiving a commit/abort?
 - Needs to communicate either with the coordinator, or possibly with another site to find out what the outcome was
 - Helps if coordinator tells each subordinate who the other subordinates are



2PC and comm failures

- ❖ If coordinator loses contact with subordinate after sending decision but before receiving ack?
 - Cannot garbage-collect transaction from coordinator state
 - When communication reestablished, can verify that subordinate knows about decision and then garbage-collect



2PC and site failures

- ❖ Now suppose the network is fine, but either the coordinator or a subordinate fails
- ❖ Need to remember sufficient state to allow recovery
- ❖ Will use a log for this
 - Coordinator and subordinate both log crucial steps in 2PC (and force log to disk so it survives a crash)
- ❖ A couple of variants for what is logged
 - Here we follow your course textbook on this
 - You may see minor differences elsewhere

Two-Phase Commit

Coordinator

Send prepare

Wait for all responses

Force-write abort or commit

Send abort or commit

Wait for all ACKs

Write (not force-write) end record

Subordinate

Make local decision

Force-write prepare or abort
Send yes or no

Force-write abort or commit
Send ACK



Restart after a failure

- ❖ Node crashes, comes back up
- ❖ Examines all in-progress 2PC transactions
 - Could be coordinator for some, subordinate for others
- ❖ Course of action based on last log record
- ❖ **Desired behavior:** if coordinator wrote commit to log, transaction is considered committed, else should abort



Restart after a failure

- ❖ Determine whether node was coordinator or subordinate
- ❖ Carry out recovery accordingly



Coordinator restart after a failure

- ❖ If have **end** log record, nothing to do
- ❖ If have **commit** or **abort** log record (but no end log record)
 - put transaction back into in-memory transaction table
 - know what the decision was; notify subordinates
 - wait for acks, clean up state and write end log record



Coordinator restart after a failure

- ❖ If don't have any log records
 - Can't have broadcast decision to subordinates
 - Decide to abort
 - If subordinates contact you asking for decision, can tell them it was abort
 - Could enter transaction back into transaction table, but no need
 - ◆ Default behavior: if you don't know anything about the transaction and a subordinate asks, tell them it was aborted.




Subordinate restart after a failure

- ❖ If have no (2PC-related) log entries, abort unilaterally
 - global decision can't have been a commit
- ❖ If have a **commit** or **abort** record, proceed accordingly



Subordinate restart after a failure

- ❖ If have **prepare** record but nothing else, cannot decide unilaterally
 - Site crashed in its **uncertainty period**
 - Needs to contact coordinator or other subordinates for what to do



2PC Optimizations

- ❖ Possible to optimize by reducing the number of messages and forced log entries in certain cases
- ❖ 2 optimizations:
 - Presumed Abort
 - Presumed Commit
- ❖ The XA standard for Distributed Transactions is 2PC with Presumed Abort

Two-Phase Commit

Coordinator

Send prepare

Wait for all responses

Force-write abort or commit

Send abort or commit

Wait for all ACKs

Write (not force-write) end record

Subordinate

Make local decision

Force-write prepare or abort
Send yes or no

Force-write abort or commit
Send ACK




Coordinator restart after a failure

- ❖ If have **end** log record, there is nothing to do
- ❖ If have **commit** or **abort** log record (but no end log record)
 - know what the decision was, proceed accordingly
- ❖ If don't have a commit or abort log record
 - decide to abort



2PC Presumed Abort

- ❖ If coordinator has no log records of transaction, it decides to **abort**
- ❖ So, if we decide to abort in a non-failure setting, can optimize by simply forgetting transaction (remove from transaction table)
 - **No need to force-write abort log record** at coordinator and subordinates
 - **No need for acks from subordinates after abort**
- ❖ But for commit, we proceed as in normal 2PC



2PC/PA, commit case

Coordinator

Send prepare

Wait for all responses

Force-write **commit**

Send commit

Wait for all ACKs

Write (not force-write) end record

Subordinate

Make local decision

Force-write **prepare**

Send yes

Force-write **commit**

Send ACK

2PC/PA, abort case (subordinate voted no)

Coordinator

Send prepare

Wait for all responses

Write **abort**

Send abort and forget transaction

Subordinate

Make local decision

Write **abort**

Send no


Write **abort**

No need to send ACK




2PC Presumed Abort

- ❖ If subordinate crashes and queries coordinator on what to do, and coord. has garbage-collected transaction, will reply abort
- ❖ Of course, coordinator cannot garbage collect a **committed** transaction until it has received acks from all subordinates



Presumed Commit: Motivation

- ❖ Commit is the more common case! Let's optimize for it, not for abort
 - Require ack for ABORT not COMMIT
 - Subordinates force ABORT records, not COMMIT records
 - No information in transaction table: presume commit!
- ❖ We can do this, but the coordinator *must* force some extra records for correctness



2PC/PC, abort case

Coordinator

Send prepare

Wait for all responses

Force-write **abort**

Send abort

Wait for all ACKs

Write (not force-write) end record

Subordinate


Make local decision

Force-write **abort**

Send no

Force-write **abort**

Send ACK



2PC/PC, commit case, first try

Coordinator

Send prepare

Wait for all responses

Write **commit**

Send commit and forget state

Subordinate

Make local decision

Write **prepare**

Send yes or no


Write **commit**

No need to send ACK




Presumed Commit

- ❖ Suppose coordinator crashes and comes back up; needs to figure out what to do
- ❖ If there are no log records, does it mean that it decided to commit before crash?
 - Or does it mean that it only sent PREPARES and didn't decide to commit yet?
 - Need to be able to distinguish between the two because actions to be taken are different!!



The solution

- ❖ Coordinator **force-writes begin/prepare record** (upon start of protocol) **AND commit record** (upon decision to commit)
- ❖ Subordinates do not need to force commit log records
- ❖ Now after crash recovery, either:
 - ◆ Coord has begin but no commit -> rollback
 - ◆ Coord has begin and commit -> commit



2PC/PC, commit case

Coordinator

Send prepare

Force-write **begin**

Wait for all responses

Force-write **commit**

Send commit and forget state

Subordinate


Make local decision

Write **prepare**

Send yes or no

Write **commit**


No need to send ACK






Additional 2PC Optimization

- ❖ Subordinates who only read send READ votes instead of YES votes
 - No log writes!
- ❖ Coordinator logic
 - READ & YES = YES
 - READ & NO = NO
 - READ & READ = READ
- ❖ If READ at coordinator, no need for second phase! Else, only contact non-READs.



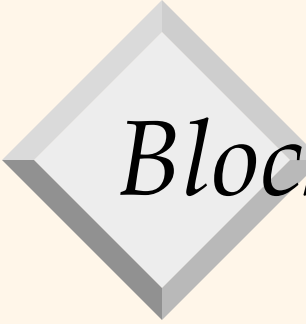
2PC *summary*

- ❖ Basic version
- ❖ Handling comm and site failures
 - A subordinate cannot always unilaterally recover
 - If failure occurred during its **uncertainty period**
- ❖ Optimizations to reduce messages, logging
 - Presumed Abort/Commit
 - Special treatment of readers



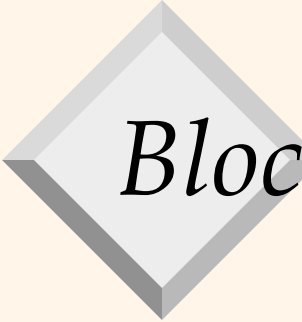
2PC *and blocking*

- ❖ 2PC is a protocol that may block even when a portion of the nodes are up (non-total failure)
- ❖ Blocks if a subordinate is in its uncertainty period and can only contact other subordinates who are in their uncertainty period
 - Could block indefinitely until they can finally reach someone who knows what to do
 - ◆ Coordinator or a subordinate not in its uncertainty period



Blocking protocols

- ❖ There are theoretical limitations on our ability to avoid blocking in a commit protocol
 - (while still retaining correctness)
- ❖ For an in-depth discussion see Phil Bernstein's textbook, Chapter 7



Blocking protocols

- ❖ There are protocols which reduce the probability of blocking
- ❖ Example: 3PC (Three-phase commit)
- ❖ If no comm failures, 3PC will not block as long as a majority of sites are operational
 - Think about why 2PC does not guarantee that!



Three Phase Commit

- ❖ Phase 1: Voting as before
- ❖ Phase 2: Dissemination of results
 - If coordinator gets all "yes" votes, sends "precommit" message
 - When coordinator gets acks from a majority of the sites, actually makes decision to commit
- ❖ Phase 3: Termination as before



- ❖ Reduces chance of blocking
- ❖ Phase 2 makes sure that the **decision to commit is recorded** on a majority of sites before the final order to commit is issued



- ❖ Recovery: no comm failures, majority of sites are up
- ❖ If no-one has a "precommit" message, coordinator cannot have issued final order to commit
 - safe to abort
- ❖ If someone has "precommit" message, knows decision was going to be commit
 - So safe to commit!