



"NewSQL:" C-Store, Spanner



Where We've Been ...

- ❖ an interesting recent system
- ❖ C-Store (commercial fork: Vertica)
 - Store tables by columns rather than rows to optimize performance for OLAP queries



Where We're Going ...

- ❖ Two more interesting recent systems
- ❖ H-Store (commercial fork: VoltDB)
 - Custom architecture tuned for OLTP workloads
 - Heavy use of main memory
- ❖ Spanner (an internal system used at Google)
 - SQL-like language over a key-value store
 - Strong consistency (Paxos!!)

H-Store/VoltDB



- ❖ What about OLTP workloads?
- ❖ We can improve performance on these too!
- ❖ Exploit modern hardware
- ❖ Exploit unique features of OLTP workloads
- ❖ Example: H-Store (commercialized as VoltDB)



Today's OLTP workloads

- ❖ Data is relatively small, < 100 GB typically
- ❖ Transactions are short-running, no "user stalls"
 - This is no longer the 1970s where you input SQL at a terminal
 - When you buy something on Amazon, typically split into several underlying transactions



Today's computers

- ❖ Memory is no longer tiny
 - 100GB of RAM? Sure, you can outfit a machine with that
- ❖ Your database no longer sits on a single box
 - Have available infrastructure that is distributed and replicated for fault-tolerance



H-Store design decisions

- ❖ Run everything **in memory**
 - Could rely on replication and failover for durability
 - ◆ So, only need to log for undo purposes (not for redo)
 - Though VoltDB does use periodic disk snapshots




H-Store design decisions

- ❖ Run all transactions **serially**
 - Hey, they're short anyway
- ❖ Result: saves on some major overhead
 - Disk accesses
 - Synchronization/concurrency



OLTP Workloads

- ❖ Make use of the fact that OLTP workloads are not ad-hoc
- ❖ Require all possible transaction classes to be predefined and registered with the system
 - Can be pre-optimized
 - For distributed transactions, can identify which of them really require inter-site communication/2PC
- ❖ Allows a better DB design as we know the entire workload up front (data partitioning etc.)



Summary so far

- ❖ Custom solutions and architectures for particular classes of applications
- ❖ Column stores for read-mostly, OLAP style workloads
- ❖ H-Store and similar systems for OLTP workloads
 - In-memory
 - Transactions run serially
 - Optimized for a fully pre-specified workload

Google Spanner

- Distributed multiversion database
 - General-purpose transactions (ACID)
 - SQL query language
 - Schematized tables
 - Semi-relational data model
- Running in production
 - Storage for Google's ad data
- Presented at OSDI (major systems conference) in 2012

Overview

- Supports lock-free distributed read transactions using snapshots
- Guarantees external consistency (linearizability) of distributed transactions
 - A combination of serializability and linearizability
 - If T1 commits before T2 starts, then T1 is serialized before T2
 - First system at global scale to enforce this

Read Transactions

- In a social network, generate a page of friends' recent posts
 - Consistent view of friend list and their posts

Why consistency matters

1. Remove untrustworthy person X as friend
2. Post P: “My government is repressive...”

Synchronizing snapshots

- Read snapshots must be consistent across multiple sites
- Implementation relies on appropriate use of transaction timestamps
 - In distributed commit!


Assigning Timestamps

- Strict two-phase locking for write transactions
- Assign timestamp while locks are held



Some Timestamp Guarantees

- For conflicting transactions, timestamp order == serialization order



- T_4 starts after T_3 ends $\Rightarrow T_4$ has larger timestamp





TrueTime API

- “Global wall-clock time” with bounded uncertainty



Timestamps and TrueTime



Distributed Timestamps w/2PC

- Coord sends **Prepare**
- Sub replies **Yes(\sim TT.now().latest)**
- Coord computes $s = \max(\{\text{replies}\}, \text{TT.now().latest})$
- Coord waits until $\text{TT.now.earliest} > s$
- Coord sends **Commit(s)**
- Coord waits for ACKs

Subtle Issues

- Message propagation delay
- Sub reply must be $>$ any previous proposal by that sub
- Reader with snapshot greater than some outstanding proposal must block

Summary

- Lock-free read transactions across datacenters
- External consistency
 - A very strong formal guarantee
- TrueTime
 - Uncertainty in time can be waited out
- More details (e.g. how to actually implement consistent reads at a time/version) in paper



Slide credits

- ❖ VLDB 2009 tutorial on Column stores
 - http://www.cs.yale.edu/homes/dna/talks/Column_Store_Tutorial_VLDB09.pdf
- ❖ H-Store slides
 - <http://hstore.cs.brown.edu/slides/hstore-vldb2007.pdf>
- ❖ Google Spanner Slides
 - <http://research.google.com/archive/spanner-osdi2012.pptx> (substantially modified)