

*Map Reduce*



# *Map Reduce*

- ❖ Not a data model, but a processing/programming model
- ❖ Focus: the abstraction/model and how to use it
- ❖ Original Google paper
  - <http://research.google.com/archive/mapreduce.html>
- ❖ A great textbook available on-line
  - <http://lintool.github.io/MapReduceAlgorithms/ed1n.html>



# *Map Reduce*

- ❖ Programming model introduced by Google (2004)
- ❖ Not really a new concept
  - Dates back to ideas from functional programming
- ❖ Very useful for processing large datasets
- ❖ Infrastructure
  - Hadoop, Amazon Elastic Map Reduce,...
- ❖ Also treated as a programming pattern
  - MongoDB and other systems support it



# *What is Map Reduce?*

- ❖ A programming model and infrastructure for parallel programming
- ❖ Doing very large-scale data processing requires parallelization



# *What is Map Reduce?*

- ❖ Not trivial to parallelize arbitrary task:
  - What are the subproblems?
  - How do we get the data to/from the workers working on each subproblem?
  - How do we synchronize and share results as needed between workers?

# *Example*



- ❖ You get a job at Google
- ❖ Your boss says: hey, we have a large corpus of documents
  - All the pages on the Web
- ❖ We need some statistics - a list of word frequency occurrences across the whole corpus
- ❖ We need the results fast, so don't do it all on one machine.
- ❖ What do you do?



# *Typical Workflow*

- ❖ Iterate over a large number of records
  - ❖ *Extract something of interest from each*
  - ❖ Bring together intermediate results
  - ❖ *Aggregate intermediate results*
  - ❖ Generate final output
- 
- ❖ Most of the real "computation" occurs in the two blue phases



# *Map Reduce*

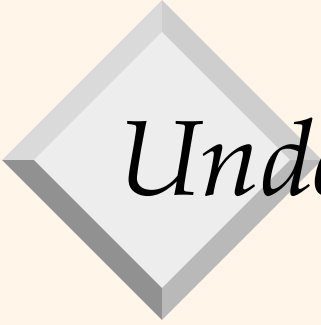
- ❖ (Or map-reduce, mapreduce, etc.)
- ❖ A general framework for writing parallel programs that follow the workflow we saw
- ❖ Idea:
  - You write the code for the two blue phases
    - ◆ Because that's what is unique to your computation
  - System takes care of the rest





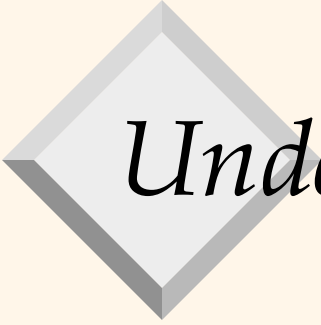
# *Typical Workflow*

- ❖ Iterate over a large number of records
- ❖ Map: Extract something of interest from each
- ❖ Bring together intermediate results
  - In some standardized way
- ❖ Reduce: Aggregate intermediate results
- ❖ Generate final output




# *Understanding Map Reduce*

- ❖ Key to understanding Map Reduce is the third point in workflow
  - How are intermediate results brought together?
  - The framework does it for you, so you have to understand how it does it



# *Understanding Map Reduce*

- ❖ Fundamental idea: key-value model
  - This is the data model for Map Reduce jobs
  - Helps provide a unified interface for bringing results together



# *MR and key-value model*

- ❖ Key-value is simple data model
- ❖ Map-Reduce uses that as input and output format
- ❖ Map function takes one key-value pair as input and outputs a set of key-value pairs
  - The input and output keys can be different



## *Example: Word Count*

```
map(String key, String value):  
    // key: document id  
    // value: document contents  
    for each word w in value:  
        EmitIntermediate(w, "1");
```

A mapper utility can apply this in parallel to a whole lot of documents



# *Now what?*


- ❖ Have a whole lot of key-value pairs after the map
- ❖ Need to bring them together and aggregate
- ❖ Would like to bucketize/ GROUP BY something so can compute in parallel again
  - What to bucketize by?
    - ◆ The English word (i.e. the output key of the mapper)



## *Example: Word Count*

```
reduce(String key, Iterator values):  
    // key: a word  
    // values: a list of counts  
    int result = 0;  
    for each v in values:  
        result += ParseInt(v);  
    Emit(key, AsString(result));
```

Can also run this in parallel



# *So what do we have so far?*

- ❖ See how to write map and reduce functions
  - Work on a key-value model
- ❖ Believe that both map and reduce functions can be executed in parallel
- ❖ But what about the middle step?
  - Bucketize output of mapper based on value of output key...
- ❖ Fortunately, the exact same middle step is useful for a lot of other problems
  - So map reduce frameworks have it built-in!
  - You don't need to write this logic yourself





# *Map Reduce*

❖ You only need to specify two functions:

**map**  $(k, v) \rightarrow [(k', v')]$

**reduce**  $(k', [v']) \rightarrow [(k'', v'')]$

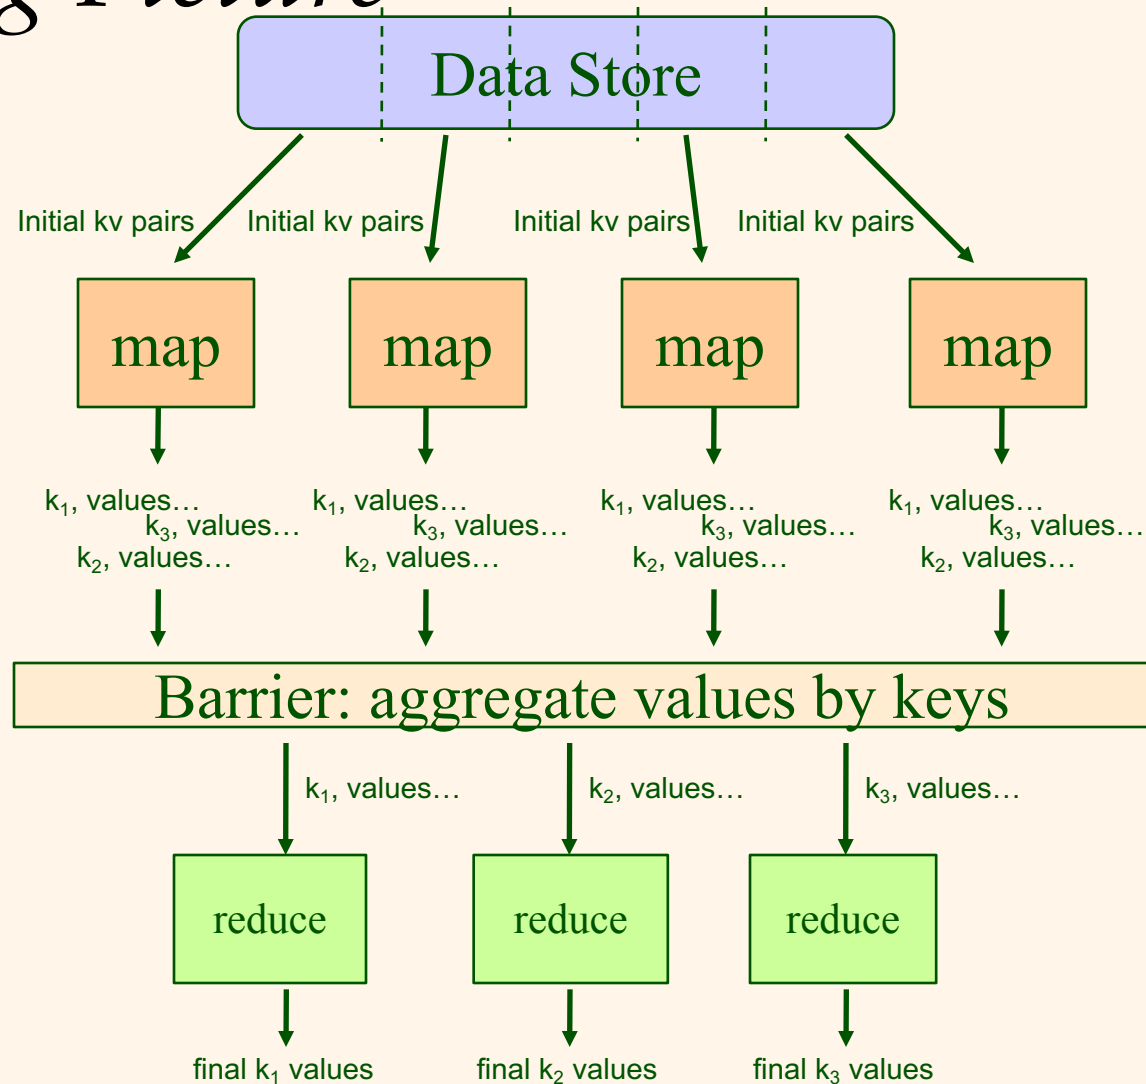
–  $[ \dots ]$  denotes a list



# *Map Reduce*

- ❖ Framework takes care of the actual execution:
  - Applies your map function to every initial (k,v)
  - Reshuffles the output of the map to group by k'
  - Applies your reduce function

# *The Big Picture*





# *Example: Word Count*

```
map(String key, String value):  
    // key: document id  
    // value: document contents  
    for each word w in value:  
        EmitIntermediate(w, "1");
```

```
reduce(String key, Iterator values):  
    // key: a word  
    // values: a list of counts  
    int result = 0;  
    for each v in values:  
        result += ParseInt(v);  
    Emit(key, AsString(result));
```




# *Map and Reduce Functions*

- ❖ Do not have to be purely functional
- ❖ Can keep internal state across multiple inputs
- ❖ Can also have external side effects
  - E.g. write to files
- ❖ Should be careful about using external resources
  - E.g. if have multiple mappers and/or reducers contending for same database, could become a bottleneck




# *Map and Reduce Functions*

- ❖ Possible to have programs without a reduce
  - Mappers just apply some computation in parallel to a dataset
- ❖ Impossible to have programs without a map
  - But map could be identity function
  - Use framework to re-sort and re-group key-value pairs before feeding to reducers
- ❖ Could have reducer identity function too
  - Computation occurs in map phase
  - Framework used to re-sort and re-group output of mapper
- ❖ Could even have both mapper and reducer as identity functions



# *Map Reduce Frameworks*


- ❖ Various frameworks to support this programming pattern
  - Google's own internal implementation, Hadoop, Amazon EMR
  - Also supported in MongoDB



# *Map Reduce Frameworks*

- ❖ Implementations vary a bit in what exact functionality they allow/support
  - E.g. whether reducer input and output keys must be the same
  - Or whether they support additional functions like partitioners and combiners
  - Or what happens in corner cases (e.g. MongoDB won't run reducer if there is only one value for a reducer input key)





# *Programming in Map Reduce*

- ❖ Powerful abstraction, but requires different way of thinking about programming
- ❖ Things to figure out:
  - How to impose key-value structure on problem?
  - What can be parallelized?
  - How to deal with the fact that there is no "global state" anymore (or at least that you should avoid using global state?)
- ❖ Next: a few examples to get you more comfortable with doing things in MR



# *Inverted Index*


- ❖ How to use Map Reduce to build an inverted index?
- ❖ Given a set of documents as input
- ❖ Want as output a set of entries of the form:  
*<word, list of documents it appears in>*
- ❖ May assume documents have unique ids
- ❖ Very useful in practice e.g. in search engines



# *Inverted Index*

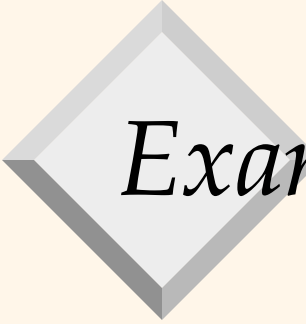
```
map(String docid, String contents):  
    for each word w in contents:  
        EmitIntermediate(w, docid);
```

```
reduce(String word, Iterator docids):  
    List result = new ArrayList();  
    for each d in docids:  
        result.add(d);  
    Emit(word, result);
```



# *Inverted Index*

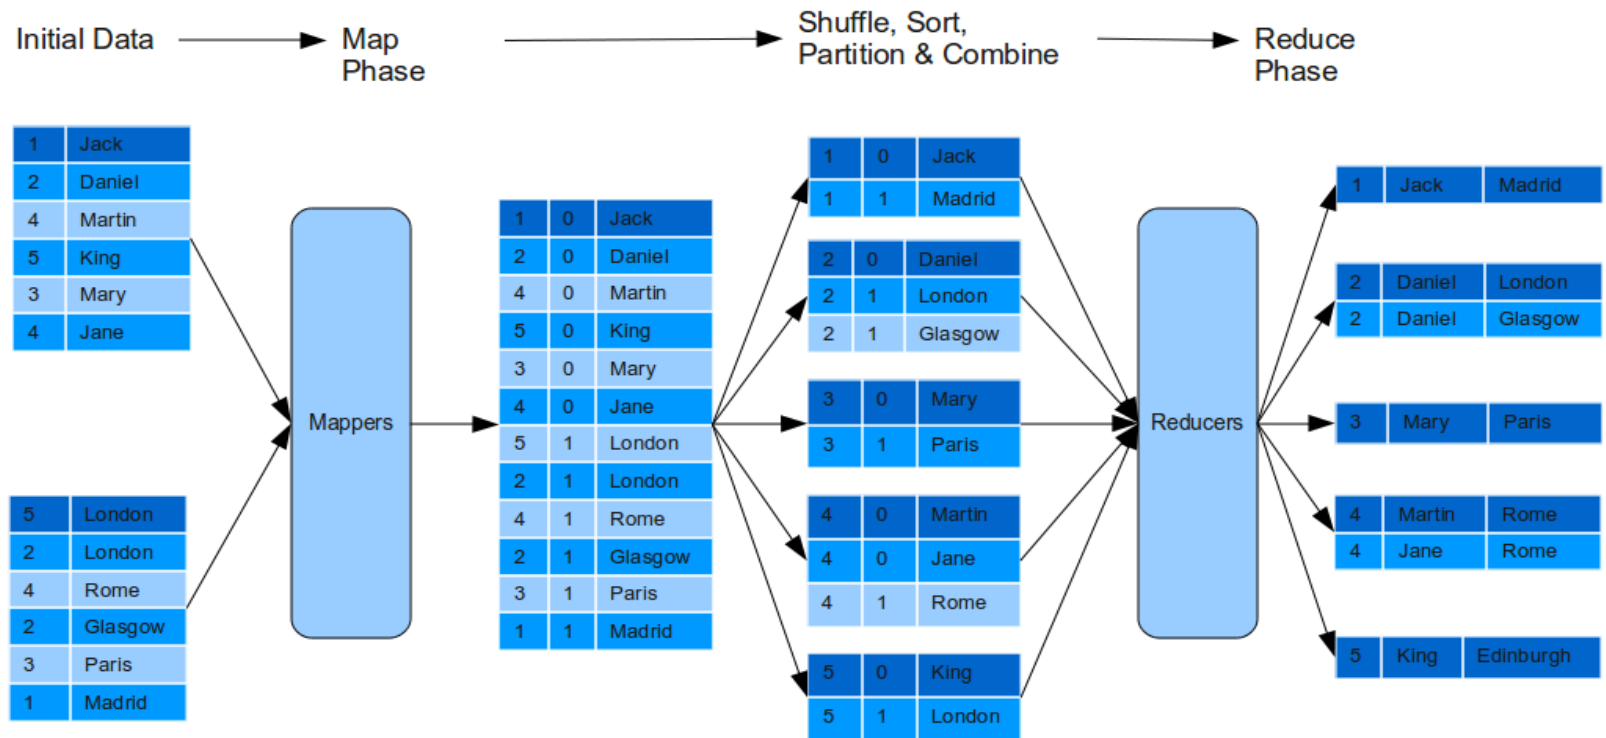
- ❖ Easy to modify this to make it more fancy
  - Keep track of word positions within documents
  - Etc...




## *Example 2: Joins*

- ❖ Have two large datasets representing two large relations
  - In some reasonable format, e.g. with one line per row
- ❖ Need to perform a join on a particular column
- ❖ How to do it with Map Reduce??

# Example: Join






# *Iterative Map Reduce*

- ❖ You are not limited to a single MR run
- ❖ Output of one MR job can serve as input to the next one
  - MR pipelines
  - Need to make sure your formats are compatible!
- ❖ Iterative algorithms also possible and useful
- ❖ First case study: graph processing

# Representing Graphs



	$n_1$	$n_2$	$n_3$	$n_4$	$n_5$
$n_1$	0	1	0	1	0
$n_2$	0	0	1	0	1
$n_3$	0	0	0	1	0
$n_4$	0	0	0	0	1
$n_5$	1	1	1	0	0

adjacency matrix

$n_1$  [  $n_2$ ,  $n_4$  ]  
 $n_2$  [  $n_3$ ,  $n_5$  ]  
 $n_3$  [  $n_4$  ]  
 $n_4$  [  $n_5$  ]  
 $n_5$  [  $n_1$ ,  $n_2$ ,  $n_3$  ]

adjacency lists






# *Single Source Shortest Path*

- ❖ **Problem:** find shortest path from a source node to one or more target nodes
- ❖ Single processor machine: Dijkstra's Algorithm
- ❖ MapReduce: parallel Breadth-First Search (BFS)


# SSSP




# Dijkstra's Algorithm Review

```
1: DIJKSTRA( $G, w, s$ )
2:    $d[s] \leftarrow 0$ 
3:   for all vertex  $v \in V$  do
4:      $d[v] \leftarrow \infty$ 
5:    $Q \leftarrow \{V\}$ 
6:   while  $Q \neq \emptyset$  do
7:      $u \leftarrow \text{EXTRACTMIN}(Q)$ 
8:     for all vertex  $v \in u.\text{ADJACENCYLIST}$  do
9:       if  $d[v] > d[u] + w(u, v)$  then
10:         $d[v] \leftarrow d[u] + w(u, v)$ 
```


# Dijkstra's Algorithm Example




# Dijkstra's Algorithm Example




# *Dijkstra's Algorithm Example*



# Dijkstra's Algorithm Example




# *Dijkstra's Algorithm Example*





# *Dijkstra's Algorithm Example*






# *What about Map Reduce?*

- ❖ Suppose we have a very, very, very big graph
- ❖ And would like to compute this information quickly and in parallel
  - Using the magic of Map Reduce
- ❖ Definitely can't run Dijkstra's algorithm directly
  - Can't have a global queue!



# *SSSP in Map Reduce*


- ❖ Can't run Dijkstra's algorithm directly
  - Can't have a global queue!
- ❖ Another way to do it: Parallel BFS
- ❖ Start by assuming all edge weights are equal
  - Will relax this later




# *Finding the Shortest Path*

- ❖ Intuition: process all nodes at each step
- ❖ Some nodes have no information (distance = infinity)
  - So can't do much
- ❖ But other nodes do know something
  - E.g. source knows it is at distance 0
  - So can pass this fact on to its out-neighbors
    - ◆ Who now know they are at distance 1!
  - At the next iteration, these neighbors know they're at distance 1
    - ◆ So can tell *their* out-neighbors they're at distance 2.


# *Parallel BFS*





# *From Intuition to Algorithm*

- ❖ A map task receives
  - Key: node  $n$
  - Value:  $D$  (distance from start), points-to (list of nodes reachable from  $n$ )
- ❖  $\forall p \in \text{points-to: emit } (p, D+1)$
- ❖ The reduce task gathers possible distances to a given  $p$  and selects the minimum one
- ❖ Possible through the magic of the "sort and shuffle" between Map and Reduce
  - Map processes node and updates distances of **out-neighbors**
  - Reduce processes node based on info from its **in-neighbors**



# *Multiple Iterations Needed*

- ❖ Each Map Reduce task advances the “known frontier” by one hop
  - Subsequent iterations include more reachable nodes as frontier advances
  - Multiple iterations are needed to explore entire graph
  - Feed output back into the same MapReduce task



# *Multiple Iterations Needed*

- ❖ Passing along the graph structure:
  - Next iteration of Map needs points-to list again
  - So need to "carry" it with us as we run the algorithm





```
class MAPPER
  method MAP(nid  $n$ , node  $N$ )
     $d \leftarrow N.DISTANCE$ 
    EMIT(nid  $n$ ,  $N$ )
    for all nodeid  $m \in N.ADJACENCYLIST$  do
      EMIT(nid  $m$ ,  $d + 1$ )
```

# Reduce

```
class REDUCER
  method REDUCE(nid  $m$ , [ $d_1, d_2, \dots$ ])
     $d_{min} \leftarrow \infty$ 
     $M \leftarrow \emptyset$ 
    for all  $d \in \text{counts } [d_1, d_2, \dots]$  do
      if IsNode( $d$ ) then
         $M \leftarrow d$ 
      else if  $d < d_{min}$  then
         $d_{min} \leftarrow d$ 
     $M.\text{DISTANCE} \leftarrow d_{min}$ 
    EMIT(nid  $m$ , node  $M$ )
```



# *Termination*


- ❖ Eventually, all nodes will be discovered, all edges will be considered (in a connected graph)
- ❖ Stop when there are no nodes with a distance of infinity
  - Can be checked by the driver/harness/program that runs the outer loop and schedules each Map Reduce job




# *Weighted Edges*

- ❖ Now add positive weights to the edges
- ❖ Simple change: points-to list in map task includes a weight  $w$  for each pointed-to node
  - emit  $(p, D+w_p)$  instead of  $(p, D+1)$  for each node  $p$
- ❖ Termination behavior different
  - Just because we've reached a node doesn't mean we've found the shortest path to it!

# Node Exploration Process




# Node Exploration Process





# *Termination*

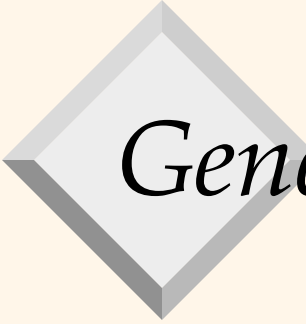
- ❖ When distances have not changed during an iteration, safe to stop



# *Comparison to Dijkstra*

- ❖ Dijkstra's algorithm is more efficient
  - At any step it only pursues edges from the minimum-cost path inside the frontier
  - Only processes each node once
- ❖ MapReduce explores all paths in parallel
  - Does a lot of recomputation
    - ◆ Not a bug, need it to handle situations where the "shortest" path contains more edges than another available path
  - But can be done in parallel






# *General Approach*

- ❖ Graph algorithms with MapReduce:
  - Each map task receives a node and its outlinks
  - Map task compute some function of the link structure, emits value with target as the key
  - Reduce task collects keys (target nodes) and aggregates
- ❖ Iterate multiple MapReduce cycles until some termination condition



# *PageRank*

- ❖ Google's famous algorithm for ranking Web Pages
  - A measure of "quality reputation" of a page
  - Useful for ranking/ordering search results



# *Random Surfer Model*

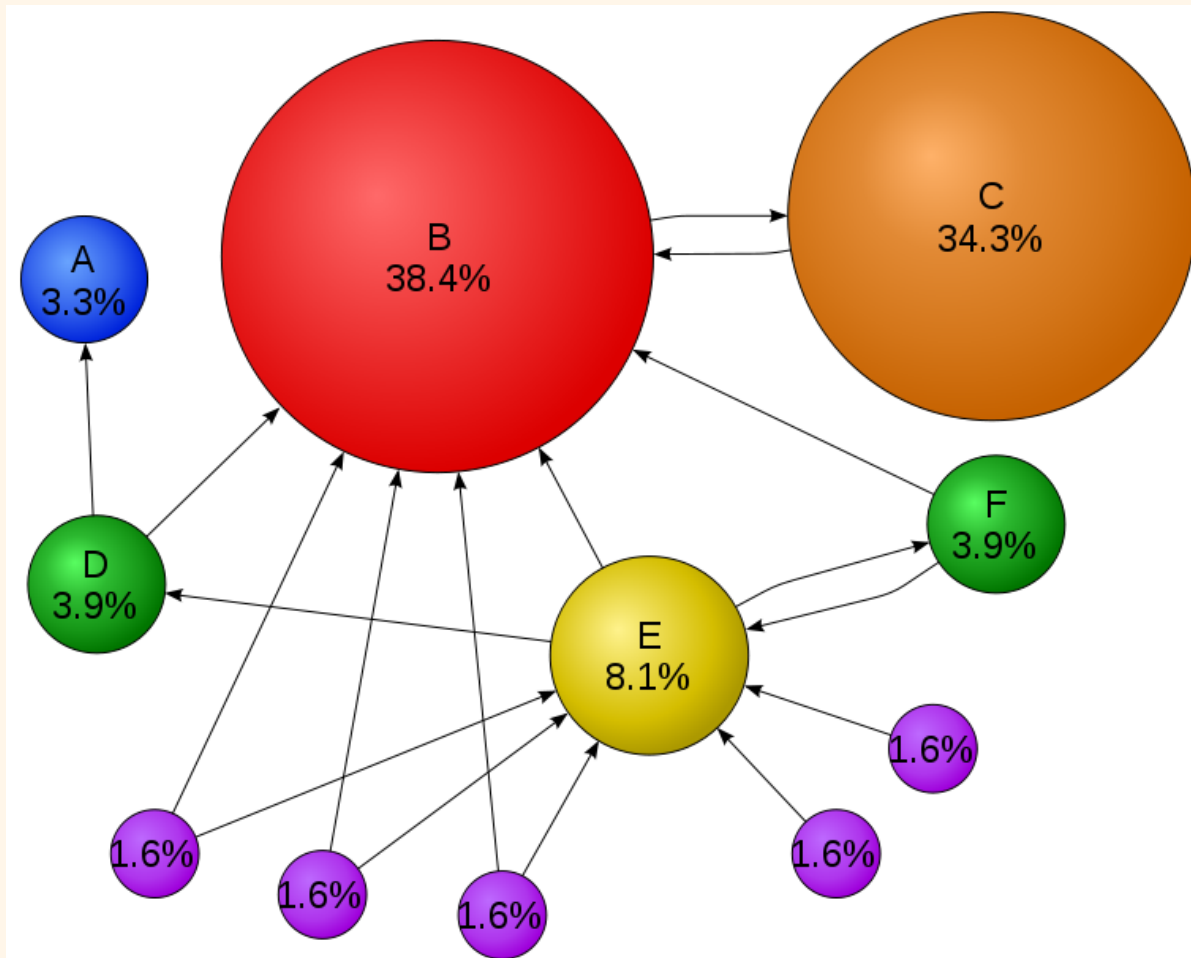
- ❖ Intuition for PageRank
- ❖ Imagine a surfer who starts on a randomly chosen page and then follows outgoing links at random
  - Markov process
- ❖ PageRank is probability that user will arrive at a given page during this random walk



## *A little more complex!*

- ❖ Model assumes that surfer doesn't always follow a link, but sometimes e.g. bookmarks instead.
- ❖ Before each move, surfer flips a coin
  - With probability  $1 - \alpha$ , follows an out-link
  - With probability  $\alpha$ , teleports to a (uniformly chosen) random page

# PageRank Example






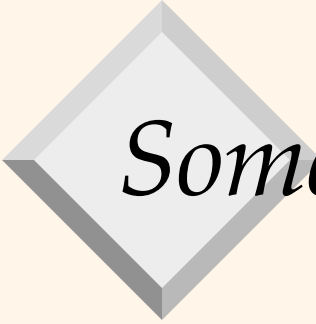
# *Computing PageRank*

- ❖ Let's start with some simplifying assumptions
  - Assume all nodes have at least one outlink
  - And surfer never does a random restart using bookmarks
- ❖ Will talk about how to lift these assumptions soon

# *Simplified PageRank Intuition*

- ❖ PageRank of a page is based on the PageRank of the pages which link to it
- ❖ A page divides its PageRank equally among all its outgoing links






*Somewhat more formally*

$$P(n) = \sum_{m \in L(n)} \frac{P(m)}{C(m)}$$

- ❖  $L(n)$  is the set of pages that link to  $n$  and  $C(m)$  is the number of out-neighbors of page  $m$







# *Iterative computation of SPR*

- ❖ How do we compute this?
- ❖ Various methods, but we are interested in Map Reduce
- ❖ Idea:
  - Initialize everything to the same PageRank ( $1/\text{number of nodes}$ )
  - "pass around" PageRank contributions from nodes to their out-neighbors



# Example

## Iteration 1



# Example

## Iteration 2





```
class MAPPER
```

```
  method MAP(nid  $n$ , node  $N$ )
```

```
     $p \leftarrow N.\text{PAGERANK} / |N.\text{ADJACENCYLIST}|$ 
```

```
    EMIT(nid  $n$ ,  $N$ ) ▷ Pass along graph structure
```

```
    for all nodeid  $m \in N.\text{ADJACENCYLIST}$  do
```

```
      EMIT(nid  $m$ ,  $p$ ) ▷ Pass PageRank mass to neighbors
```

# Reduce

```
class REDUCER
  method REDUCE(nid  $m$ ,  $[p_1, p_2, \dots]$ )
     $M \leftarrow \emptyset$ 
    for all  $p \in \text{counts } [p_1, p_2, \dots]$  do
      if ISNODE( $p$ ) then
         $M \leftarrow p$  ▷ Recover graph structure
      else
         $s \leftarrow s + p$  ▷ Sum incoming PageRank contributions
     $M.\text{PAGERANK} \leftarrow s$ 
    EMIT(nid  $m$ , node  $M$ )
```



# *Iterate*

- ❖ Full algorithm is iterative
- ❖ Initialize the nodes to uniform distribution
- ❖ Run the two MR jobs described iteratively
- ❖ Until convergence (no change)



## *Now, back to dangling nodes*

- ❖ If any PageRank is lost due to nodes with no out-neighbors, redistribute that PageRank uniformly throughout the graph for next iteration
- ❖ In the Map Reduce model: keep track of any lost PageRank
  - E.g. by using a special reserved intermediate key, or using a counter (i.e. storing it somewhere)



# *Solution*

- ❖ After the MR task is done, do a cleanup pass
- ❖ Deal both with "missing mass" and with random restart factor





## *Cleanup pass*

- ❖ Adjust the PageRank of each node to be

$$p' = \alpha \left( \frac{1}{|G|} \right) + (1 - \alpha) \left( \frac{m}{|G|} + p \right)$$

- ❖ Where  $p$  is the current PageRank,  $m$  is the mass lost due to sinks, and  $|G|$  is the number of nodes in the graph
- ❖ This can be done using a map job (no reduce)



# *The full PageRank Algorithm*

- ❖ Initialize the nodes to uniform distribution
- ❖ Run the two MR jobs described iteratively
- ❖ Until convergence (no change)