

# *Final Exam Discussion*



# *Logistics*

- ❖ Mon, Dec 12th, 7-9pm
- ❖ Take-home exam in CMS
  - Available shortly before 7PM
  - Scanned submissions accepted
  - Late submissions accepted
    - ◆ for a few hours, appropriate late penalty
- ❖ There will be office hours next week
  - Consult schedule in CMS
- ❖ See me ASAP if you have a conflict!



# *Exam Details*

- ❖ Cumulative
- ❖ Emphasis on material not in the prelim
- ❖ Time is very limited
  - so some material will end up not being covered
  - many questions will be short-answer




# *Exam Details*

- ❖ BUT I may still ask things like:
  - Write a query in SQL or Relational Algebra
  - Solve an optimization/cost calculation question
  - Questions on indexing
  - ... and other stuff from early in the course



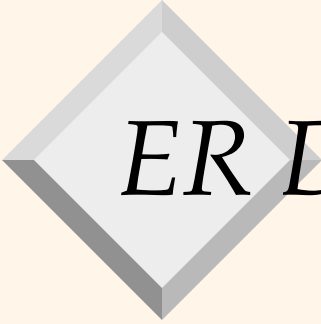
# *Recommendation*

- ❖ Go over all the homeworks!!
  - And the prelim
- ❖ You should be familiar with **everything that was in any homework**, even if you chose to split up work so that your partners did a particular question...



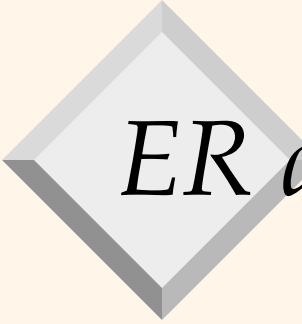
## *New topics for Final (not in Prelim)*

- ❖ Functional Dependencies/ER Diagrams
- ❖ Concurrency and Recovery
- ❖ Distributed Systems Topics
  - 2PC, Paxos, Eventual Consistency, ...
- ❖ Column Stores
- ❖ MapReduce



# *ER Diagrams*


- ❖ Be familiar with the ER diagrams introduced in class and in the textbook
  - Including extended features like weak entities
- ❖ Understand how to translate/map between ER diagrams and relational tables



# *ER diagrams*


- ❖ Typical questions you can expect:
  - Draw an ER diagram for this real-world scenario
  - Convert this ER diagram to a SQL schema or vice versa





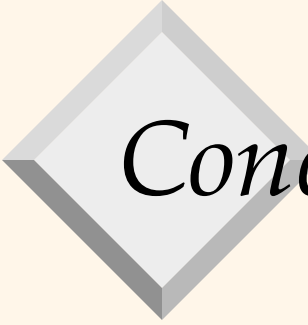
# *Functional dependencies*

- ❖ Understand why redundancy is bad and how FDs relate to redundancy
- ❖ Understand FDs
  - Definition
  - Relationship to keys
  - Armstrong's Axioms, soundness, completeness etc
  - What a closure, attribute closure and cover are
  - Be able to do proofs on the above



# *Decompositions and Normal Forms*

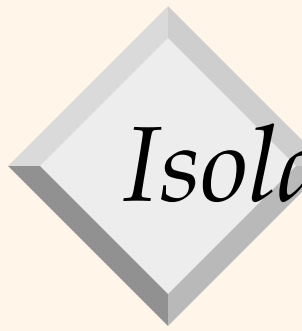
- ❖ Understand problems with decomposition
  - Performance hit
  - May be lossy
  - May not be dependency preserving
- ❖ Know definitions of BCNF and 3NF, how to check if a relation is in one of these, and how to decompose



# *Concurrency and Transactions*

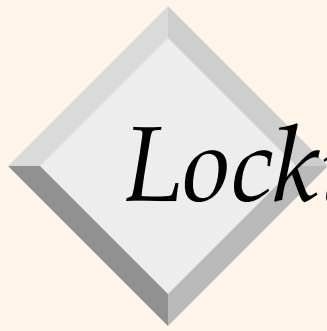
## ❖ ACID properties

- Atomicity
- Durability
- Consistency
- Isolation



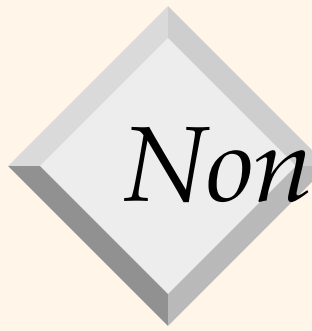
# *Isolation/concurrency*

- ❖ Formal/theoretical concepts:
  - Isolation anomalies
  - Various kinds of serializability
  - Recoverable/ACA/strict schedules



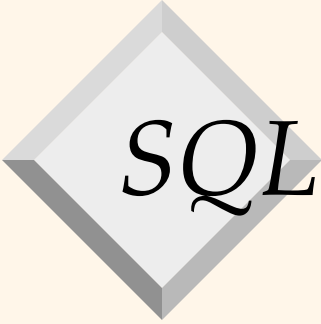
# *Locking Protocols*

- ❖ 2PL and its variants (conservative, strict, both)
  - How they work
  - Pros and cons
- ❖ Advanced/custom locking protocols
  - Phantoms and index locking
  - Locking in B+-trees
  - Multiple granularity locks (IS, IX, S, X)
- ❖ Deadlock and solutions



# *Nonlocking protocols*

- ❖ Optimistic
- ❖ Version-based
  - Timestamp-based multiversion CC
  - Snapshot isolation



# *SQL isolation levels*

Level	Dirty Read	Unrepeatable Read	Phantom
READ UNCOMMITTED	Possible	Possible	Possible
READ COMMITTED	No	Possible	Possible
REPEATABLE READ	No	No	Possible
SERIALIZABLE	No	No	No

# Recovery and ARIES



## LogRecords

prevLSN  
XID  
type  
pageID  
length  
offset  
before-image  
after-image



## Data pages

each  
with a  
pageLSN

## master record



## Xact Table

lastLSN  
status

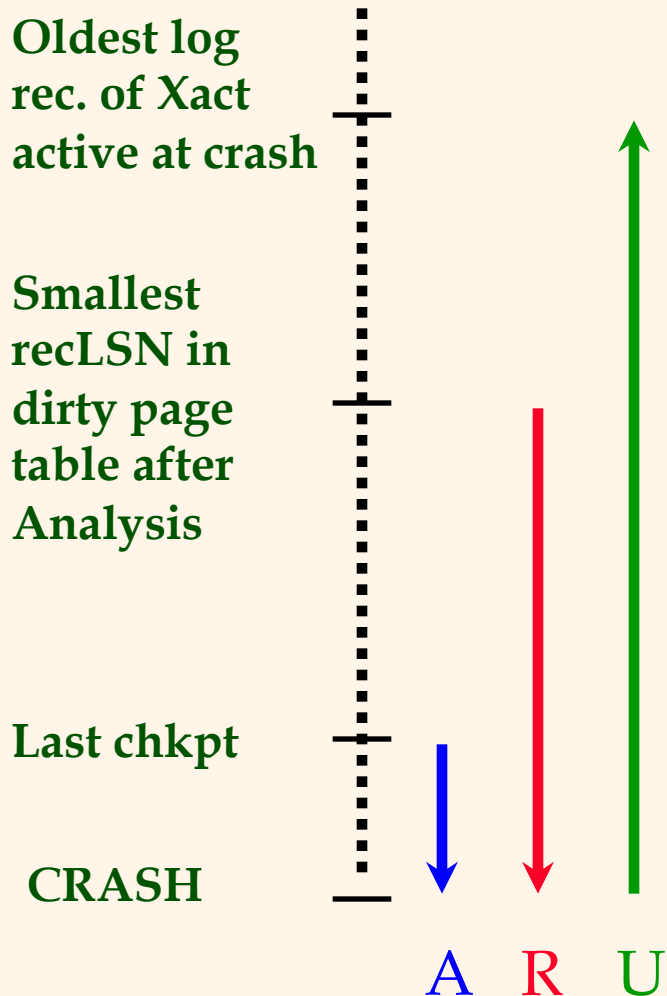
## Dirty Page Table

recLSN


## flushedLSN



# Crash Recovery: Big Picture



- ❖ Start from a **checkpoint** (found via **master** record).
- ❖ Three phases. Need to:
  - Figure out which Xacts committed since checkpoint, which failed (**Analysis**).
  - **REDO** *all* actions.
    - ◆ (repeat history)
  - **UNDO** effects of failed Xacts.



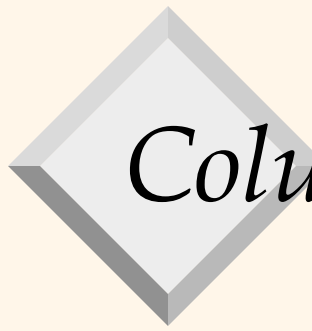
# *Distributed systems*

- ❖ Query processing
  - New join algorithms (semijoin, Bloomjoin)
- ❖ Consensus protocols
  - 2PC and variants
  - 3PC
  - Paxos




# *Distributed systems*

- ❖ Eventual Consistency
  - Understand Vector Clocks
- ❖ Spanner
  - Commit time via 2PC
- ❖ Basically at the short-answer question level



# *Column Stores*

- ❖ Performance advantages for OLAP (vs OLTP)
- ❖ Basically at the short-answer question level



# *MapReduce*

- ❖ Understand basic processing model
- ❖ May ask you to provide pseudocode to solve a particular problem in Map Reduce
- ❖ Look at the examples
  - especially graph processing
- ❖ Common error: solving everything in one mapper/reducer, i.e. not using parallelism

*That's All Folks!*

- ❖ See you in office hours ...
- ❖ Have a great semester break!

